




**BRENT OZAR**  
UNLIMITED®

# Fundamentals of Index Tuning

Up first: intros & logistics.

Module 1 Slide 1



**BRENT OZAR**  
UNLIMITED

99-05: dev, architect, DBA  
05-08: DBA, VM, SAN admin  
08-10: MCM, Quest Software  
Since: consulting DBA

[www.BrentOzar.com](http://www.BrentOzar.com)  
[Help@BrentOzar.com](mailto:Help@BrentOzar.com)



Module 1 Slide 2

## Introduce yourself in Slack:

	Developer	Development DBA	Production DBA
Write C#, Java code	Daily		
Build queries, tables	Daily	Sometimes	
Tune queries	Sometimes	Daily	
Design indexes		Daily	
Monitor performance		Daily	Sometimes
Troubleshoot outages			Daily
Manage backups, jobs			Daily
Install, config SQL			Sometimes
Install, config OS			Sometimes

Module 1 Slide 3



## Slack pro tips

Accidentally close your browser? Want to share screenshots? Lots of pro tips: [BrentOzar.com/slack](https://BrentOzar.com/slack)

To share code or T-SQL, click the + sign next to where you type text in, and choose “code or text snippet.”

To share files, upload at Imgur.com, paste URL here.

No direct messages please.

Keep the questions on-topic.

Module 1 Slide 4



## Instant Replay & lab scripts

For a year from your purchase, paid students can:

- Watch the class recordings
- Download the scripts
- Re-run the labs on your local machine

Problems or questions?

Leave a comment on the relevant module.

Module 1 Slide 5



## Today, we'll cover:

- How to pick order of keys in an index
- How to design indexes for a query without a plan
- How parameter values can change index needs
- Where to find index recommendations in plans, DMVs, and sp\_BlitzIndex
- How SQL Server's index recommendations are built, and why they're often wrong

Module 1 Slide 6



## We're using Stack Overflow data.

Open source, licensed with Creative Commons

XML dump: [archive.org/details/stackexchange](https://archive.org/details/stackexchange)

SQL Server db: [BrentOzar.com/go/querystack](https://BrentOzar.com/go/querystack)

I'm using the StackOverflow2013 database (50GB)

- You can use smaller or larger ones depending on your hardware
- Index creations may just take longer
- Exact logical page reads will be different

Module 1 Slide 7



## About my setup

SQL Server 2022, currently patched

SQL Server Management Studio 19

StackOverflow2013 50GB database

4 cores, 30GB RAM, SSDs  
(to make index creation fast)

Module 1 Slide 8





**BRENT OZAR**  
UNLIMITED®

# Fundamentals of Index Tuning

## Part 1: The WHERE Clause

Module 1 Slide 9

## This module's agenda

Building indexes for the WHERE clause

Understanding how selectivity determines key order

Learn how to visualize an index's contents



Module 1 Slide 10

# WHERE

## with 1 equality search

Module 1 Slide 11



**Design an index for this query.**

```
SELECT Id, DisplayName, Location  
FROM dbo.Users  
WHERE DisplayName = 'alex';
```

Module 1 Slide 12



## Check your logical reads before...

```
/* Design an index for this: */
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex';
```

200 %

Results Messages Execution plan

(3488 rows affected)

Table 'Users'. Scan count 1, logical reads 44530

Module 1 Slide 13



## Clippy has an idea:

```
101 /* Design an index for this: */
102 SELECT Id, DisplayName, Location
103 FROM dbo.Users
104 WHERE DisplayName = 'alex';
```

150 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT [Id],[DisplayName],[Location] FROM [dbo].[Users] WHERE [DisplayName]=@1

Missing Index (Impact 99.8523): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[Users] ([DisplayName])

Parallelism (Gather Streams)  
Cost: 2 %  
0.017s  
3488 of 3336 (104%)

Clustered Index Scan (Clustered)  
[Users].[PK\_Users\_Id]  
Cost: 98 %  
0.129s  
3488 of 3336 (104%)

I'm not cutting that screenshot off, either:  
Clippy didn't suggest that we include Location.

Module 1 Slide 14







## Clippy's hints are a gift.

But they're just a lucky byproduct of query plan optimization.

They're not Clippy's main job.

Later today, we'll explain how he builds them and why they're often wrong.

For now, focus on building your own.

Module 1 Slide 15



## I bet you could do better. Try it.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex';
```

Module 1 Slide 16





## Did you come up with this?

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex';

CREATE INDEX IX_DisplayName_Includes
ON dbo.Users(DisplayName)
INCLUDE (Location);
```

Module 1 Slide 17



## Should you include Id?

How to Think Like the Engine explained that the clustering key on a table is always included.

There's no extra cost whether you include it or not: it doesn't get stored twice.

I only include it if my query needs it in the output, and I suspect somebody's gonna come behind me and change the clustering key later.

Here, I'm fine either way.

Module 1 Slide 18



## Create the index. Does it get used?

```
CREATE INDEX IX_DisplayName_Includes
ON dbo.Users(DisplayName)
INCLUDE (Location);

SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex';
```

200 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT [Id],[DisplayName],[Location] FROM [dbo].[Users] WHERE [DisplayName]='alex'

Index Seek (NonClustered)  
[Users].[IX\_DisplayName\_Includes]  
Cost: 100 %  
0.001s  
3488 of  
3488 (100%)

SELECT  
Cost: 0 %

Yep!

## Does it do less logical reads?

Before: 44,530 logical page reads

Now, with the index:

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex';
```

100 %

Results Messages Execution plan

(3488 rows affected)

Table 'Users'. Scan count 1, logical reads 16

Great!



# WHERE

## with 2 equality searches

Module 1 Slide 21



**Now try this query.**

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location = 'Seattle, WA';
```

Module 1 Slide 22



## A couple of common solutions

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
    AND Location = 'Seattle, WA';

CREATE INDEX IX_DisplayName_Location
ON dbo.Users(DisplayName, Location);


CREATE INDEX IX_Location_DisplayName
ON dbo.Users(Location, DisplayName);
```

Module 1 Slide 23



## But remember that last index?

```
1 CREATE INDEX IX_DisplayName_Includes
2   ON dbo.Users(DisplayName) INCLUDE (Location);
3 GO
4 SELECT Id, DisplayName, Location
5   FROM dbo.Users
6  WHERE DisplayName = N'alex'
7  AND Location = N'Seattle, WA';
```



SQL Server can use that last index we created.

Sure, Location isn't sorted in order, but we can still:

1. Seek to Alex
2. Scan through them, checking their locations

Module 1 Slide 24



## How many reads does it do?

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
AND Location = 'Seattle, WA';
```

100 %

Results Messages Execution plan

(5 rows affected)

Table 'Users'. Scan count 1, logical reads 16

Just 16 8KB pages! That's pretty good.

Module 1 Slide 25



## Could it be even better? Sure.

Hover your mouse over the Index Seek in this plan:

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
AND Location = 'Seattle, WA';
```

100 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT [Id],[DisplayName],[Location] FROM [dbo].[Users] WHERE [Dis

**Index Seek (NonClustered)**  
Scan a particular range of rows from a nonclustered index.

Physical Operation	Index Seek
Logical Operation	Index Seek
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows Read	3488
Actual Number of Rows for All Executions	5
Actual Number of Batches	0
Estimated Operator Cost	0.0192425 (100%)
Estimated I/O Cost	0.0152487
Estimated Subtree Cost	0.0192425
Estimated CPU Cost	0.0039938
Estimated Number of Executions	1
Number of Executions	1
Estimated Number of Rows for All Executions	5.15641
Estimated Number of Rows to be Read	3488
Estimated Number of Rows Per Execution	5.15641
Estimated Row Size	42 B
Actual Rebinds	0
Actual Rewinds	0
Ordered	True
Node ID	0

**Index Seek (NonClustered)**  
[Users].[IX\_DisplayName\_Includes]  
Cost: 100 %  
0.000s  
5 of  
5 (100%)

Index Seek (NonClustered)	
Scan a particular range of rows from a nonclustered index.	
Physical Operation	Index Seek
Logical Operation	Index Seek
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Number of Rows Read	3488
Actual Number of Rows for All Executions	5
Actual Number of Batches	0
Estimated Operator Cost	0.0192425 (100%)
Estimated I/O Cost	0.0152487
Estimated Subtree Cost	0.0192425
Estimated CPU Cost	0.0039938
Estimated Number of Executions	1
Number of Executions	1
Estimated Number of Rows Per Execution	134.316
Estimated Number of Rows to be Read	3488
Estimated Row Size	42 B
Actual Rebinds	0
Actual Rewinds	0
Ordered	True
Node ID	0
Predicate	[StackOverflow2013].[dbo].[Users].[Location]=[@2]
Object	[StackOverflow2013].[dbo].[Users].[IX_DisplayName_Includes]
Output List	[StackOverflow2013].[dbo].[Users].Id, [StackOverflow2013].[dbo].[Users].DisplayName, [StackOverflow2013].[dbo].[Users].[Location]
Seek Predicates	Seek Keys[1]: Prefix: [StackOverflow2013].[dbo].[Users].DisplayName = Scalar Operator([@1])

## Decoding the popup

**Seek predicate** on DisplayName: we're able to jump straight to the Alexes.

**Predicate** on Location: but note that it doesn't have the word "seek" in front of it. This is called a residual predicate.



## What "index seek" really means

SQL Server seeked to a specific value on the index's **first column**.

It doesn't necessarily mean we seeked on subsequent columns: we may have scanned them.

Predicate	[StackOverflow2013].[dbo].[Users].[Location]=[@2]
Object	[StackOverflow2013].[dbo].[Users].[IX_DisplayName_Includes]
Output List	[StackOverflow2013].[dbo].[Users].Id, [StackOverflow2013].[dbo].[Users].DisplayName, [StackOverflow2013].[dbo].[Users].[Location]
Seek Predicates	Seek Keys[1]: Prefix: [StackOverflow2013].[dbo].[Users].DisplayName = Scalar Operator([@1])



Module 1 Slide 28

## To find out why, visualize the index

Write a query to exactly show what's inside this index:

- How it's ordered: ORDER BY
- What columns are on it: INCLUDE

```
CREATE INDEX IX_DisplayName_Includes
ON dbo.Users(DisplayName)
INCLUDE (Location);
```

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex';
```

Module 1 Slide 29



## If the visualization query is right

Then the visualization query's plan will ONLY have an index scan on that index – no other indexes, no sorts:

```
CREATE INDEX IX_DisplayName_Includes
ON dbo.Users(DisplayName)
INCLUDE (Location);

/* Visualize the index: */
SELECT DisplayName, Location
FROM dbo.Users
ORDER BY DisplayName;
```

200 %  
Results Messages Execution plan  
Query 1: Query cost (relative to the batch): 100%  
SELECT DisplayName, Location FROM dbo.Users ORDER BY DisplayName

Index Scan (NonClustered)  
[Users].[IX\_DisplayName\_Includes]  
Cost: 100 %  
0.590s  
2465713 of  
2465710 (100%)

SELECT  
Cost: 0 %

Module 1 Slide 3



## Gotchas with this technique

You may have to use an estimated plan:  
if your table is big, a select w/o a where can take days

You can use a WHERE to get results faster,  
but only on the first column in the index!

If you use a WHERE on the first column,  
your plan will be an index seek,  
but only a seek: no sorts. If you see sorts,  
the visualization query isn't right.

Module 1 Slide 31



## Testing our query & index

We're trying to see why this query has a residual  
(scan) predicate with our index:

```
CREATE INDEX IX_DisplayName_Includes  
ON dbo.Users(DisplayName)  
INCLUDE (Location);
```

```
SELECT Id, DisplayName, Location  
FROM dbo.Users  
WHERE DisplayName = 'alex';
```

Module 1 Slide 32



```

CREATE INDEX IX_DisplayName_
ON dbo.Users(DisplayName)
INCLUDE (Location);

/* Visualize the index: */
SELECT DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
ORDER BY DisplayName;

```

	DisplayName	Location
1	Alex	Springfield, MO
2	Alex	NULL
3	Alex	NULL
4	alex	NULL
5	Alex	Netherlands
6	Alex	
7	Alex	Seattle, WA
8	Alex	Brooklyn, NY
9	Alex	Oberhausen, Germany
10	Alex	Toulouse, France
11	Alex	Belarus
12	alex	United States
13	Alex	France
14	Alex	NULL
15	Alex	Wakefield, United Kingdom
16	Alex	Bucharest, Romania

## Visualize the index

The index is sorted on DisplayName, and Location is only included.

It's not sorted by Location.

So we can seek to Alex, but then we have to read all Alexes in all locations.

That's why we had 16 reads even with few Alexes in Seattle.



## Maybe your indexes are better. Let's create 'em and try 'em.

```

SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
AND Location = 'Seattle, WA';

CREATE INDEX IX_DisplayName_Location
ON dbo.Users(DisplayName, Location);

CREATE INDEX IX_Location_DisplayName
ON dbo.Users(Location, DisplayName);

```



```

/* Test 'em with index hints: */
SET STATISTICS IO ON;
GO
SELECT Id, DisplayName, Location
  FROM dbo.Users WITH (INDEX = 1) /* Clustered index scan */
 WHERE DisplayName = N'alex'
    AND Location = N'Seattle, WA';

SELECT Id, DisplayName, Location
  FROM dbo.Users WITH (INDEX = IX_DisplayName_Includes)
 WHERE DisplayName = N'alex'
    AND Location = N'Seattle, WA';

SELECT Id, DisplayName, Location
  FROM dbo.Users WITH (INDEX = IX_DisplayName_Location)
 WHERE DisplayName = N'alex'
    AND Location = N'Seattle, WA';

SELECT Id, DisplayName, Location
  FROM dbo.Users WITH (INDEX = IX_Location_DisplayName)
 WHERE DisplayName = N'alex'
    AND Location = N'Seattle, WA';

```

## Test 'em

I don't like index hints for long-term usage because your query will simply fail if the index disappears or is renamed. Hints are great for checking logical reads though.

## Survey says...

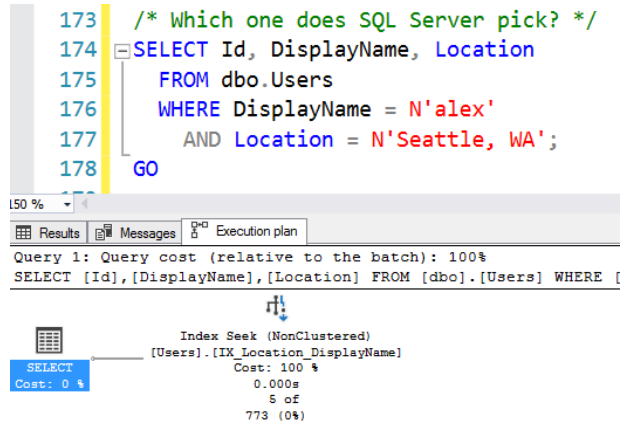
Index	Logical Reads
Clustered index (white pages)	45,184
IX_DisplayName_Includes	16
IX_DisplayName_Location	4
IX_Location_DisplayName	5

But don't quibble over a handful of logical reads.

All of the indexes are pretty good!



## Without a hint, what gets used?



SQL Server uses the Location, DisplayName index.

Don't read too much into that yet though: it doesn't mean it's the one you should keep.

Module 1 Slide 37



## Before you pick a winner...

You rarely see SQL Servers that only run one query.

What if this query still runs sometimes?

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex';
```

You wouldn't want one index for this query, and a different index for Location, DisplayName.

Module 1 Slide 38



# WHERE

## with both equality and inequality searches

Module 1 Slide 39



## What SQL Server calls equality

Equality searches	Inequality searches
=	<>, >, <, >=, <=
IS NULL	IS NOT NULL
IN (one value)	IN (two or more values)
	LIKE, NOT LIKE

Module 1 Slide 40



## Now try this query.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
AND Location <> 'Seattle, WA';
```

The <> is really important: it changes the game.

Module 1 Slide 41



## Think back to your 2 earlier indexes.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
AND Location <> 'Seattle, WA';

CREATE INDEX IX_DisplayName_Location
ON dbo.Users(DisplayName, Location);

CREATE INDEX IX_Location_DisplayName
ON dbo.Users(Location, DisplayName);
```

Module 1 Slide 42



## Your execution plan...

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA';

CREATE INDEX IX_DisplayName_Location
ON dbo.Users(DisplayName, Location);
```

“I’ll seek to DisplayName = Alex,  
then read through them, looking at their locations,  
returning all the ones who aren’t in Seattle.”

## But if the index starts with Location?

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA';

CREATE INDEX IX_Location_DisplayName
ON dbo.Users(Location, DisplayName);
```

“I can’t just jump to the Alexes.”

“I’m gonna have to scan most of the index.”



```

SET STATISTICS IO ON;
GO
SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = 1) /* Clustered index scan */
WHERE DisplayName = N'alex'
AND Location <> N'Seattle, WA';

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_DisplayName_Includes)
WHERE DisplayName = N'alex'
AND Location <> N'Seattle, WA';

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_DisplayName_Location)
WHERE DisplayName = N'alex'
AND Location <> N'Seattle, WA';

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_Location_DisplayName)
WHERE DisplayName = N'alex'
AND Location <> N'Seattle, WA';
GO

```

## Test 'em

Note the <>.

## Survey says...


Index	Logical Reads	Total Pages in the Index
Clustered index (white pages)	45,184	45,184
IX_DisplayName_Includes	16	12,577
IX_DisplayName_Location	13	12,701
IX_Location_DisplayName	4,566	13,183

Let's compare the actual execution plans for the bottom two.




# Remember, “seek” just means we seeked on the first column

Query 1: Query cost (relative to the batch): 0%  
SELECT Id, DisplayName, Location FROM dbo.Users WITH (INDEX = IX\_DisplayName\_Location)



Hover your mouse over each seek to see details...

Query 2: Query cost (relative to the batch): 100%  
SELECT Id, DisplayName, Location FROM dbo.Users WITH (INDEX = IX\_Location\_DisplayName)



Module 1 Slide 47



Index Seek (NonClustered)	
Scan a particular range of rows from a nonclustered index.	
Physical Operation	Index Seek
Logical Operation	Index Seek
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Number of Rows Read	703
Actual Number of Rows	703
Actual Number of Batches	0
Estimated I/O Cost	0.006088
Estimated Operator Cost	0.0071571 (100%)
Estimated CPU Cost	0.0010691
Estimated Subtree Cost	0.0071571
Estimated Number of Executions	1
Number of Executions	1
Estimated Number of Rows	829.184
Estimated Number of Rows to be Read	829.184
Estimated Row Size	42.8
Actual Rebinds	0
Actual Rewinds	0
Ordered	True
Node ID	0
Object	[StackOverflow2013].[dbo].[Users].[IX_DisplayName_Location]
Output List	[StackOverflow2013].[dbo].[Users].Id, [StackOverflow2013].[dbo].[Users].DisplayName, [StackOverflow2013].[dbo].[Users].Location
Seek Predicates	[1] Seek Keys[1]: Prefix: [StackOverflow2013].[dbo].[Users].DisplayName = Scalar Operator(N'alex'), End: [StackOverflow2013].[dbo].[Users].Location < Scalar Operator(N'Seattle, WA), [2] Seek Keys[1]: Prefix: [StackOverflow2013].[dbo].[Users].DisplayName = Scalar Operator(N'alex'), Start: [StackOverflow2013].[dbo].[Users].Location > Scalar Operator(N'Seattle, WA)

Index Seek (NonClustered)	
Scan a particular range of rows from a nonclustered index.	
Physical Operation	Index Seek
Logical Operation	Index Seek
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Number of Rows Read	586161
Actual Number of Rows	703
Actual Number of Batches	0
Estimated Operator Cost	2.89176 (100%)
Estimated I/O Cost	2.24683
Estimated Subtree Cost	2.89176
Estimated CPU Cost	0.644934
Estimated Number of Executions	1
Number of Executions	1
Estimated Number of Rows	829.184
Estimated Number of Rows to be Read	586161
Estimated Row Size	42.8
Actual Rebinds	0
Actual Rewinds	0
Ordered	True
Node ID	0
Predicate	[StackOverflow2013].[dbo].[Users].[DisplayName]=N'alex'
Output List	[StackOverflow2013].[dbo].[Users].[IX_Location_DisplayName]
Seek Predicates	[1] Seek Keys[1]: End: [StackOverflow2013].[dbo].[Users].Location < Scalar Operator(N'Seattle, WA), [2] Seek Keys[1]: Start: [StackOverflow2013].[dbo].[Users].Location > Scalar Operator(N'Seattle, WA)

## Different:

# of rows read

The one on the right has a residual predicate (Alex) that we can't seek to.

## So what's the lesson?

When you have both equality and inequality searches, you might think it's important to put the equality fields first in the index key order so that you can seek directly to the rows you want.

```
WHERE DisplayName = 'alex'  
      AND Location <> 'Seattle, WA';
```

But that's not necessarily true. Hold that thought.

Module 1 Slide 49



## Field order isn't about equality.

Field order is about **selectivity**, the ability to reduce the amount of work we're about to do.

Sometimes that's about reducing row counts by filtering down the number of rows we're going to pass on to the next operator in a plan.

Other times, it's about pre-sorting data to avoid sorts.

Module 1 Slide 50



# Selectivity

## How selective is Alex?

```
SELECT Id, DisplayName, Location  
FROM dbo.Users  
WHERE DisplayName = 'alex';
```

Total rows in table: 2,465,713

Rows where DisplayName = alex: **3,488 (0.14%)**





## Comparing selectivity

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
AND Location = 'Seattle, WA';
```

Total rows in table: 2,465,713

- DisplayName = alex: 3,488 (0.14%)
- Location = Seattle, WA: 3,345 (0.14%)
- Location <> Seattle, WA: 586,161 (24%)  
(remainder: nulls)

Module 1 Slide 53



**People think “selectivity” is  
the uniqueness of each value.**

(They’re totally wrong.)



```
SELECT * FROM dbo.Booze  
WHERE BrandName LIKE '%e%';
```



## What selectivity really means

It's not about how unique each row is by itself.

It's about your query,  
and how small a percentage of the table  
you're searching for.

When evaluating column order for indexes,  
don't think about how unique each column is.  
Think about what percentage you're searching for.



## Testing it out

When designing indexes for a query, craft a separate SELECT query for each filter in the WHERE clause, and test to see how selective it is.

```
SELECT Id, LastAccessDate, DownVotes  
FROM dbo.Users  
WHERE LastAccessDate <= GETDATE()  
AND Reputation = 0;
```



# Re-cap



## Recap

If your queries only have equality searches, key field order isn't all that important.

When you have inequality searches, though, key field order matters a LOT. The first fields in the key need to help reduce the amount of rows you scan.

Just because you see a “seek” doesn't mean you're seeking to a specific row: residual predicates indicate a seek, followed by a scan of an area of the index.

Module 1 Slide 59



## Picking key order

Fields in the WHERE clause usually\* need to go first

Selective ~~fields~~ query filters go first:  
reduce the amount of data you're searching through

Commonly filtered-on fields go first:  
maximize the number of queries that can use an index

\* We'll break this rule in the next module

Module 1 Slide 60



## Testing it out

When designing indexes for a query, craft a separate `SELECT` query for each filter in the `WHERE` clause, and test to see how selective it is.

Module 1 Slide 61



## Fundamentals of Index Tuning

Lab 1: let's see what you learned about `WHERE` filters.

Module 1 Slide 62

## Lab requirements

Download any Stack Overflow database:

- [BrentOzar.com/go/querystack](https://BrentOzar.com/go/querystack)
- I'm using the 50GB Stack Overflow 2013 (but any year is fine, even the 10GB one)

Desktop/laptop requirements:

- Any supported SQL Server version will work
- The faster your machine, the faster your indexes will get created

Module 1 Slide 63



## Working through the lab

Read the first query, execute it, do your work inline, creating and dropping indexes where directed

10 minutes: I'll start the lab with you

45 minutes: you work through the rest, asking questions in Slack as you go

30 minutes: I work through it onscreen

Module 1 Slide 64

