



BRENT OZAR
UNLIMITED®

Fundamentals of Index Tuning

Indexing for ORDER BY

Module 3 Slide 1

Agenda

How ORDER BY comes into play

Combining WHERE and ORDER BY

TOP exceptions: when ORDER BY goes first

How parameters affect key order



Module 3 Slide 2

I'd like to place an **ORDER BY** after two equality searches

Module 3 Slide 3



Bring some order around here

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location = 'Seattle, WA'
ORDER BY Reputation;
```

Module 3 Slide 4



Think back to your 2 earlier indexes.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location = 'Seattle, WA'
ORDER BY Reputation;

CREATE INDEX IX_DisplayName_Location
ON dbo.Users (DisplayName, Location);

CREATE INDEX IX_Location_DisplayName
ON dbo.Users (Location, DisplayName);
```

Module 3 Slide 5



Add new versions with Reputation

```
CREATE INDEX IX_DisplayName_Location_Reputation
ON dbo.Users (DisplayName, Location, Reputation);

CREATE INDEX IX_Location_DisplayName_Reputation
ON dbo.Users (Location, DisplayName, Reputation);

/* Plus a third idea: */
CREATE INDEX IX_Reputation_DisplayName_Location
ON dbo.Users (Reputation, DisplayName, Location);
```

Module 3 Slide 6



```

SET STATISTICS IO ON;
GO
SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = 1) /* Clustered index scan */
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_DisplayName_Location_Reputation)
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_Location_DisplayName_Reputation)
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;

SELECT Id, DisplayName, Location
FROM dbo.Users WITH (INDEX = IX_Reputation_DisplayName_Location)
WHERE DisplayName = N'alex'
AND Location = N'Seattle, WA'
ORDER BY Reputation;
GO

```

Test 'em

Survey says...

| Index | Logical Reads | Total Pages in the Index |
|------------------------------------|---------------|--------------------------|
| Clustered index (white pages) | 45,184 | 45,184 |
| IX_DisplayName_Location_Reputation | 4 | 13,995 |
| IX_Location_DisplayName_Reputation | 4 | 14,486 |
| IX_Reputation_DisplayName_Location | 13,996 | 13,996 |

Ouch. Putting reputation first meant no seeking at all, and we scanned the whole thing. (Still better than a table scan though.)



Which one does SQL Server pick?

```
90  /* Which one does SQL Server pick? */
91  SELECT Id, DisplayName, Location
92  FROM dbo.Users
93  WHERE DisplayName = 'alex'
94  AND Location = 'Seattle, WA'
95  ORDER BY Reputation;
96  GO
```

150 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT [Id],[DisplayName],[Location] FROM [dbo].[Users] WHERE

Index Seek (NonClustered)

[Users].[IX_DisplayName_Location_Re...

Cost: 100 %

0.000s

5 of

3225 (0%)

SELECT

Cost: 0 %

The one that leads with
DisplayName.

Module 3 Slide 9



ORDER BY

after an INequality search

Module 3 Slide 10



Your last query:

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location = 'Seattle, WA'
ORDER BY Reputation;
```

Module 3 Slide 11



Let's go anywhere BUT Seattle

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA'
ORDER BY Reputation;
```

What's the perfect index for this?
How selective is each part of the filter?

Module 3 Slide 12



Survey says...

| Index | Logical Reads | Total Pages in the Index |
|------------------------------------|---------------|--------------------------|
| Clustered index (white pages) | 45,184 | 45,184 |
| IX_DisplayName_Location_Reputation | 13 | 13,995 |
| IX_Location_DisplayName_Reputation | 4,864 | 14,486 |
| IX_Reputation_DisplayName_Location | 13,996 | 13,996 |

Ouch. Putting reputation first meant no seeking at all, and we scanned the whole thing. (Still better than a table scan though.)

Module 3 Slide 13



So the perfect index for it:

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA'
ORDER BY Reputation;
```

```
CREATE INDEX IX_DisplayName_Location_Reputation ON
dbo.Users(DisplayName, Location, Reputation);
```

Step 1: seek to Alex

Step 2: scan through,
returning everyone
EXCEPT Seattle

Step 3: read them out
sorted by Reputation,
except...they're not.

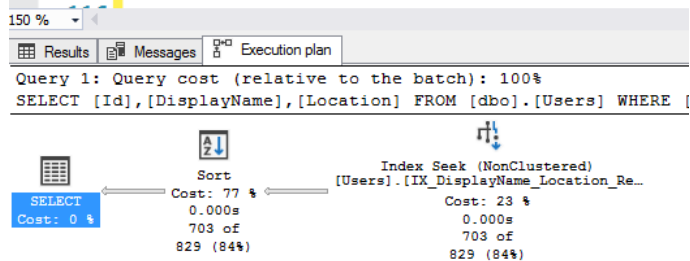
Module 3 Slide 14



Our index gets used, but...

```
110 SELECT Id, DisplayName, Location
111 FROM dbo.Users
112 WHERE DisplayName = 'alex'
113 AND Location <> 'Seattle, WA'
114 ORDER BY Reputation;
115 GO
```

The plan has a Sort even though the data in the index is sorted in order – isn't it?



Module 3 Slide 15



Write a query to visualize the index

```
CREATE INDEX IX_DisplayName_Location_Reputation ON  
dbo.Users(DisplayName, Location, Reputation);
```

```
SELECT DisplayName, Location, Reputation, Id  
FROM dbo.Users  
ORDER BY DisplayName, Location, Reputation;
```

| | DisplayName | Location | Reputation | Id |
|---|-------------|------------------------|------------|--------|
| 1 | guipoo | London, United Kingdom | 14747 | 389099 |
| 2 | µBio | California | 8999 | 9796 |
| 3 | µilad | Tehran, Iran | 5 | 136691 |
| 4 | 0_ | NULL | 48302 | 515054 |
| 5 | 0_o | NULL | 1 | 418884 |
| 6 | 0_o | NULL | 1 | 438437 |
| 7 | 0_o | NULL | 3 | 406169 |

Module 3 Slide 16

Seek down to Alex

```
SELECT DisplayName, Location, Reputation, Id
FROM dbo.Users
ORDER BY DisplayName, Location, Reputation;
```

| | DisplayName | Location | Reputation | Id |
|------|-------------|------------------|------------|--------|
| 8... | Aletubby | Japan | 2191 | 466432 |
| 8... | aletzo | Paris, France | 2099 | 269493 |
| 8... | aleung | Guangzhou, China | 4977 | 94148 |
| 8... | alewando | NULL | 16 | 166640 |
| 8... | Alex | NULL | 1 | 72895 |
| 8... | Alex | NULL | 1 | 198105 |
| 8... | Alex | NULL | 1 | 198345 |
| 8... | Alex | NULL | 1 | 202170 |
| 8... | Alex | NULL | 1 | 210507 |
| 8... | Alex | NULL | 1 | 213524 |
| 8... | Alex | NULL | 1 | 213698 |
| 8... | Alex | NULL | 1 | 216795 |
| 8... | Alex | NULL | 1 | 218333 |
| 8... | Alex | NULL | 1 | 228772 |
| 8... | Alex | NULL | 1 | 232257 |
| 8... | Alex | NULL | 1 | 233858 |
| 8... | Alex | NULL | 1 | 236336 |
| 8... | Alex | NULL | 1 | 237014 |
| 8... | Alex | NULL | 1 | 246305 |

Remember, we need them ordered by Reputation.

At first it looks like this will work, but...



Module 3 Slide 17

```
SELECT DisplayName, Location, Reputation, Id
FROM dbo.Users
ORDER BY DisplayName, Location, Reputation;
```

| | DisplayName | Location | Reputation | Id |
|------|-------------|-----------------------------|------------|--------|
| 9... | Alex | | 739 | 275559 |
| 9... | Alex | | 793 | 288023 |
| 9... | Alex | | 1133 | 94819 |
| 9... | Alex | | 1189 | 91631 |
| 9... | Alex | | 1818 | 179962 |
| 9... | Alex | | 13733 | 85185 |
| 9... | Alex | "The Cloud" | 435 | 291796 |
| 9... | Alex | Ahmadabad, India | 951 | 499711 |
| 9... | Alex | Ahokio, NC | 3 | 88242 |
| 9... | Alex | Amsterdam, The Nethe... | 121 | 464950 |
| 9... | Alex | Annecey, France | 96 | 383173 |
| 9... | alex | Athens, Greece | 563 | 423581 |
| 9... | Alex | Atlanta, GA | 1490 | 207090 |
| 9... | Alex | Auckland, New Zealand | 6 | 474011 |
| 9... | Alex | Austin, TX | 148 | 485381 |
| 9... | Alex | Austin, TX | 240 | 177746 |
| 9... | Alex | Australia | 817 | 167011 |
| 9... | Alex | Australia | 1473 | 211869 |
| 9... | Alex | Austria | 1935 | 275837 |
| 9... | alex | Barcelona, Spain | 4567 | 26787 |
| 9... | Alex | Bay Area, CA, United ... | 188 | 389050 |
| 9... | Alex | Belarus | 137 | 515369 |
| 9... | Alex | Belarus | 593 | 19081 |
| 9... | Alex | Belgium | 81 | 452522 |
| 9... | alex | Bermuda | 1 | 363997 |
| 9... | Alex | Billingshurst, United Ki... | 1364 | 181739 |
| 9... | Alex | Boca Raton, FL | 5726 | 257404 |
| 9... | Alex | Borås, Sweden | 71 | 464570 |
| 9... | Alex | Boston, MA | 1130 | 499643 |
| 9... | Alex | Boston, MA | 1965 | 417291 |
| 9... | Alex | Brasilia, Brazil | 1855 | 312808 |
| 9... | Alex | Breda, The Netherlands | 329 | 236813 |
| 9... | Alex | Brighton, United Kingd... | 53 | 289608 |
| 9... | Alex | Brooklyn, NY | 1770 | 12204 |

Reputation isn't sorted.

We're going to skip everyone who isn't in Seattle.

That means we need all the Alexes on this screen, plus more.

And they're not sorted by Reputation.

The fact that Reputation is "sorted" isn't helping here.



Ordering Reputation doesn't help.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
AND Location <> 'Seattle, WA'
ORDER BY Reputation;
```

```
CREATE INDEX IX_DisplayName_Location_Reputation ON
dbo.Users(DisplayName, Location, Reputation);
```

Step 1: seek to Alex

Step 2: scan through,
returning everyone
EXCEPT Seattle

Step 3: read them out
sorted by Reputation,
except...they're not.

Module 3 Slide 19



To prove it, create another index:

```
CREATE INDEX IX_DisplayName_Location_Reputation
ON dbo.Users
(DisplayName, Location, Reputation);
```

```
CREATE INDEX IX_DisplayName_Location_Includes ON
dbo.Users
(DisplayName, Location) INCLUDE (Reputation);
```

Module 3 Slide 20



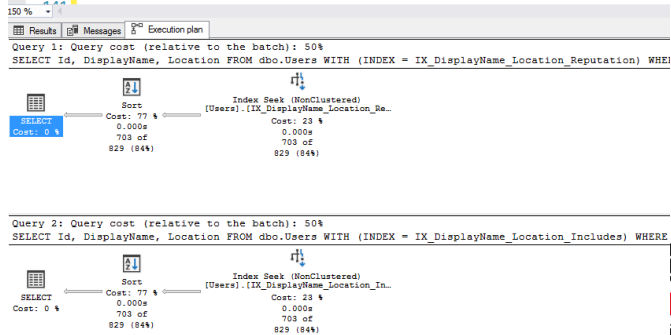
They both get the same plan

```
129 SELECT Id, DisplayName, Location
130 FROM dbo.Users WITH (INDEX = IX_DisplayName_Location_Reputation)
131 WHERE DisplayName = 'alex'
132 AND Location <> 'Seattle, WA'
133 ORDER BY Reputation;
134
135 SELECT Id, DisplayName, Location
136 FROM dbo.Users WITH (INDEX = IX_DisplayName_Location_Includes)
137 WHERE DisplayName = 'alex'
138 AND Location <> 'Seattle, WA'
139 ORDER BY Reputation;
140 GO
```

Use an index hint to test both indexes separately.

Both do the sort.

And both have the same number of logical reads.



Inequality searches make it tricky.

```
WHERE DisplayName = 'alex'
      AND Location <> 'Seattle, WA'
ORDER BY Reputation;
```

After you do an inequality search on a field, the sorting of subsequent fields in the index are usually less useful.

(That's a mouthful.)



Putting Reputation SECOND helps.

```
SELECT Id, DisplayName, Location
FROM dbo.Users
WHERE DisplayName = 'alex'
    AND Location <> 'Seattle, WA'
ORDER BY Reputation;

CREATE INDEX IX_DisplayName_Reputation_Location ON
dbo.Users (DisplayName, Reputation, Location);
```

Step 1: seek to Alex

Step 2: the sort isn't
needed: they're sorted

Step 3: Skip the users
who aren't in Seattle

Module 3 Slide 23



The sort is gone with this trick.

```
144 /* Promote Reputation one level: */
145 CREATE INDEX IX_DisplayName_Reputation_Location
146 ON dbo.Users (DisplayName, Reputation, Location);
147 GO
148
149 /* And the sort is gone: */
150 SELECT Id, DisplayName, Location
151 FROM dbo.Users WITH (INDEX = IX_DisplayName_Reputation_Location)
152 WHERE DisplayName = 'alex'
153 AND Location <> 'Seattle, WA'
154 ORDER BY Reputation;
```

Query 1: Query cost (relative to the batch): 100%

SELECT Id, DisplayName, Location FROM dbo.Users WITH (INDEX = IX_DisplayName_Reputation_Location) WHERE DisplayName = 'alex' AND Location <> 'Seattle, WA' ORDER BY Reputation;

Execution plan

Index Seek (NonClustered)

(Users).IX_DisplayName_Reputation_Location

Cost: 100 %

0.000s

703 of 829 (84%)

Obscure trick. To get it, key on:

1. Equality fields, then
2. Sort fields, then
3. Inequality fields

Module 3 Slide 24



SQL Server picks it, too.

```
157  /* Which one does SQL Server pick? */
158  SELECT Id, DisplayName, Location
159  FROM dbo.Users
160  WHERE DisplayName = 'alex'
161        AND Location <> 'Seattle, WA'
162  ORDER BY Reputation;
163  GO
164
```

150 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT [Id],[DisplayName],[Location] FROM [dbo].[Users] WHERE [Disp

Index Seek (NonClustered)
(Users).[IX_DisplayName_Reputation_...
Cost: 100 %
0.000s
703 of
829 (84%)

SELECT
Cost: 0 %

If we don't hint the query, here it picks the index that removes the sort.



Module 3 Slide 25

What we've learned so far

Indexes help by pre-sorting rows to prep them for:

- WHERE: finding the rows we want
- ORDER BY: sorting them on the way out the door
- GROUP BY, FROM, JOINS, CTEs:
more on these later

And so far, it kinda seems like you want to put keys in that same order: WHERE first, then ORDER BY. But that's not exactly how it works.



Module 3 Slide 26

S TOP

me if you've heard this one before

Module 3 Slide 27



Start by dropping your indexes.

We're going to tackle a new set of queries, and I don't want to confuse SQL Server's hints with existing indexes.

```
EXEC DropIndexes;
```

Get the code:

BrentOzar.com/go/dropindexes

Module 3 Slide 28



Design an index for this:

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

Module 3 Slide 29



Which field should we lead with?

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

```
CREATE INDEX IX_Reputation_CreationDate
ON dbo.Users(Reputation, CreationDate);
```

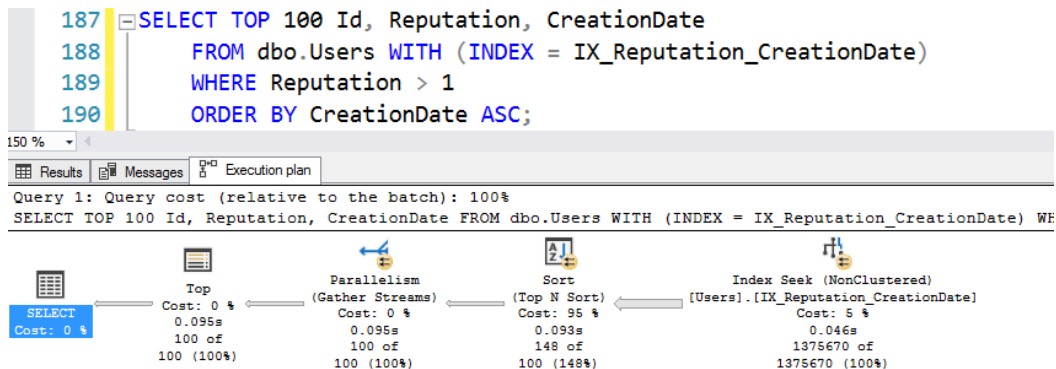
```
CREATE INDEX IX_CreationDate_Reputation
ON dbo.Users(CreationDate, Reputation);
```

Module 3 Slide 30



If we lead with Reputation...

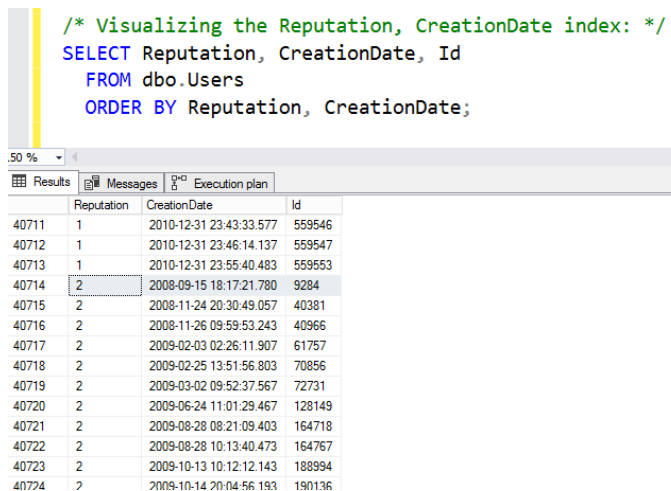
We seek to 2, but then we find 1.4M users that match!
We have to sort 'em all by CreationDate.



Module 3 Slide 31



Visualize the index contents



Module 3 Slide 32



When the index is on Reputation, CreationDate, we can seek to 2, but...are the first 10 users we find the lowest CreationDates overall?

Or just the lowest for Reputation = 2?

Easier way to see it

```
/* Visualizing the Reputation, CreationDate, Id
SELECT Reputation, CreationDate, Id
FROM dbo.Users
ORDER BY Reputation, CreationDate;
```

| | Reputation | CreationDate | Id |
|--------|------------|-------------------------|--------|
| 277696 | 3819 | 2010-03-25 08:32:56.610 | 301514 |
| 277697 | 3819 | 2010-06-14 02:28:25.290 | 365977 |
| 277698 | 3819 | 2010-07-19 15:48:21.980 | 395975 |
| 277699 | 3820 | 2008-09-24 08:03:37.333 | 21537 |
| 277700 | 3820 | 2009-07-30 09:26:28.910 | 147695 |
| 277701 | 3820 | 2010-10-08 09:07:30.453 | 470062 |
| 277702 | 3821 | 2008-10-05 08:08:20.540 | 25234 |
| 277703 | 3821 | 2008-12-30 00:22:36.087 | 50025 |
| 277704 | 3821 | 2010-04-12 15:33:11.823 | 314670 |
| 277705 | 3821 | 2010-12-29 08:26:55.093 | 556899 |
| 277706 | 3822 | 2009-03-25 01:36:44.977 | 82333 |
| 277707 | 3822 | 2010-01-22 20:12:36.790 | 257065 |
| 277708 | 3822 | 2010-09-01 10:51:24.297 | 436853 |

It's more obvious when we page down to higher Reputation numbers.

The CreationDate keeps resetting with each new Reputation.

The sort on the second field is less useful when we're scanning.

Module 3 Slide 33



What if we lead with CreationDate?

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

```
CREATE INDEX IX_Reputation_CreationDate
ON dbo.Users(Reputation, CreationDate);
```



```
CREATE INDEX IX_CreationDate_Reputation
ON dbo.Users(CreationDate, Reputation);
```

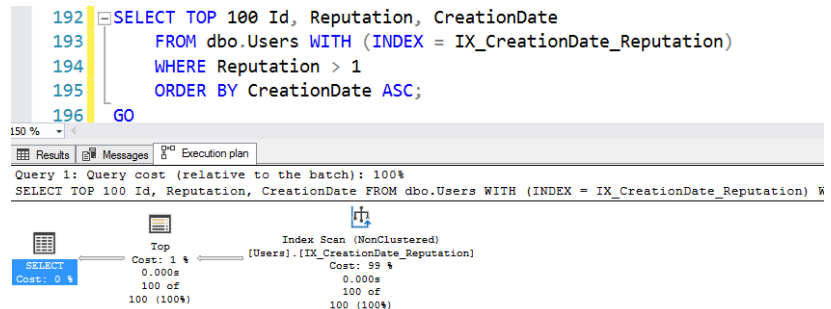
Module 3 Slide 34



We “scan” the index, but...

Remember from How to Think Like the Engine: scan just means we start at one end of the index, and we read until we find the rows that match.

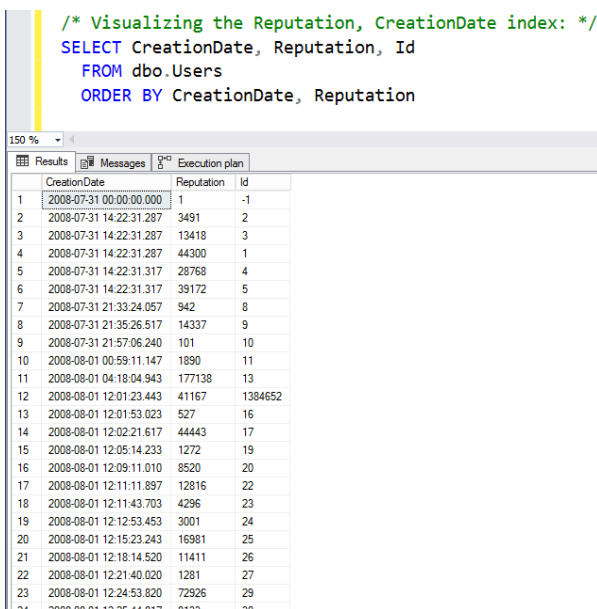
And there’s no sort! The data is already sorted.



Module 3 Slide 35



Visualize the index contents



When the index is on CreationDate, Reputation, we start reading, looking for 100 users with Reputation > 1.

They almost all match!

As soon as we read 100 rows that match, we’re done. No need to scan the whole index.



Survey says...

| Index | Logical Reads | Total Pages in the Index |
|-------------------------------|---------------|--------------------------|
| Clustered index (white pages) | 45,184 | 45,184 |
| IX_Reputation_CreationDate | 3,805 | 6,812 |
| IX_CreationDate_Reputation | 3 | 6,817 |

Module 3 Slide 37



In this case, the ORDER BY field should go first in the index.

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

```
CREATE INDEX IX_Reputation_CreationDate
ON dbo.Users(Reputation, CreationDate);
```

```
CREATE INDEX IX_CreationDate_Reputation
ON dbo.Users(CreationDate, Reputation);
```

Module 3 Slide 38



Remember selectivity?

TOP is kinda like a WHERE clause.

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1
ORDER BY CreationDate ASC;
```

That's kinda like saying:

```
SELECT stuff
FROM dbo.Users
WHERE (user is in the top ~100) by CreationDate
```



So let's keep just this one for now

```
DropIndexes;  
  
CREATE INDEX IX_CreationDate_Reputation  
ON dbo.Users(CreationDate, Reputation);
```

Let's say we decided to just keep this one.

Module 3 Slide 41



Now run this.

```
SELECT TOP 100 Id, Reputation, CreationDate  
FROM dbo.Users  
WHERE Reputation > 1000000  
ORDER BY CreationDate ASC;
```

There aren't a lot of rows with
Reputation > 1,000,000.

Module 3 Slide 42



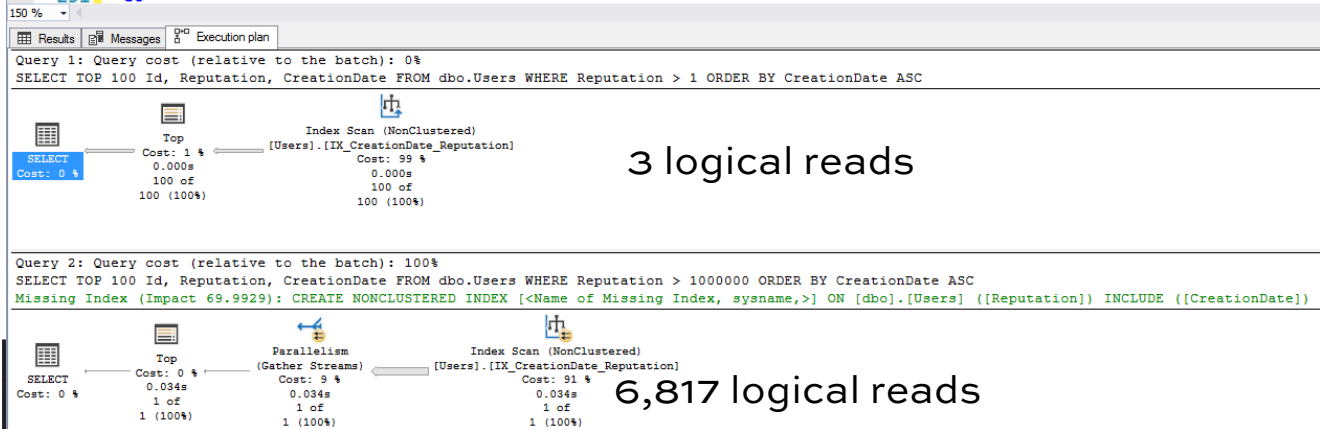
```

220  /* The original query: */
221  SELECT TOP 100 Id, Reputation, CreationDate
222  FROM dbo.Users
223  WHERE Reputation > 1
224  ORDER BY CreationDate ASC;
225
226  /* The new one looking for Jon Skeet: */
227  SELECT TOP 100 Id, Reputation, CreationDate
228  FROM dbo.Users
229  WHERE Reputation > 1000000
230  ORDER BY CreationDate ASC;
231  GO

```

Both old & new queries use the index...

But the index isn't as good of a fit for the second query. Why?



Jon Skeet isn't in the first 100.

```

SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1000000
ORDER BY CreationDate ASC;

```


The TOP 100 by CreationDate is only selective IF the person you're looking for is in that list.

In this case, WHERE Reputation > 1000000 is much more selective – that should go first.





**Indexing requires
compromises and choices.**



Index for the WHERE
to reduce reads

Index for the ORDER BY
to reduce sorts

**Indexing requires
compromises and choices.**

These are just 2 inequality searches.

```
SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > 1000000
ORDER BY CreationDate ASC;
```

It comes down to:

- Which ones are the most selective
- And whether you want to cut reads or cut sorts
- Which parameters run the most often

Module 3 Slide 47



Say this is a stored procedure.

```
CREATE PROC usp_SearchUsers
    @SearchReputation INT AS

SELECT TOP 100 Id, Reputation, CreationDate
FROM dbo.Users
WHERE Reputation > @SearchReputation
ORDER BY CreationDate ASC;
GO
```

Module 3 Slide 48




```
CREATE PROC usp_SearchUsers  
    @SearchReputation INT AS  
  
SELECT TOP 100 Id, Reputation, CreationDate  
    FROM dbo.Users  
    WHERE Reputation > @SearchReputation  
    ORDER BY CreationDate ASC;  
GO
```

When @SearchReputation = 1, lots of data matches,
so it's better to index on CreationDate, then Reputation.

When @SearchReputation = 1,000,000, then only 1 person matches,
so it's better to index on Reputation, then CreationDate.



Re-cap

Recap

If your WHERE clause is filtering just for equalities, then add the ORDER BY fields into the index key, and the index will handle all the sorting for you.

Out here in the real world, though, your query will have a mix of equality and inequalities.

Different parameter values affect key order too.

Our goal: get a good enough combination of keys to cover as many queries as practical.

Module 3 Slide 51



BRENT OZAR
UNLIMITED®

Fundamentals of Index Tuning

Lab 2: let's see what you learned about ORDER BY.

Module 3 Slide 52

Lab requirements

Download any Stack Overflow database:

- BrentOzar.com/go/querystack
- I'm using the 50GB Stack Overflow 2013 (but any year is fine, even the 10GB one)

Desktop/laptop requirements:

- Any supported SQL Server version will work
- The faster your machine, the faster your indexes will get created

Module 3 Slide 53



Working through the lab

Read the first query, execute it, do your work inline, taking notes as you go

2 hours: you work through the lab, asking questions in Slack as you go, and get lunch (either lunch first, or after your work)

The live stream will be off during lunch.

After lunch: I work through it onscreen

Module 3 Slide 54

