



**BRENT OZAR**  
UNLIMITED®

# Fundamentals of Index Tuning

The built-in missing index recommendations

# Index hints are a gift.

They're a byproduct of plan compilation, but they're not the main deliverable.

- Shown in execution plans
- Tracked over time in DMVs like `sys.dm_db_missing_index_details`
- Shown in tools like `sp_BlitzIndex`



## But they're not perfect gifts.

Suggests super wide indexes

Doesn't de-duplicate requests

Don't get thrown for all queries

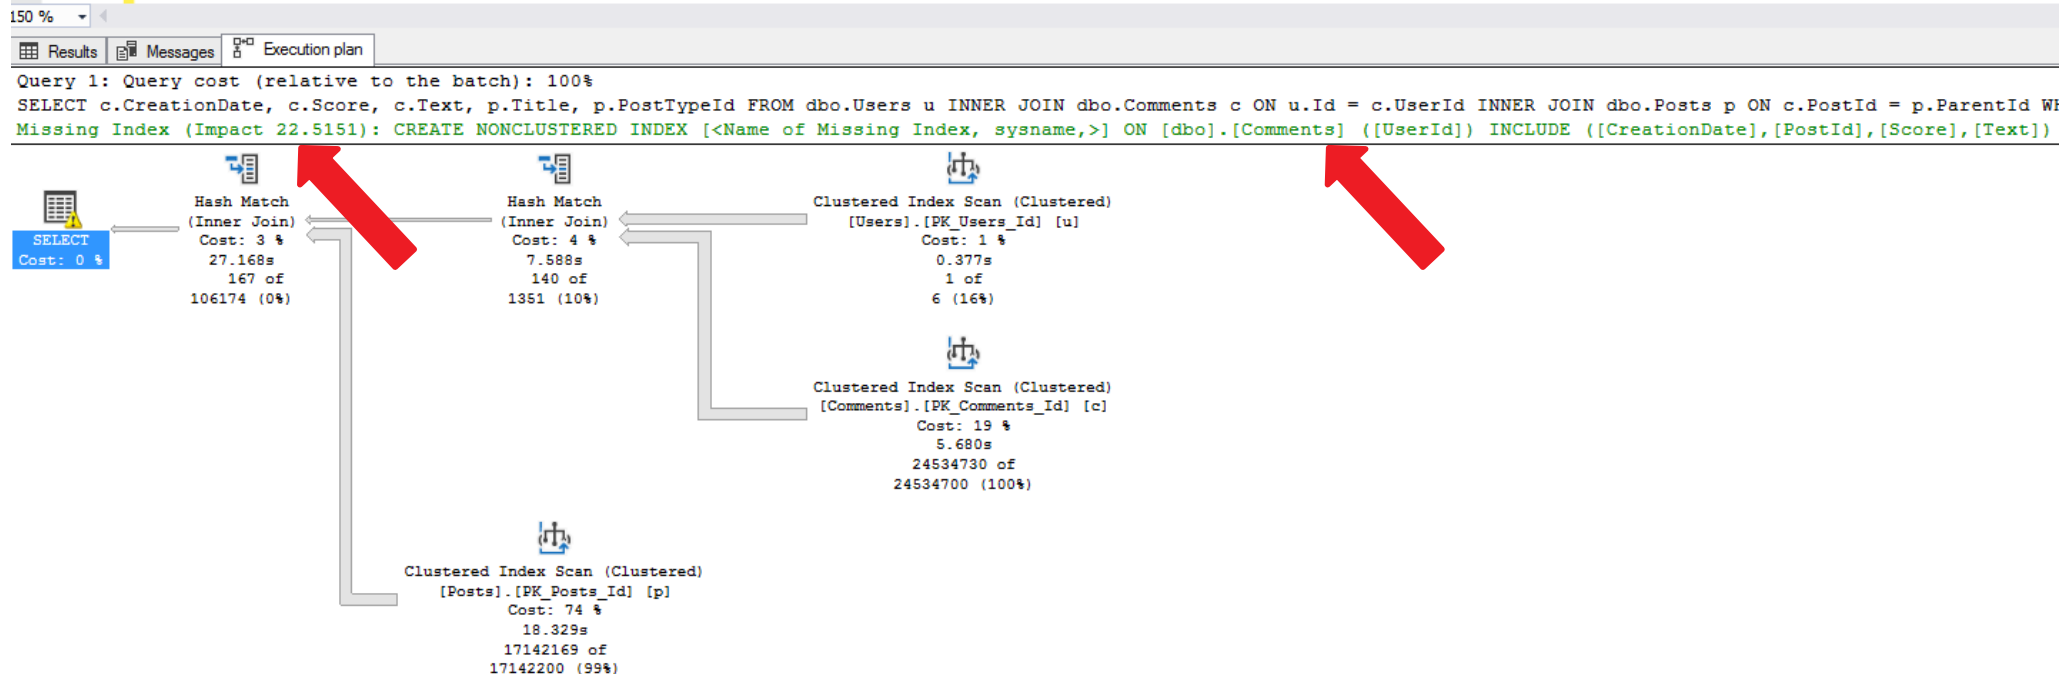
Get cleared at tricky times

Doesn't recommend filtered,  
columnstore, indexed views,  
XML, spatial, in-memory OLTP



# In plans, only the first one shows

```
29  /* What missing index does this ask for? Are you sure? */
30  SELECT c.CreationDate, c.Score, c.Text, p.Title, p.PostTypeId
31  FROM dbo.Users u
32  INNER JOIN dbo.Comments c ON u.Id = c.UserId
33  INNER JOIN dbo.Posts p ON c.PostId = p.ParentId
34  WHERE u.DisplayName = 'Brent Ozar';
35  GO
```



```

1 |<?xml version="1.0" encoding="utf-16"?>
2 |<ShowPlanXML xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xmlns:xsd="http://www.w3.org/2001/XM
3 |  <BatchSequence>
4 |    <Batch>
5 |      <Statements>
6 |        <StmtSimple StatementCompId="1" StatementEstRows="106174" StatementId="1" StatementOptmLevel="F
7 |          <StatementSetOptions ANSI_NULLS="true" ANSI_PADDING="true" ANSI_WARNINGS="true" ARITHABORT="t
8 |          <QueryPlan DegreeOfParallelism="1" MemoryGrant="46248" CachedPlanSize="88"
9 |            <MissingIndexes>
10 |              <MissingIndexGroup Impact="22.5151">
11 |                <MissingIndex Database="[StackOverflow2013]" Schema="[dbo]" Table="[Comments]">
12 |                  <ColumnGroup Usage="EQUALITY">
13 |                    <Column Name="[UserId]" ColumnId="6" />
14 |                  </ColumnGroup>
15 |                  <ColumnGroup Usage="INCLUDE">
16 |                    <Column Name="[CreationDate]" ColumnId="2" />
17 |                    <Column Name="[PostId]" ColumnId="3" />
18 |                    <Column Name="[Score]" ColumnId="4" />
19 |                    <Column Name="[Text]" ColumnId="5" />
20 |                  </ColumnGroup>
21 |                </MissingIndex>
22 |              </MissingIndexGroup>
23 |              <MissingIndexGroup Impact="76.6096">
24 |                <MissingIndex Database="[StackOverflow2013]" Schema="[dbo]" Table="[Posts]">
25 |                  <ColumnGroup Usage="EQUALITY">
26 |                    <Column Name="[ParentId]" ColumnId="15" />
27 |                  </ColumnGroup>
28 |                  <ColumnGroup Usage="INCLUDE">
29 |                    <Column Name="[PostTypeId]" ColumnId="16" />
30 |                    <Column Name="[Title]" ColumnId="19" />
31 |                  </ColumnGroup>
32 |                </MissingIndex>
33 |              </MissingIndexGroup>
34 |            </MissingIndexes>
35 |          <Warnings>

```

SSMS shows the FIRST one

But not the rest



```

1 |<?xml version="1.0" encoding="utf-16"?>
2 |<ShowPlanXML xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xmlns:xsd="http://www.w3.org/2001/XM
3 |   <BatchSequence>
4 |     <Batch>
5 |       <Statements>
6 |         <StmtSimple StatementCompId="1" StatementEstRows="106174" StatementId="1" StatementOptmLevel="F
7 |           <StatementSetOptions ANSI_NULLS="true" ANSI_PADDING="true" ANSI_WARNINGS="true" ARITHABORT="t
8 |           <QueryPlan DegreeOfParallelism="1" MemoryGrant="46248" CachedPlanSize="88"
9 |             <MissingIndexes>
10 |               <MissingIndexGroup Impact="22.5151">
11 |                 <MissingIndex Database="[StackOverflow2013]" Schema="[dbo]" Table="[Comments]">
12 |                   <ColumnGroup Usage="EQUALITY">
13 |                     <Column Name="[UserId]" ColumnId="6" />
14 |                   </ColumnGroup>
15 |                   <ColumnGroup Usage="INCLUDE">
16 |                     <Column Name="[CreationDate]" ColumnId="2" />
17 |                     <Column Name="[PostId]" ColumnId="3" />
                     <Column Name="[Score]" ColumnId="4" />
                     <Column Name="[Text]" ColumnId="5" />
                   </ColumnGroup>
                 </MissingIndex>
               </MissingIndexGroup>
               <MissingIndexGroup Impact="76.6096">
                 <MissingIndex Database="[StackOverflow2013]" Schema="[dbo]" Table="[Posts]">
                   <ColumnGroup Usage="EQUALITY">
                     <Column Name="[ParentId]" ColumnId="15" />
                   </ColumnGroup>
                   <ColumnGroup Usage="INCLUDE">
                     <Column Name="[PostTypeId]" ColumnId="16" />
                     <Column Name="[Title]" ColumnId="19" />
                   </ColumnGroup>
                 </MissingIndex>
               </MissingIndexGroup>
             </MissingIndexes>
           </Warnings>

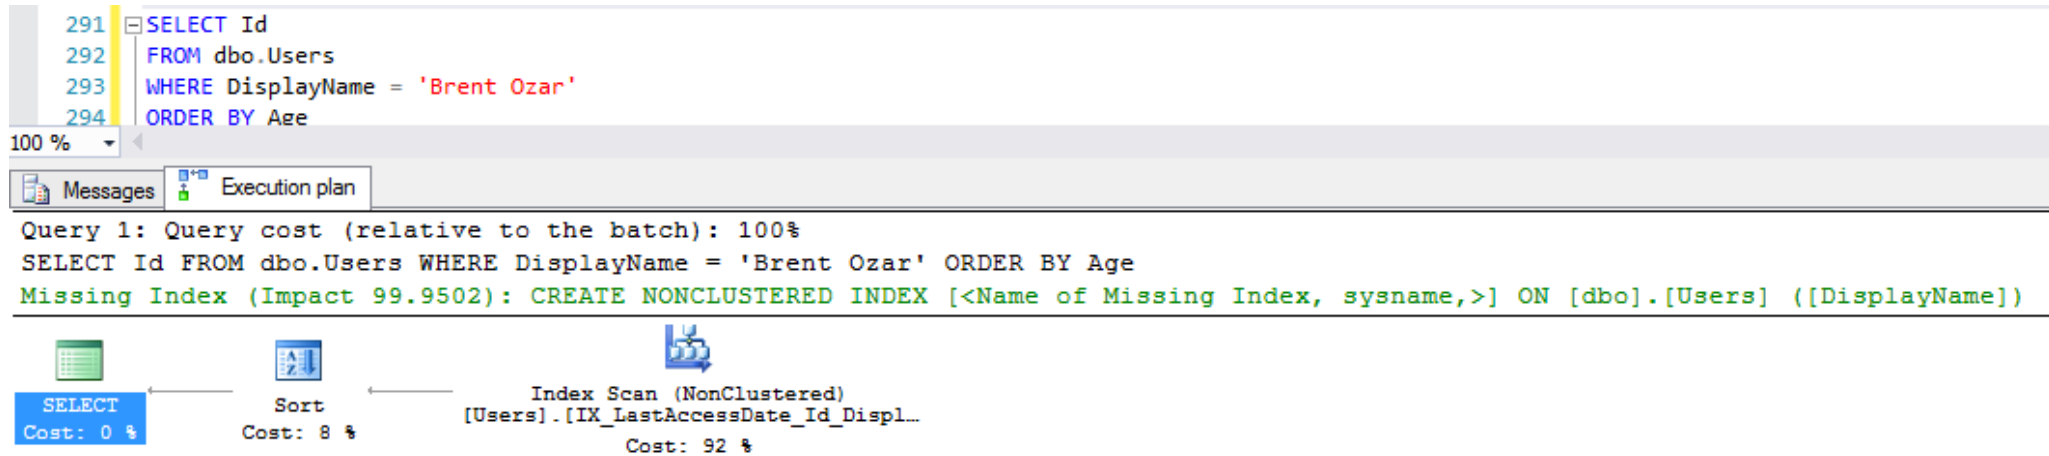
```

SSMS shows the FIRST one

But not the rest



# We're sorting by Age, but...



The screenshot shows a SQL query in the 'Messages' pane and its corresponding execution plan in the 'Execution plan' pane.

**Query 1:** Query cost (relative to the batch): 100%

```
SELECT Id FROM dbo.Users WHERE DisplayName = 'Brent Ozar' ORDER BY Age
```

**Missing Index (Impact 99.9502):** CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[Users] ([DisplayName])

**Execution Plan:**


- Index Scan (NonClustered)** [Users].[IX\_LastAccessDate\_Id\_Displ...]  
Cost: 92 %
- Sort**  
Cost: 8 %
- SELECT**  
Cost: 0 %

## Limitations of the Missing Indexes Feature

SQL Server 2008 R2 | [Other Versions](#) | This topic has not yet been rated - [Rate this topic](#)

The missing index feature has the following limitations:

- It is not intended to fine tune an indexing configuration.





## Limitations of the missing index feature

When the query optimizer generates a query plan, it analyzes what the best indexes are for a particular filter condition. If the best indexes don't exist, the query optimizer still generates a query plan using the least-costly access methods available, but also stores information about these indexes. The missing indexes feature enables you to access that information about best possible indexes so you can decide whether they should be implemented.

Query optimization is a time sensitive process, so there are limitations to the missing index feature. Limitations include:

- Missing index suggestions are based on estimates made during the optimization of a single query, prior to query execution. Missing index suggestions aren't tested or updated after query execution.
- The missing index feature suggests only nonclustered disk-based rowstore indexes. **Unique** and **filtered indexes** aren't suggested.
- **Key columns** are suggested, but the suggestion doesn't specify an order for those columns. For information on ordering columns, see the [Apply missing index suggestions](#) section of this article.
- **Included columns** are suggested, but SQL Server performs no cost-benefit analysis regarding the size of the resulting index when a large number of included columns are suggested.
- Missing index requests may offer similar variations of indexes on the same table and column(s) across queries. It's important to [review index suggestions and combine where possible](#).
- Suggestions aren't made for trivial query plans.
- Cost information is less accurate for queries involving only inequality predicates.
- Suggestions are gathered for a maximum of 500 missing index groups. After this threshold is reached, no more missing index group data is gathered.

Due to these limitations, missing index suggestions are best treated as one of several sources of information when performing index analysis, design, tuning, and testing. Missing index suggestions are not prescriptions to create indexes exactly as suggested.

*And these apply to both the missing indexes in query plans, AND missing index DMVs.*





Query optimization is a time sensitive process, so there are limitations to the missing index feature.


Limitations include:

- Missing index suggestions are based on estimates made during the optimization of a single query, prior to query execution. Missing index suggestions aren't tested or updated after query execution.
- The missing index feature suggests only nonclustered disk-based rowstore indexes. [Unique](#) and [filtered indexes](#) aren't suggested.
- [Key columns](#) are suggested, but the suggestion doesn't specify an order for those columns. For information on ordering columns, see the [Apply missing index suggestions](#) section of this article.
- [Included columns](#) are suggested, but SQL Server performs no cost-benefit analysis regarding the size of the resulting index when a large number of included columns are suggested.
- Missing index requests may offer similar variations of indexes on the same table and column(s) across queries. It's important to [review index suggestions and combine where possible](#).
- Suggestions aren't made for trivial query plans.
- Cost information is less accurate for queries involving only inequality predicates.
- Suggestions are gathered for a maximum of 500 missing index groups. After this threshold is reached, no more missing index group data is gathered.



Query optimization is a time sensitive process, so there are limitations to the missing index feature. Limitations include:

- Missing index suggestions are based on estimates made during the optimization of a single query, prior to query execution. Missing index suggestions aren't tested or updated after query execution.
- The missing index feature suggests only nonclustered disk-based rowstore indexes. Unique and filtered indexes aren't suggested.
- **Key columns** are suggested, but the suggestion doesn't specify an order for those columns. For information on ordering columns, see the [Apply missing index suggestions](#) section of this article.
- **Included columns** are suggested, but SQL Server performs no cost-benefit analysis regarding the size of the resulting index when a large number of included columns are suggested.
- Missing index requests may offer similar variations of indexes on the same table and column(s) across queries. It's important to [review index suggestions and combine where possible](#).
- Suggestions aren't made for trivial query plans.
- Cost information is less accurate for queries involving only inequality predicates.
- Suggestions are gathered for a maximum of 500 missing index groups. After this threshold is reached, no more missing index group data is gathered.



**Let's see how  
he does it.**



# Create table w/10M identical rows

```
19 CREATE TABLE dbo.DiningRoom
20     (FirstColumn INT,
21       SecondColumn INT,
22       ThirdColumn INT,
23       FourthColumn INT,
24       FifthColumn INT,
25       SixthColumn INT
26     );
27 INSERT INTO dbo.DiningRoom
28     (FirstColumn, SecondColumn, ThirdColumn, FourthColumn, FifthColumn, SixthColumn)
29     SELECT TOP 10000000 1, 1, 1, 1, 1, 1
30     FROM sys.all_columns ac1
31     CROSS JOIN sys.all_columns ac2
32     CROSS JOIN sys.all_columns ac3;
33 GO
```

200 %

Results Messages

	FirstColumn	SecondColumn	ThirdColumn	FourthColumn	FifthColumn	SixthColumn
1	1	1	1	1	1	1
2	1	1	1	1	1	1
3	1	1	1	1	1	1
4	1	1	1	1	1	1

# Single-column equality search

```
44  /* Turn on actual execution plans, and check the missing index requests: */
45  SET STATISTICS TIME, IO ON;
46  GO
47  SELECT 'Hi Mom!'
48  FROM dbo.DiningRoom
49  WHERE FirstColumn = 0;
```


200 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

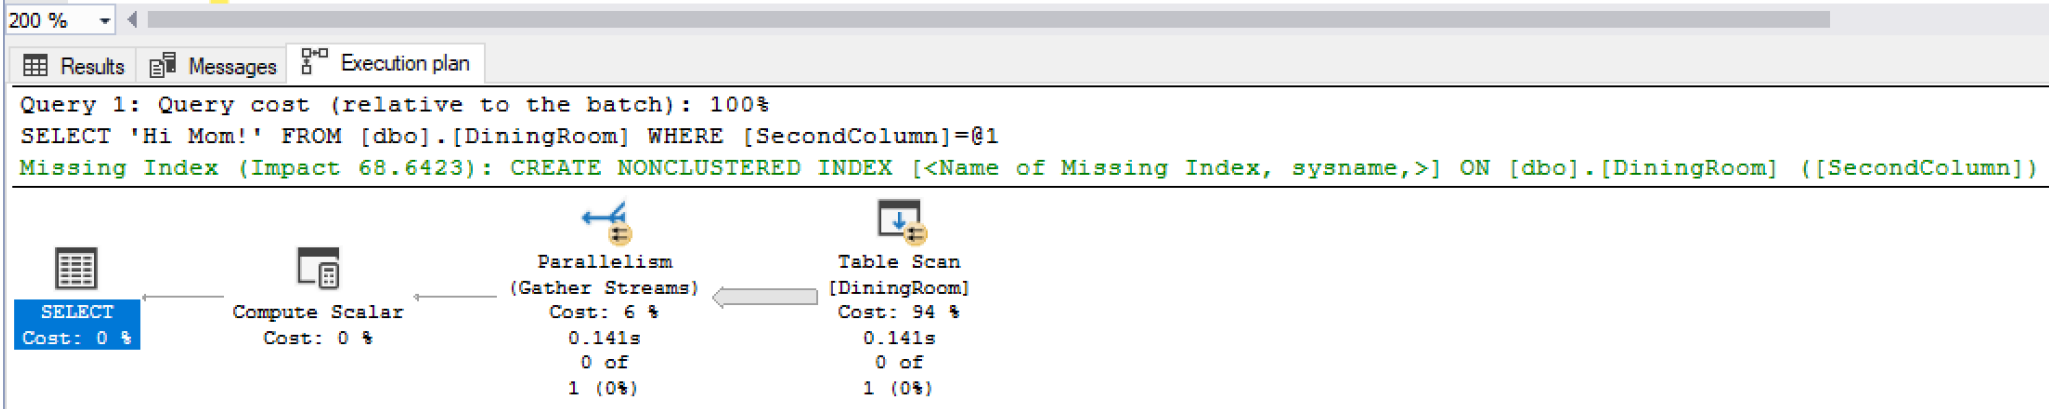
SELECT 'Hi Mom!' FROM [dbo].[DiningRoom] WHERE [FirstColumn]=0

Missing Index (Impact 68.6423): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[DiningRoom] ([FirstColumn])



## Also works if we look for column 2

```
51 SELECT 'Hi Mom!'  
52 FROM dbo.DiningRoom  
53 WHERE SecondColumn = 0;  
54 GO
```



# And if we look for both columns...

```
59 SELECT 'Hi Mom!'
60 FROM dbo.DiningRoom
61 WHERE FirstColumn = 0
62        AND SecondColumn = 0;
63
```


200 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT 'Hi Mom!' FROM [dbo].[DiningRoom] WHERE [FirstColumn]=@1 AND [SecondColumn]=@2

Missing Index (Impact 59.3985): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[DiningRoom] ([FirstColumn],[SecondColumn])





# So far, not bad.



## And if we flip the WHERE clause?

What if we put SecondColumn first?

```
64 SELECT 'Hi Mom!'  
65 FROM dbo.DiningRoom  
66 WHERE SecondColumn = 0  
67     AND FirstColumn = 0;  
68 GO
```



# Hmm...what's determining order?

```
64 SELECT 'Hi Mom!'
65 FROM dbo.DiningRoom
66 WHERE SecondColumn = 0
67 AND FirstColumn = 0;
68 GO
69
```


200 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

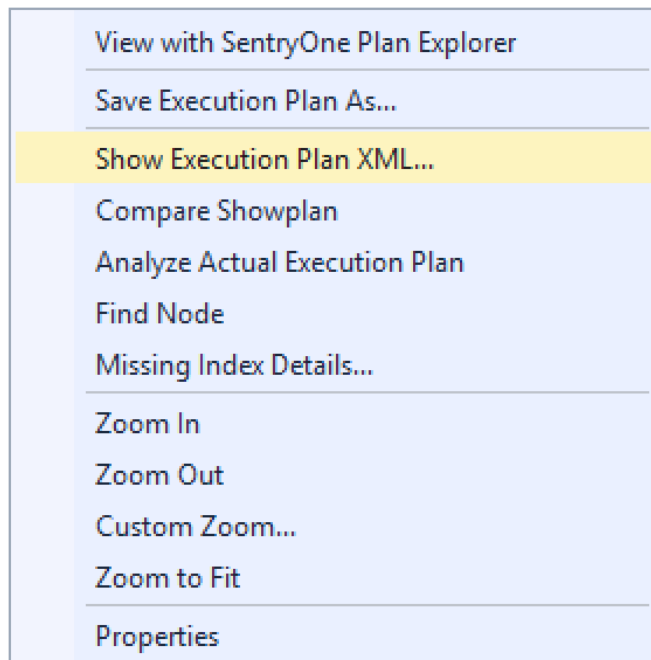
SELECT 'Hi Mom!' FROM [dbo].[DiningRoom] WHERE [SecondColumn]=@1 AND [FirstColumn]=@2

Missing Index (Impact 59.3985): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[DiningRoom] ([FirstColumn],[SecondColumn])



# View the execution plan XML

```
dbo].[DiningRoom] ([FirstColumn],[SecondColumn])
```



To see how order is calculated, right-click on the plan and view the XML:



# Clippy uses the table order.

The first column in the table goes first,  
second goes second, and so forth.

```
</InreadStat>
```

```
<MissingIndexes>
```

```
Group Impact="59.3985">
```

```
Database="[StackOverflow2013]" Schema="[dbo]" Table="[DiningRoom]">
```

```
Usage="EQUALITY">
```

```
Name="[FirstColumn]" ColumnId="1" />
```

```
Name="[SecondColumn]" ColumnId="2" />
```


```
Group>
```

```
Index>
```

```
Group>
```

```
</>
```

```
SerialRequiredMemory="0" SerialDesiredMemory="0" RequiredMemory="72"
```



# It's just a little bit more complex...

Clippy picks key order using:

- Equality searches  
(=, IS NULL, IN a list of 1)  
ordered by the column they are in the table
- Inequality search columns  
(<, >, LIKE, IS NOT NULL, IN a list of 2 or more)  
ordered by the column they are in the table




# Clippy can't consider

How often you filter on a field

How selective your filter clause is

The size of the field

What you do further upstream  
(joining, grouping, ordering)

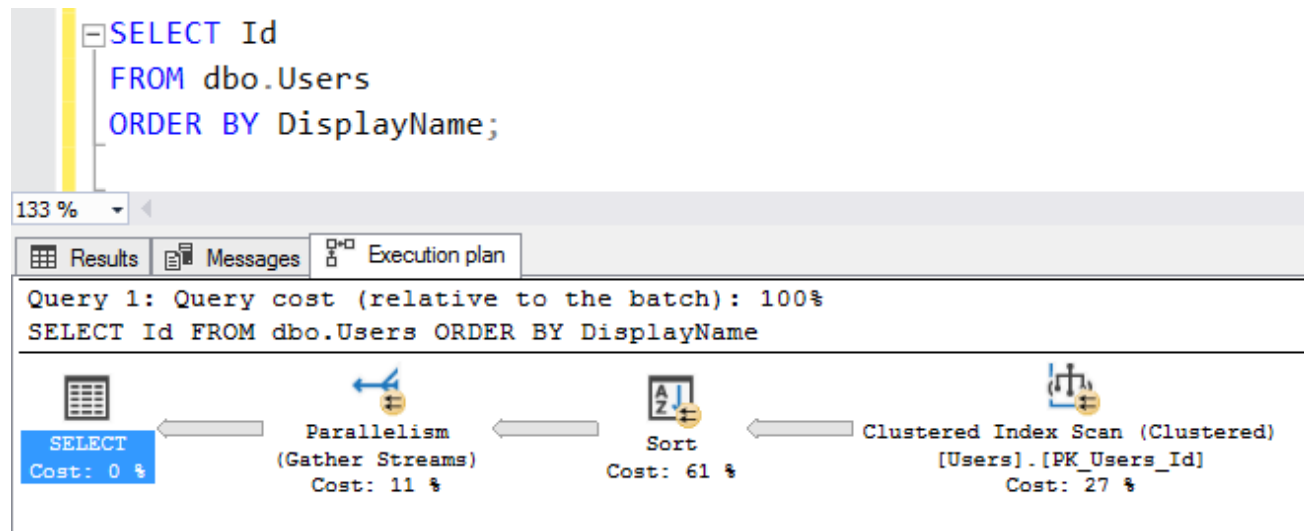




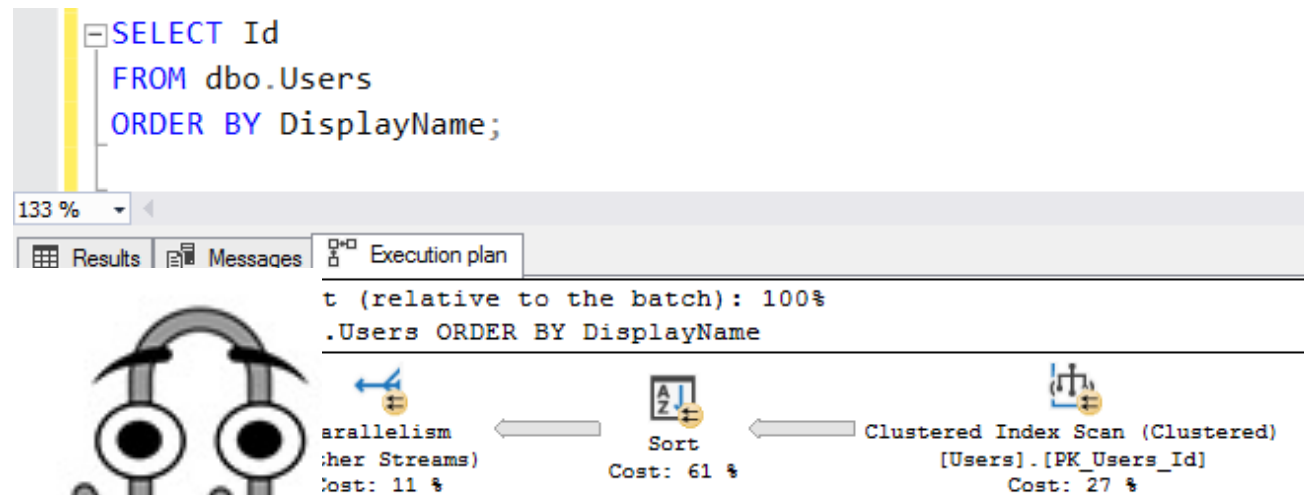
**He's focused on WHERE,  
not GROUP BY or ORDER BY.**



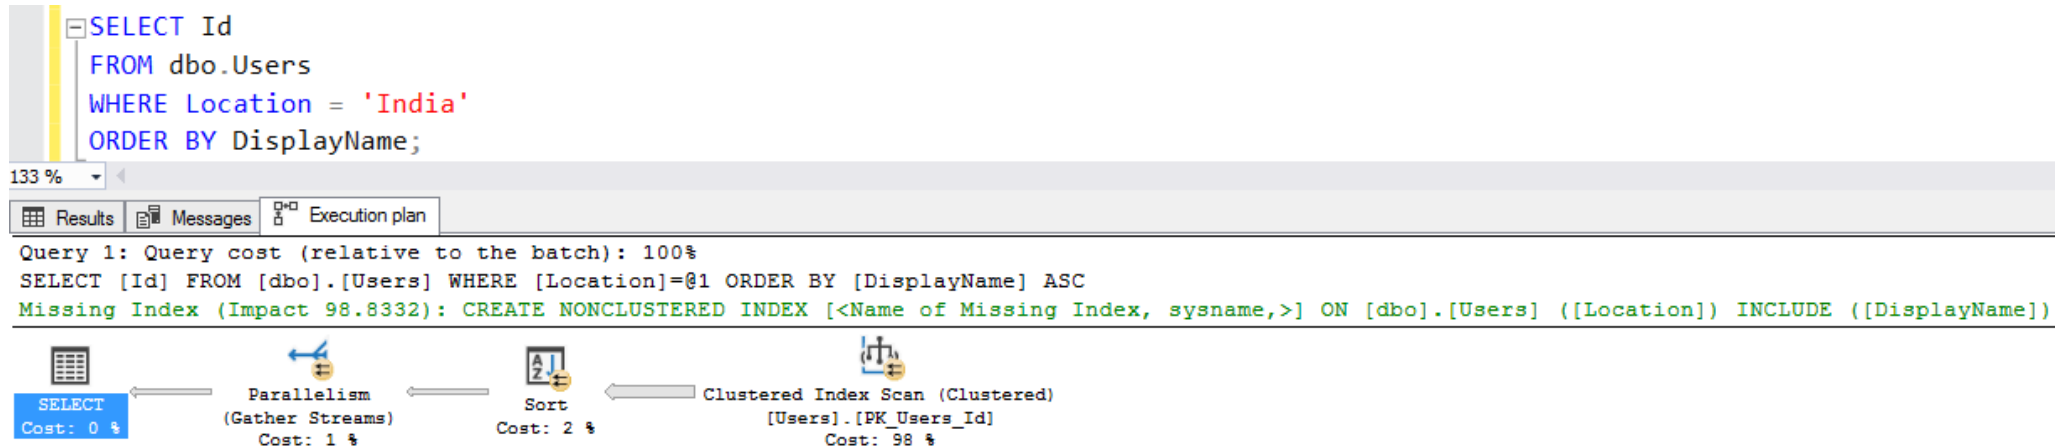
# Order the whole table



# Order the whole table



# Filter, then order by



Clippy just INCLUDEs DisplayName, figuring he's going to sort all of the people in India by name, every single time this query runs. Another blind spot.

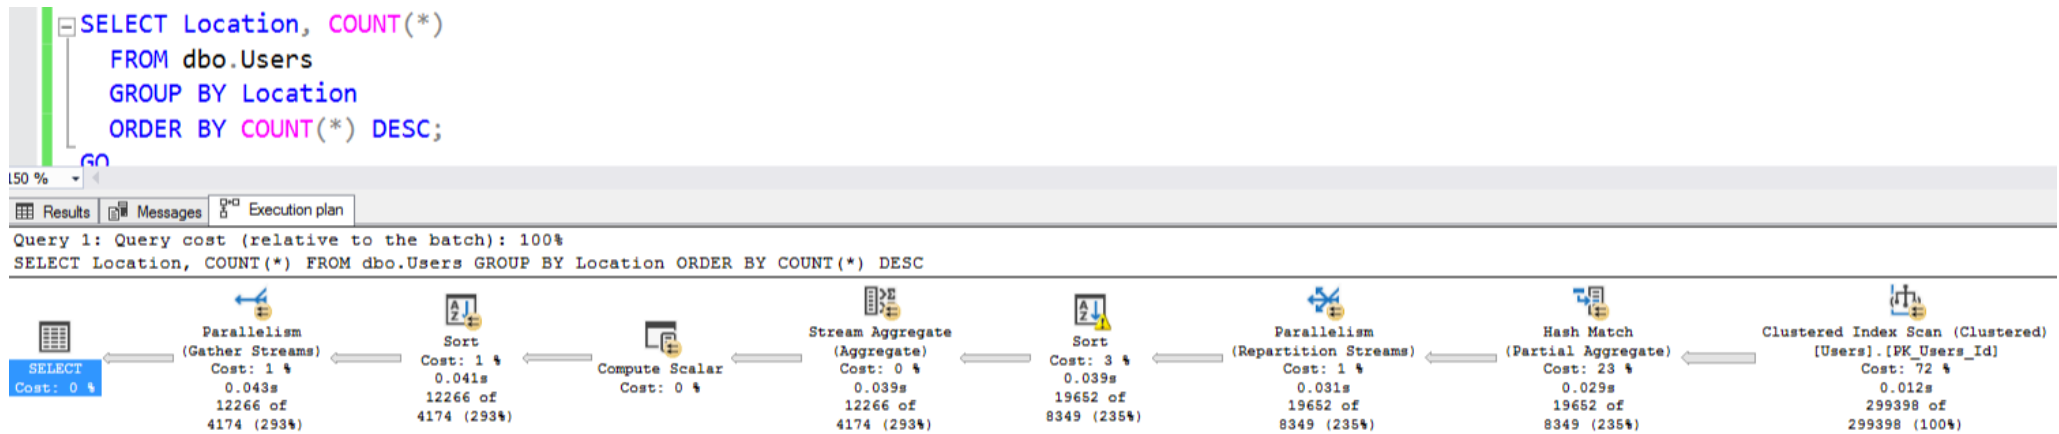


## What's he suggest for this?

```
SELECT Location, COUNT(*)  
FROM dbo.Users  
GROUP BY Location  
ORDER BY COUNT(*) DESC;
```



# Seems like a lot of work



Scan the whole table, dump locations into buckets, go parallel across threads, sort them, spill to disk...

But no index recommendation?



## Try creating one by hand.


```
CREATE INDEX IX_Location  
ON dbo.Users(Location);
```





## Try creating one by hand.

```
CREATE INDEX IX_Location  
ON dbo.Users(Location);
```



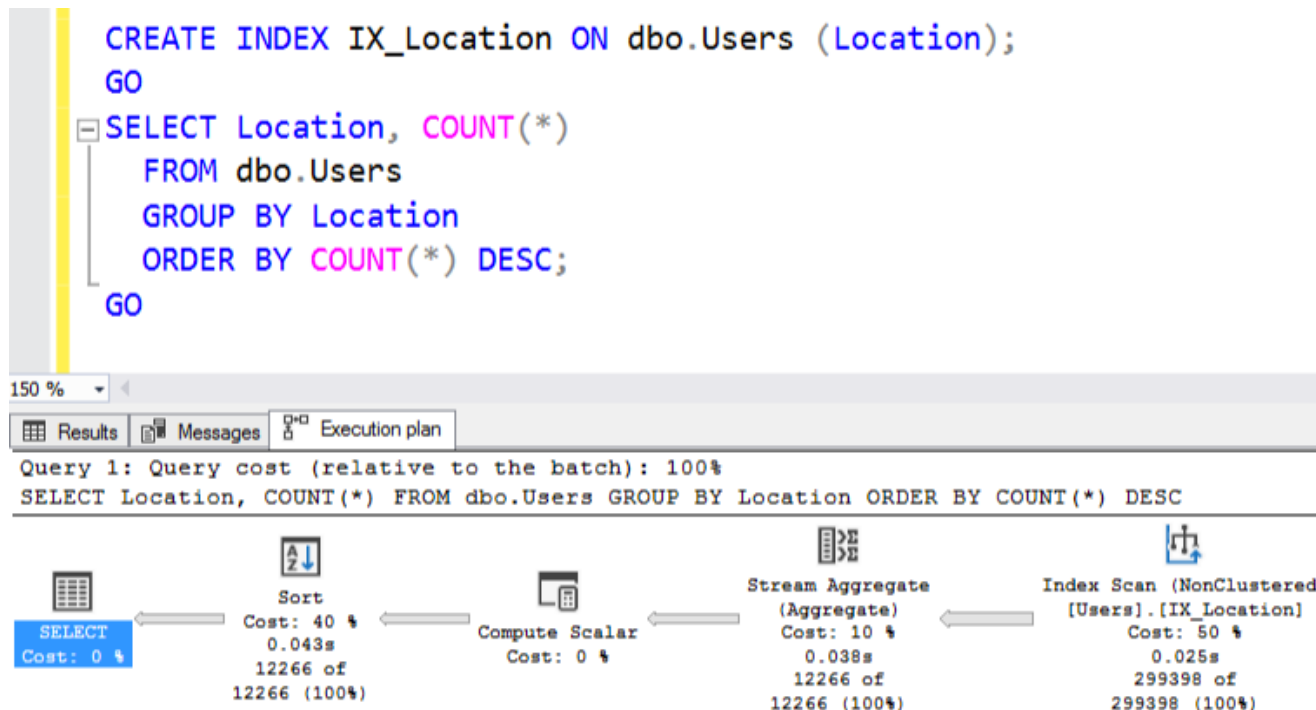
# He uses the index

Way faster


Single-threaded

Great estimates

No spills to disk



**Adding Clippy's indexes can  
even make things worse.**



# Disclaimer: reproing this is tricky.

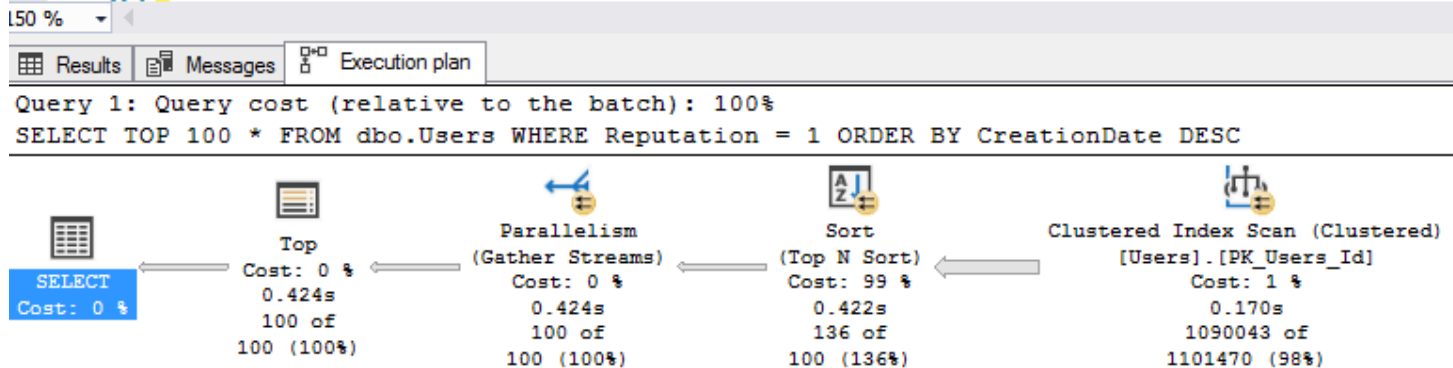
The exact index suggestions will vary based on:

- Your Stack database size  
(10GB, 50GB, 300+GB)
- Your SQL Server version
- Cost Threshold for Parallelism



# Try this query with no indexes.

```
57 DropIndexes;
58 GO
59 SELECT TOP 100 *
60 FROM dbo.Users
61 WHERE Reputation = 1
62 ORDER BY CreationDate DESC;
63 GO
64
```



# Add an index, and it's fast!

```
65 CREATE INDEX IX_CreationDate ON dbo.Users(CreationDate);
66 GO
67 SELECT TOP 100 *
68 FROM dbo.Users
69 WHERE Reputation = 1
70 ORDER BY CreationDate DESC;
71 GO
```


150 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT TOP 100 \* FROM dbo.Users WHERE Reputation = 1 ORDER BY CreationDate DESC

Missing Index (Impact 61.1101): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[Users] ([Reputation]) INCLUDE ([AboutMe],[Age],[CreationDate])



# Add an index, and it's fast, but...

```
65 CREATE INDEX IX_CreationDate ON dbo.Users(CreationDate);
66 GO
67 SELECT TOP 100 *
68 FROM dbo.Users
69 WHERE Reputation = 1
70 ORDER BY CreationDate DESC;
71 GO
```

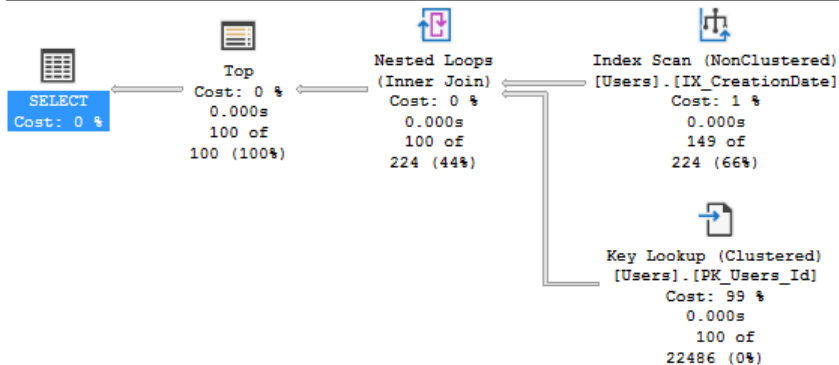
150 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT TOP 100 \* FROM dbo.Users WHERE Reputation = 1 ORDER BY CreationDate DESC

Missing Index (Impact 61.1101): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[Users] ([Reputation]) INCLUDE ([AboutMe],[Age],[CreationDate])





# Now he's got an idea.

```
65 CREATE INDEX IX_CreationDate ON dbo.Users(CreationDate);
66 GO
67 SELECT TOP 100 *
68 FROM dbo.Users
69 WHERE Reputation = 1
70 ORDER BY CreationDate DESC;
71 GO
```

150 %

```
ch): 100%
putation = 1 ORDER BY CreationDate DESC
>NCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[Users] ([Reputation]) INCLUDE ([AboutMe],[Age],[CreationDate])
```



```
Index Scan (NonClustered)
[Users].[IX_CreationDate]
Cost: 1 %
0.000s
149 of
224 (66%)
```



```
Key Lookup (Clustered)
[Users].[PK_Users_Id]
Cost: 99 %
0.000s
100 of
22486 (0%)
```



```
65 CREATE INDEX IX_CreationDate ON dbo.Users(CreationDate);
66 GO
67 SELECT TOP 100 *
68 FROM dbo.Users
69 WHERE Reputation = 1
70 ORDER BY CreationDate DESC;
71 GO
```

That's...interesting.

150 %

Results Messages Execution plan

SQL Server parse and compile time:  
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server Execution Times:  
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server parse and compile time:  
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server parse and compile time:  
CPU time = 0 ms, elapsed time = 1 ms.

(100 rows affected)

Table 'Users'. Scan count 1, logical reads 468, physical reads 0, read

(1 row affected)

SQL Server Execution Times:  
CPU time = 0 ms, elapsed time = 136 ms.

The query is already  
really, really fast, and  
does pretty few logical  
reads.

# He wants to double the table size.

```
3 The Query Processor estimates that implementing the following index could improve the query cost by 61.1101%.
4 */
5
6 /*
7 USE [StackOverflow2013]
8 GO
9 CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>]
10 ON [dbo].[Users] ([Reputation])
11 INCLUDE ([AboutMe],[Age],[CreationDate],[DisplayName],[DownVotes],[EmailHash],[LastAccessDate],[Location],[UpVotes],[Views],[WebsiteUrl],[AccountId])
```



But note the index's key.



```

75 CREATE NONCLUSTERED INDEX IX_Clippy_Reputation
76 ON [dbo].[Users] ([Reputation])
77 INCLUDE ([AboutMe],[Age],[CreationDate],[DisplayName],[DownVotes],[Email
78
79 SELECT TOP 100 *
80 FROM dbo.Users
81 WHERE Reputation = 1
82 ORDER BY CreationDate DESC;
83 GO

```

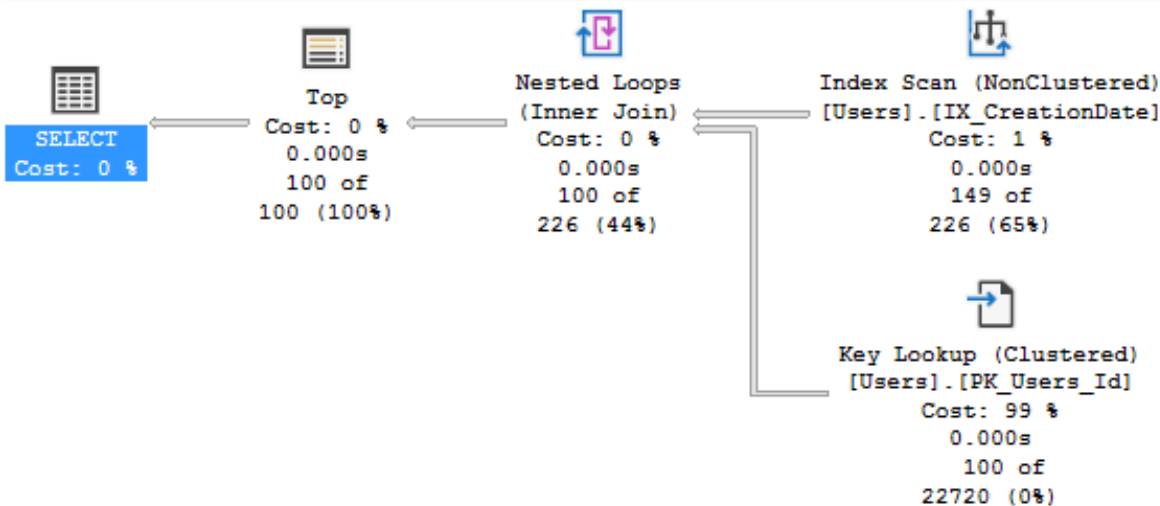
**We create it.  
It doesn't get used!**

150 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT TOP 100 \* FROM dbo.Users WHERE Reputation = 1 ORDER BY CreationDate DESC



```

75 CREATE NONCLUSTERED INDEX IX_Clippy_Reputation
76 ON [dbo].[Users] ([Reputation])
77 INCLUDE ([AboutMe],[Age],[CreationDate],[DisplayName],[DownVotes],[Email]
78
79 SELECT TOP 100 *
80 FROM dbo.Users
81 WHERE Reputation = 1
82 ORDER BY CreationDate DESC;
83 GO

```

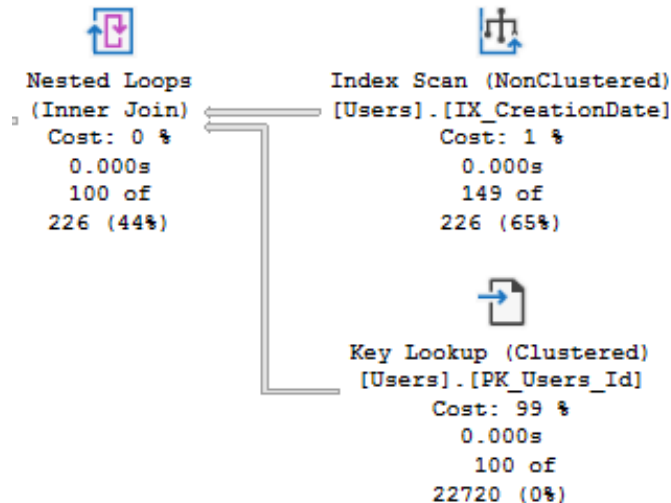

**We create it.  
It doesn't get used!**

150 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

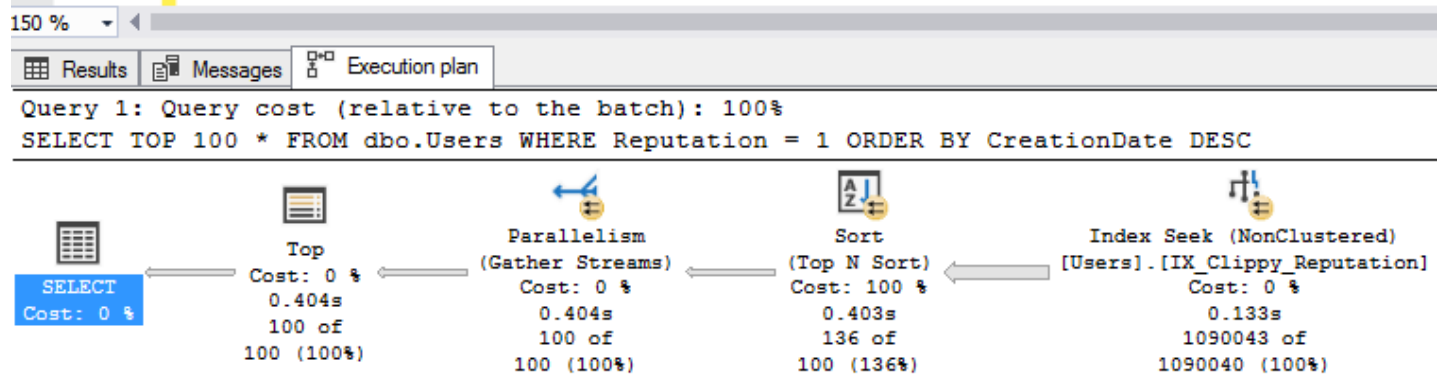
SELECT TOP 100 \* FROM dbo.Users WHERE Reputation = 1 ORDER BY CreationDate DESC



# Drop the old IX\_CreationDate...

```
86 DROP INDEX dbo.Users.IX_CreationDate;  
87 GO  
88 SELECT TOP 100 *  
89 FROM dbo.Users  
90 WHERE Reputation = 1  
91 ORDER BY CreationDate DESC;  
92 GO
```

And the index  
gets used,  
but...that sort!



# Now we're sorting 1M rows.

CPU time,  
elapsed time,  
and logical reads  
are all WORSE  
than the original query.

```
86 DROP INDEX dbo.Users.IX_CreationDate;  
87 GO  
88 SELECT TOP 100 *  
89 FROM dbo.Users  
90 WHERE Reputation = 1  
91 ORDER BY CreationDate DESC;  
92 GO
```

150 %

Results Messages Execution plan

SQL Server parse and compile time:  
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server Execution Times:  
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server parse and compile time:  
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server parse and compile time:  
CPU time = 0 ms, elapsed time = 1 ms.

(100 rows affected)


Table 'Users'. Scan count 5, logical reads 14119, physical reads 0, logical writes 0, physical writes 0.

Table 'Worktable'. Scan count 0, logical reads 0, physical reads 0, logical writes 0, physical writes 0.

(1 row affected)

SQL Server Execution Times:  
CPU time = 1577 ms, elapsed time = 520 ms.





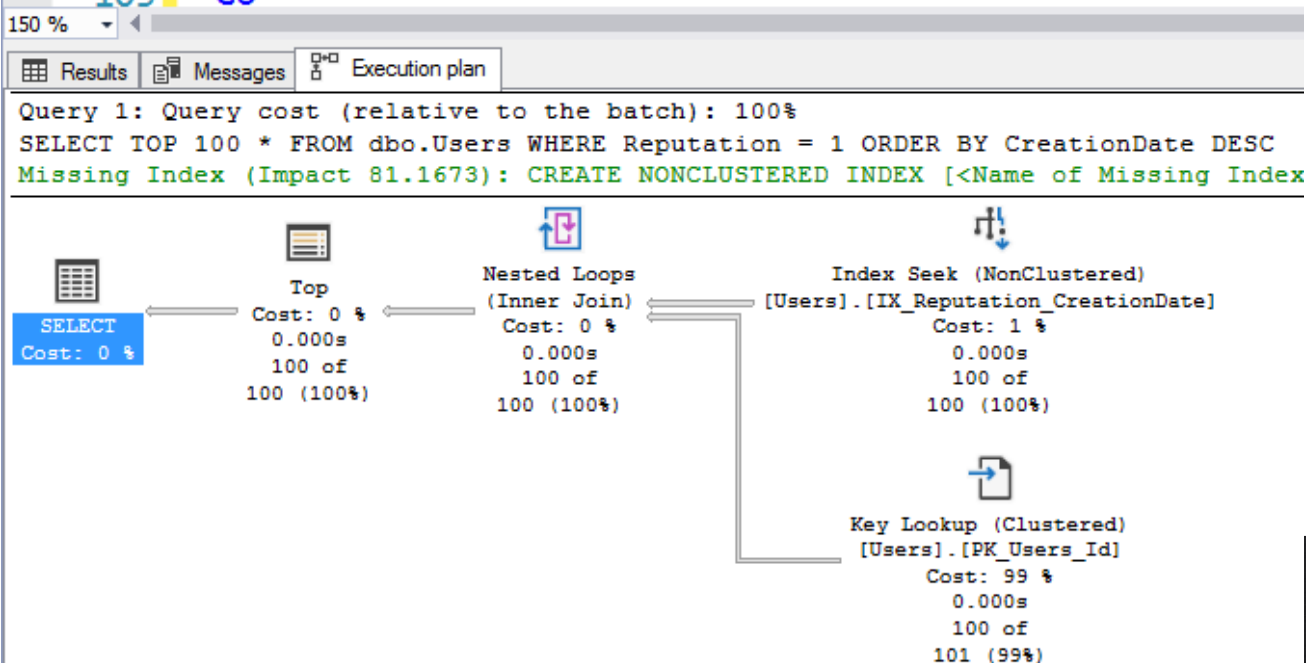


# Clippy was on to something

```
94 DropIndexes;  
95 GO  
96 CREATE INDEX IX_Reputation_CreationDate  
97 ON dbo.Users(Reputation, CreationDate);  
98 GO  
99 SELECT TOP 100 *  
100 FROM dbo.Users  
101 WHERE Reputation = 1  
102 ORDER BY CreationDate DESC;  
103 GO
```

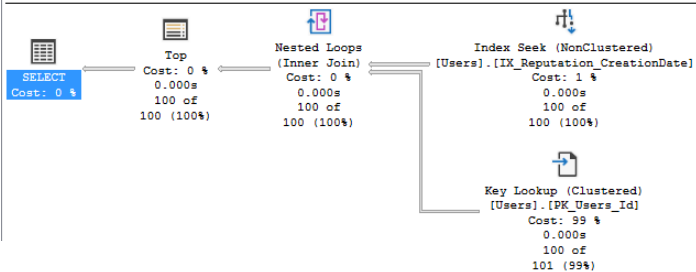
An index tweak will help – just not the index Clippy wanted.

Key on both fields, and the sort is gone.



# And he STILL wants the index.

Query 1: Query cost (relative to the batch): 100%  
SELECT TOP 100 \* FROM dbo.Users WHERE Reputation = 1 ORDER BY CreationDate DESC  
Missing Index (Impact 81.1673): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[Users] ([Reputation]) INCLUDE ([AboutMe],[Age],[CreationDate],[DisplayName],[DownVotes],[EmailHash],[LastAccessDate],[...



## What we saw

A query wasn't terribly slow,  
but SQL Server asked for an index

If this was a frequent query,  
that index might seem attractive

But the requested index had the ORDER BY column  
as an include, when it really needs to be sorted

The query was much better with that column in the key



# How to identify it

Look for high average CPU and reads on top plans

Dig into every operator

In the real world on big plans, this is time consuming

You have to rule out other things that may be the issue, such as parameter sniffing and inefficient or out of date statistics



**That's where  
tools come in.**



# sp\_BlitzIndex

Github repository: [FirstResponderKit.org](https://github.com/FirstResponderKit.org)

Psychiatrist-style analysis of indexes

But be warned: all its data comes from Clippy

- Index usage stats reset at odd times
- Missing index recommendations are derp
- Only really works in production



## Running it at the server level

```
sp_BlitzIndex @GetAllDatabases = 1;
```

I don't tune here, but I use this to get a fast overall picture of which databases & tables to focus on.



## At the table level

```
sp_BlitzIndex @SchemaName = 'dbo',  
  @TableName = 'Users';
```

This is where I spend most of my time tuning.





cap

cap



# Recap

You don't always get missing index requests.

Even when you do, Clippy's not putting much work in:

- Equality searches first, then inequality searches
- Fields ordered by their position in the table
- He's completely focused on the WHERE



Tools like `sp_BlitzIndex` get their hints from Clippy.

You can easily do better by hand.

