



BRENT OZAR
UNLIMITED®

Why 2 Out of 3 Functions Don't Scale

Let's put the fun in functions!

2.2 p1

Agenda

Bad functions:

- Scalar user-defined functions
- Multi-statement table-valued functions
- Functions in table definitions

Not-so-bad functions:

- SQL Server 2019's Froid technology
- Inline table-valued functions
- CTE and OUTER APPLY

How to find & fix the bad ones





It's not your fault.
Good developers are taught to use functions.

Developers are taught to reuse code

Functions are helpful to:

- Encapsulate code and package it for easier reuse
- Make it more test-friendly
- Don't repeat yourself (DRY)

All that stuff is true in app languages,
and functions don't get a performance overhead there either.



Your functions are no good here

- The most important takeaway is that SQL is set-based, and calculating values row by row is painful
- Same applies to cursors, they're almost always performance killers, but that's another issue.
- WHILE loops have similar performance problems, and often get used in functions for splitting strings, padding strings, removing characters, putting names in proper case, etc.



Built-in system functions can be bad, too

- Usually (but not always) have bad row estimations
- Which leads to query plans that scan rather than seek
- Can sometimes lead to dramatically higher CPU use

Examples:

- `WHERE UPPER(DisplayName) = 'BRENT'`
- `WHERE LTRIM(RTRIM(DisplayName)) = 'BRENT'`
- `WHERE ISNULL(DisplayName, '') = 'BRENT'`

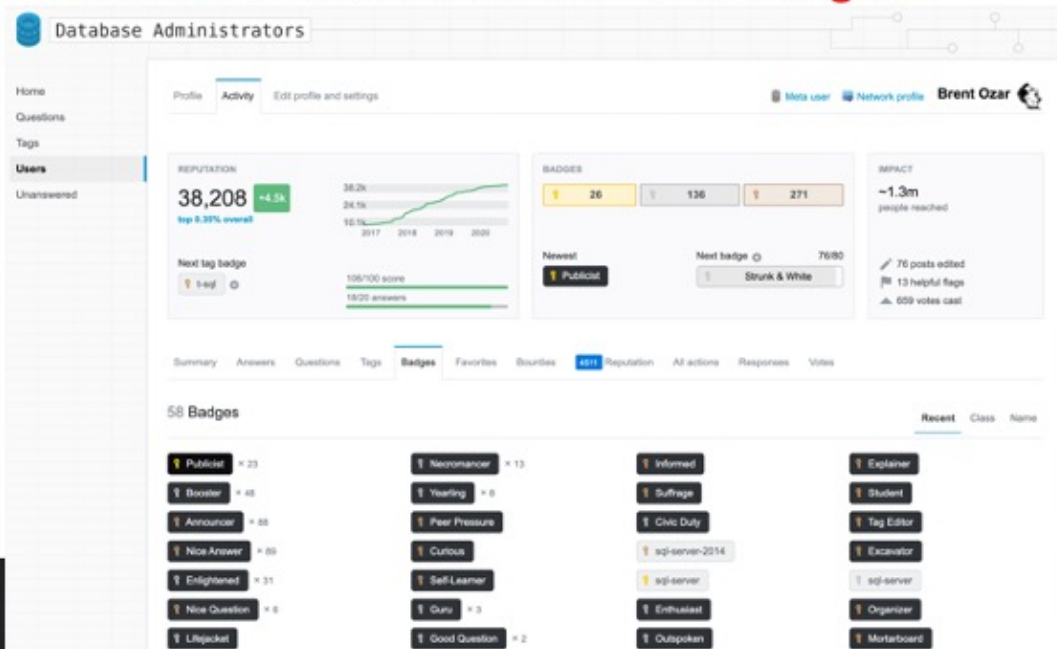
But in this class, I'm specifically talking about user-defined functions: ones we created.





Scalar user-defined functions

StackOverflow users can earn badges.



2.2 p8

We want to count badges for each user.

Code goals:

- When we show a user's info, we often want to show their # of badges
- Ideally, we code this just once, and it's reusable



The scalar function version

```
-- Scalar Function
CREATE FUNCTION dbo.ScalarFunction ( @uid INT )
RETURNS BIGINT
    WITH RETURNS NULL ON NULL INPUT,
    SCHEMABINDING
AS
BEGIN
    DECLARE @BCount BIGINT;
    SELECT @BCount = COUNT_BIG(*)
    FROM     dbo.Badges AS b
    WHERE    b.UserId = @uid
    GROUP BY b.UserId;
    RETURN @BCount;
END;
```



How you call it

```
SELECT TOP 1000  
u.DisplayName,  
dbo.ScalarFunction(u.Id)  
FROM dbo.Users AS u
```



It takes ~8 minutes to return 1,000 rows.

```
SELECT TOP 1000
  u.DisplayName,
  dbo.ScalarFunction(u.Id)
FROM dbo.Users AS u;
GO
```

150 %

Results Messages Execution plan

	DisplayName	(No column name)
1	Community	1
2	Jeff Atwood	328
3	Geoff Dalgas	90
4	Jarrod Dixon	130
5	Joel Spolsky	185
6	Jon Galloway	319
7	Eggs McLaren	24
8	Kevin Dente	87
9	Sneakers O'Toole	3

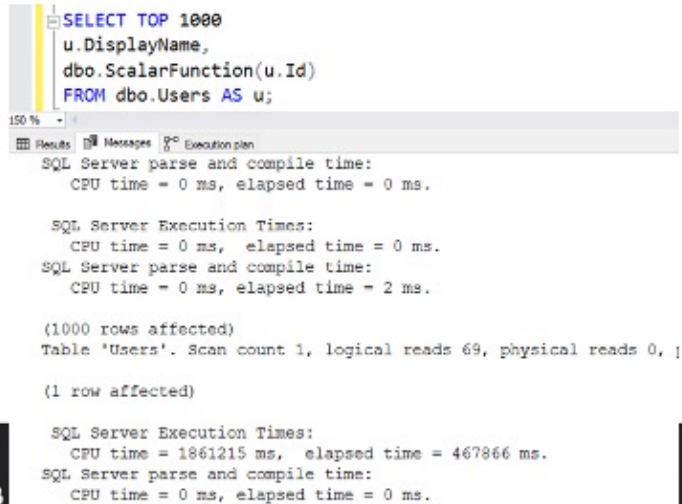
Query executed successfully. SQL2019 (15.0 RTM) SQL2019\Administrator ... StackOverflow 00:07:48 1,000 rows

2.2 p12



When a query is this slow, you wanna see what tables it read from.

But SET STATISTICS IO ON completely lies:
it doesn't show anything that scalar functions do.



```
SELECT TOP 1000
  u.DisplayName,
  dbo.ScalarFunction(u.Id)
FROM dbo.Users AS u;
```

SQL Server parse and compile time:
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server Execution Times:
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server parse and compile time:
CPU time = 0 ms, elapsed time = 2 ms.

(1000 rows affected)
Table 'Users'. Scan count 1, logical reads 69, physical reads 0, ...

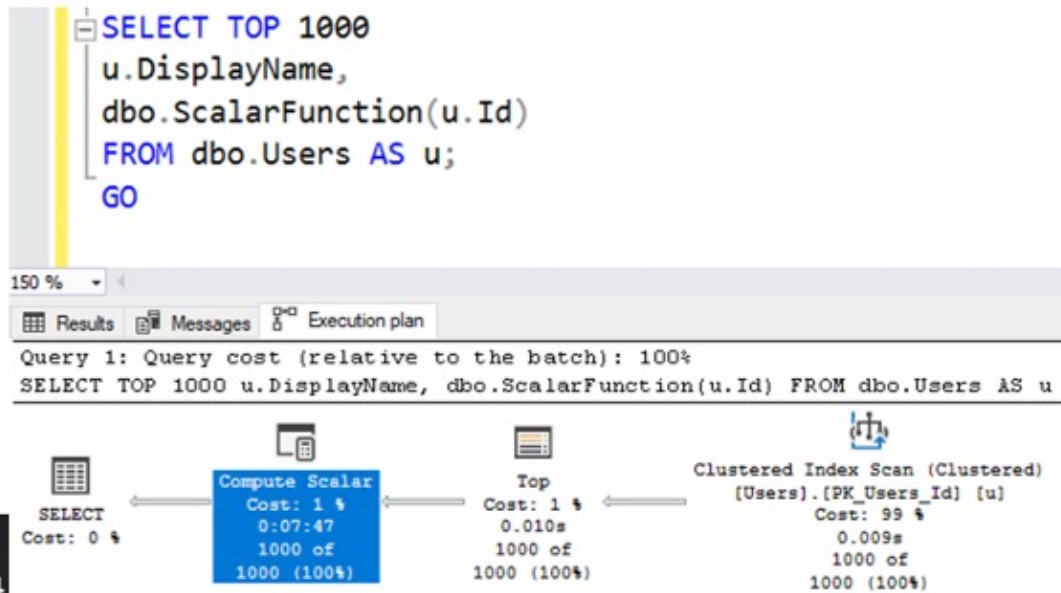
(1 row affected)

SQL Server Execution Times:
CPU time = 1861215 ms, elapsed time = 467866 ms.


SQL Server parse and compile time:
CPU time = 0 ms, elapsed time = 0 ms.

2.2 p13

What about the query plan? It gets worse.



What about the query plan? It gets worse.



SQL Server implies the clustered index scan is 99% of the cost, and that the Compute Scalar was just 1%.

That's because the estimated costs for scalars is completely made up.



```

SELECT TOP 1000
  u.DisplayName,
  dbo.ScalarFunction(u.Id)
FROM dbo.Users AS u;
GO

```

Query 1: Query cost (relative to the batch): 100%

SELECT TOP 1000 u.DisplayName, dbo.ScalarFunction(u.Id) FROM dbo.Users AS u

Cost: 1 %

0:07:47

1000 of 1000 (100%)

Compute Scalar

Compute new values from existing values in a row.

Physical Operation	Compute Scalar
Logical Operation	Compute Scalar
Actual Execution Mode	Row
Estimated Execution Mode	Row
Actual Number of Rows	1000
Actual Number of Batches	0
Estimated I/O Cost	0
Estimated Operator Cost	0.0001 (1%)
Estimated Subtree Cost	0.0162865
Estimated CPU Cost	0.0001
Number of Executions	1
Estimated Number of Executions	1
Estimated Number of Rows	1000
Estimated Row Size	59 B
Actual Rebinds	0
Actual Rewinds	0
Node ID	0

Query executed successfully.

Output List

[Stack Overflow].[dbo].[Users].[DisplayName], Expr1001

Nick was late for beer


The Estimated I/O Cost = 0

Estimated CPU Cost = 0.0001

These are purely made-up numbers that have nothing to do with the contents of the function.

Microsoft just didn't get around to coding this part of the engine.

What about the query plan? It gets worse.




SQL Server implies the clustered index scan is 99% of the cost, and that the Compute Scalar was just 1%.

That's because the estimated costs for scalars is completely made up.



You have to do some detective work.

If you look closely at the completion times for each operator:



The table scan finished 0.009 seconds into the query.

The Compute Scalar finished 7 minutes, 47 seconds in.



Another way to spot it

Look at the properties of the SELECT operator on newer versions:

The screenshot shows a SQL query in the query editor:

```
SELECT TOP 1000  
u.DisplayName,  
dbo.ScalarFunction(u.Id)  
FROM dbo.Users AS u;
```

Below the query is the execution plan for Query 1. The plan shows three operators: a SELECT operator (Cost: 0%), a Compute Scalar operator (Cost: 1%, 0:07:47, 1000 of 1000 (100%)), and a Top operator (Cost: 1%, 0:01:00, 1000 of 1000 (100%)). The Top operator is connected to a Clustered Index Scan (Clustered) operator (Cost: 99%, 0:00:59, 1000 of 1000 (100%)). A red arrow points from the Top operator to the properties window.

The properties window on the right shows the following properties for the SELECT operator:

Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	0.0162865
MemoryGrantInfo	
NonParallelPlanReason	CouldNotGenerateValidParallelPlan
Optimization Level	TRIVIAL
OptimizerHardwareDependentProperty	
ParentObjectId	0
QueryHash	0xD07CEC48EEF3112C
QueryPlanHash	0xEAF1090155A0217E
QueryTimeStats	
CpuTime	1862424
ElapsedTime	467820
UdfCpuTime	1862407
UdfElapsedTime	467802
RetrievedFromCache	true
SecurityPolicyApplied	False
SetOptions	ANSI_NULLS: True, ANSI_PADDING: True

UdfCpuTime, UdfElapsedTime show how much time was spent inside user-defined functions.



But my favorite way is sp_BlitzCache.

sp_BlitzCache

Database	Cost	Query Text	Query Type	Warnings	# Executions	Total CPU (ms)	Avg CPU (ms)	Total Duration (ms)	Avg Duration (ms)
StackOverflow	0.0162865	SELECT TOP 1000 u.DisplayName, dbo.ScalarFu...	Statement	Long Running Query, Downlevel CE, Forced Serialization, Plan created last 4hrs, Low Cost High CPU	1	1862309	1862309	467896	467896
StackOverflow	135.252	CREATE FUNCTION dbo.ScalarFunction (@But I...	Procedure or Function [dbo].[ScalarFunction]	Missing Indexes (1): Parallel, Parameter Sniffing, Downlevel CE, Plan created last 4hrs	1000	1862274	1862	467787	467
StackOverflow	135.252	SELECT @BCount = COUNT_BIG(*) FROM ...	Statement (parent [dbo].[ScalarFunction])	Missing Indexes (1): Parallel, Parameter Sniffing, Downlevel CE, Plan created last 4hrs	1000	1862239	1862	467752	467

#1 query is the SELECT TOP 1000, which calls the function.

The SELECT's CPU & duration *includes the function's work*.

Now look at the #2 query: it's the function.

The function's CPU time and duration is almost the entire time that the SELECT query worked. That's your big sign.



Scalar round up

The bad:

- Runs once per row
- Cost isn't added to the total query cost
- No info from STATISTICS IO, not much in the plan
- Inhibits parallelism, too: BrentOzar.com/go/serialudf

The good:

This section is intentionally left blank





Multi-statement table-valued user-defined functions (MSTVFs)

MSTVFs use table variables to store data.

Which means all modifications are serial, and

Your estimates will be really low (unless you recompile), and

You won't get column statistics (so SQL won't know how joins will go)



How you build it

```
-- Multi Statement Table Valued Function
CREATE FUNCTION dbo.MultiStatementTVF ( @uid INT )
RETURNS @Out TABLE ( BadgeCount BIGINT )
WITH SCHEMABINDING
AS
BEGIN
    INSERT INTO @Out (BadgeCount)
    SELECT COUNT_BIG(*) AS BadgeCount
    FROM    dbo.Badges AS b
    WHERE   b.UserId = @uid
    GROUP BY b.UserId;
    RETURN;
END;
```

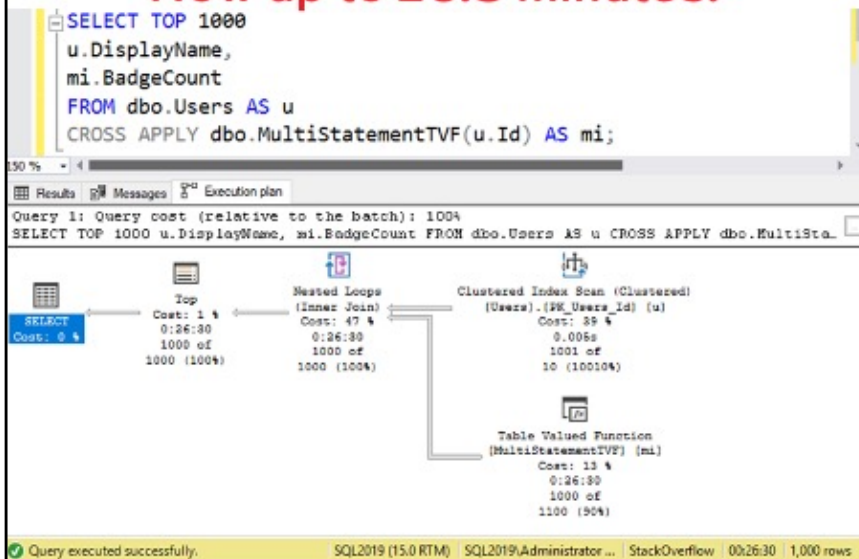


How you call it

```
SELECT TOP 1000  
u.DisplayName,  
mi.BadgeCount  
FROM dbo.Users AS u  
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi
```



Now up to 26.5 minutes.



2.2 p26



Now up to 26.5 minutes.

```
SELECT TOP 1000
u.DisplayName,
mi.BadgeCount
FROM dbo.Users AS u
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;
```

Query 1: Query cost (rel. to 1000000)
SELECT TOP 1000 u.Display

	Top
Cost: 1 %	
0:26:30	
1000 of	
1000 (100%)	



Query executed successfully. SQL2019 (15.0 RTM) SQL2019\Administrator ... StackOverflow 00:26:30 1,000 rows

2.2 p27



What's STATISTICS IO say this time?

```
SELECT TOP 1000
  u.DisplayName,
  mi.BadgeCount
FROM dbo.Users AS u
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;
```

150 %

Results Messages Execution plan

SQL Server parse and compile time:
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server Execution Times:
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server parse and compile time:
CPU time = 3 ms, elapsed time = 3 ms.

(1000 rows affected)

Table '#B1482DC7'. Scan count 1001, logical reads 1001, physical reads 0,
Table 'Users'. Scan count 1, logical reads 69, physical reads 0, page ser

(1 row affected)

150 %

2.2 p28

Query executed successfully. SQL2019 (15.0 RTM) SQL2019\Administrator ... StackOverflow 00:26:30 1,000 rows

What's STATISTICS IO say this time?

```
SELECT TOP 1000  
u.DisplayName,  
mi.BadgeCount  
FROM dbo.Users AS u  
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;
```



Table '#B1482DC7'. Scan count 1001, logical reads 1001, physical reads 0,
Table 'Users'. Scan count 1, logical reads 69, physical reads 0, page ser
(1 row affected)

2.2 p29

Query executed successfully.

SQL2019 (15.0 RTM)

SQL2019\Administrator ...

StackOverflow

00:26:30

1,000 rows

UdfCpuTime or UdfElapsedTime? No, none.

SELECT TOP 1000
u.DisplayName,
mi.BadgeCount
FROM dbo.Users AS u
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;

Query 1: Query cost (relative to the batch): 100%
SELECT TOP 1000 u.DisplayName, mi.BadgeCount FROM dbo.Users AS u CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;

Execution Plan:

- SELECT (Cost: 0%)
 - Top (Cost: 1%, 0:26:30, 1000 of 1000 (100%))
 - Nested Loops (Inner Join) (Cost: 47%, 0:26:30, 1000 of 1000 (100%))
 - Clustered Index Scan (Clustered) (Users) (PK_Users_Id) (u) (Cost: 39%, 0:00:00, 1001 of 10 (10010%))
 - Table Valued Function (MultiStatementTVF) (mi) (Cost: 13%, 0:26:30, 1000 of 1100 (90%))

QueryTimeStats:

DatabaseContextSettingsId	2
Degree of Parallelism	1
Estimated Number of Rows	1000
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	0.0086923
MemoryGrantInfo	
Optimization Level	FULL
OptimizerHardwareDependentProperties	
ParentObjectId	0
QueryHash	0x0C9CDAC7E65B4975
QueryPlanHash	0xC258F57D7A88298E
QueryTimeStats	
CpuTime	1590539
ElapsedTime	1590577
Reason For Early Termination Of State	Good Enough Plan Found
RetrievedFromCache	true
SecurityPolicyApplied	False
Set Options	ANSI_NULLS: True, ANSI_PADDING: True, ANSI_WARNINGS: True, ARITHABORT: True, CONCAT_NULL_YIELDS_NULL: True, FIPS_FLAGGER: None, NUMERIC_ROUNDABORT: False, QUOTED_IDENTIFIER: True, RECURSIVE: False, ROWCOUNT_LIMIT: 0, SET_OPTIONS: 0, ANSI_NULL_DFLT_ON: True, ANSI_NULL_DFLT_OFF: False, ANSI_WARNINGS: True, ARITHABORT: True, CONCAT_NULL_YIELDS_NULL: True, FIPS_FLAGGER: None, NUMERIC_ROUNDABORT: False, QUOTED_IDENTIFIER: True, SET_OPTIONS: 0, ANSI_NULL_DFLT_ON: True, ANSI_NULL_DFLT_OFF: False, ANSI_WARNINGS: True, ARITHABORT: True, CONCAT_NULL_YIELDS_NULL: True, FIPS_FLAGGER: None, NUMERIC_ROUNDABORT: False, QUOTED_IDENTIFIER: True, SET_OPTIONS: 0
Statement	SELECT TOP 1000 u.DisplayName, mi.BadgeCount FROM dbo.Users AS u CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;
StatementParameterizationType	0
StatementSqlHandle	0x0900B2F33A3CCFD020466668D65D99ED70
WaitStats	

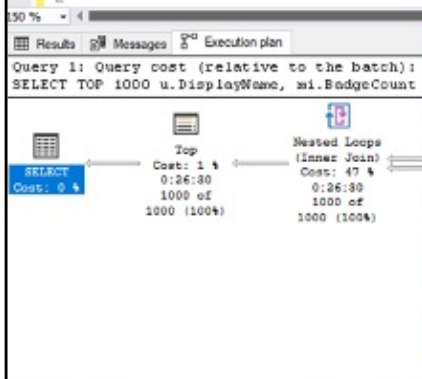
Query executed successfully. SQL2019 (15.0 RTM) SQL2019\Administrator ... StackOverflow 00:26:30 1,000 rows

2.2 p30



UdfCpuTime or UdfElapsedTime? No, none.

```
SELECT TOP 1000
  u.DisplayName,
  mi.BadgeCount
FROM dbo.Users AS u
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;
```



DatabaseContextSettingsId	2
Degree of Parallelism	1
Estimated Number of Rows	1000
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	0.0086923
MemoryGrantInfo	
Optimization Level	FULL
OptimizerHardwareDependentProperties	
ParentObjectId	0
QueryHash	0x0C9CDAC7E65B4975
QueryPlanHash	0xC258F57D7A88298E
QueryTimeStats	
CpuTime	1590539
ElapsedTime	1590577
Reason For Early Termination Of State	Good Enough Plan Found
RetrievedFromCache	true
SecurityPolicyApplied	False
Set Options	
ANSI_NULLS	True
ANSI_PADDING	True
ANSI_WARNINGS	True
ARITHABORT	True
CONCAT_NULL_YIELDS_NULL	True
FULLTEXT_LANGUAGE	1033
NUMERIC_ROUNDABORT	False
QUOTED_IDENTIFIER	True
Statement	SELECT TOP 1000 u.DisplayName, mi.BadgeCount
StatementParameterizationType	0
StatementSqlHandle	0x0900B2F33A3CCFD020466668D65D99ED71
WaitStats	

QueryTimeStats

Query executed successfully. SQL2019 (15.0 RTM) SQL2019\Administrator ... StackOverflow 00:26:30 1,000 rows

2.2 p31



But the best part...

```

SELECT TOP 1000
  u.DisplayName,
  mi.BadgeCount
FROM dbo.Users AS u
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;

```

Query 1: Query cost (relative to the batch): 100%

SELECT TOP 1000 u.DisplayName, mi.BadgeCount FROM dbo.Users AS u CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;

The execution plan shows a Nested Loops join (Inner Join) with a cost of 47%. It consists of a Top operator (cost 1%, 1000 of 1000) and a Clustered Index Scan (cost 39%, 1001 of 1001). The Table Valued Function (MultiStatementTVF) is a child of the Nested Loops operator with a cost of 13% and 1100 of 1100 rows.

DatabaseContextSettingsId	2
Degree of Parallelism	1
Estimated Number of Rows	1000
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	0.0086923
MemoryGrantInfo	
Optimization Level	FULL
OptimizerHardwareDependentProperties	
ParentObjectId	0
QueryHash	0x0C9CDAC7E65B4975
QueryPlanHash	0xC258F57D7A88298E
QueryTimeStats	
CpuTime	1590539
ElapsedTime	1590577
Reason For Early Termination Of Statement Optimization	Good Enough Plan Found
RetrievedFromCache	true
SecurityPolicyApplied	False
Set Options	ANSI_NULLS: True, ANSI_PADDING: False, ANSI_WARNINGS: On, CONCAT_NULL_YIELDS_NULL: On, FULLTEXT_LANGUAGE: English, QUOTED_IDENTIFIER: On, ARITHABORT: On, NUMERIC_ROUNDABORT: On, NOCOUNT: Off, NOEXECUTE: Off, OPTIMIZE_FOR_SEMIstructured_STORAGE: Off, OPTIMIZE_FOR_ADHOC_WORKLOAD: Off, OPTIMIZE_FOR_PARAMETERIZED_QUERIES: Off, OPTIMIZE_FOR_SEMIstructured_STORAGE: Off, OPTIMIZE_FOR_ADHOC_WORKLOAD: Off, OPTIMIZE_FOR_PARAMETERIZED_QUERIES: Off
Statement	SELECT TOP 1000 u.DisplayName, mi.BadgeCount FROM dbo.Users AS u CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi;
StatementParameterizationType	0
StatementSqlHandle	0x0900B2F33A8CCFD020466
WaitStats	

Reason For Early Termination Of Statement Optimization

Reason For Early Termination Of Statement Optimization

Query executed successfully. SQL2019 (15.0 RTM) | SQL2019\Administrator ... StackOverflow | 00:26:30 | 1,000 rows

ElapsedTime	1590577
Reason For Early Termination Of State	Good Enough Plan Found
RetrievedFromCache	true
SecurityPolicyApplied	False
<input checked="" type="checkbox"/> Set Options	ANSI_NULLS: True, ANSI_PA
Statement	SELECT TOP 1000 u.DisplayI
StatementParameterizationType	0
StatementSqlHandle	0x0900B2F33A8CCFD020466
<input checked="" type="checkbox"/> WaitStats	

ElapsedTime	1590577
Reason For Early Termination Of State	Good Enough Plan Found



I'm not even mad.

NULLS: True, ANSI_PA
CT TOP 1000 u.DisplayI
00B2F33A8CCFD020466

MSTVF round up

The bad:

- Table variable guarantees at least a serial zone in the plan
- Takes several times longer than a scalar
- Some information in the execution plan, but misleading
- Some information in STATISTICS IO, but also misleading

The good:

- They get just a little better in SQL Server 2017.





It gets worse.

Functions in table definitions are horrible.

Scalar functions can be used in more places

- Computed columns
- Check constraints

Using them in either of these makes

ALL QUERIES AGAINST THE TABLE SERIAL

This includes maintenance like index rebuilds, CHECKDB.



Blogs to prove it

Still Serial After All These Years:

<https://www.brentozar.com/archive/2016/01/still-serial-after-all-these-years/>

Scalar Functions In Computed Columns:

<https://www.brentozar.com/archive/2016/01/another-reason-why-scalar-functions-in-computed-columns-is-a-bad-idea/>

Scalar Functions In Check Constraints:

<https://www.brentozar.com/archive/2016/04/another-hidden-parallelism-killer-scalar-udfs-check-constraints/>





Let's get some relief.

SQL Server 2019, or manually rewrite 'em to go inline.

After a decade, Microsoft saw the light.

When Microsoft started hosting your databases in Azure SQL DB, they suddenly realized how CPU-intensive functions were. Go figure.

SQL Server 2019 can inline some scalar functions:

<https://docs.microsoft.com/en-us/sql/relational-databases/user-defined-functions/scalar-udf-inlining>

There are a lot of limitations, so not all functions will inline:

- Can't call time-dependent functions like GETDATE()
- The UDF can't have table variables or TVPs
- The UDF can't be referenced in a GROUP BY or ORDER BY
- The UDF can't be used in a computed column, check constraint, or partitioned function



This stuff is amazeballs complicated

White paper about the internals: <https://arxiv.org/pdf/1712.00498>

Possibly extensible to C#, Java, R, Python hosted in SQL Server

Optimization of Imperative Programs in a Relational Database*

Technical Report

Karthik Ramachandra
Microsoft Gray Systems Lab
karam@microsoft.com

Alan Halverson
Microsoft Gray Systems Lab
alanh@l.microsoft.com

Kwanghyun Park
Microsoft Gray Systems Lab
kwpark@microsoft.com

César Galindo-Legaria
Microsoft
cgalarg@microsoft.com

K. Venkatesh Emani[†]
IIT Bombay
venkateshek@cs.iitb.ac.in

Conor Cunningham
Microsoft
conorc@microsoft.com

ABSTRACT

For decades, RDBMSs have supported declarative SQL as well as imperative functions and procedures as ways for users

to write programs in various languages (such as Transact-SQL, C#, Java and R) using imperative constructs such as variable assignments, conditional branching, and loops.

2.2 p42

Aug 2019

Will your functions be a good fit?

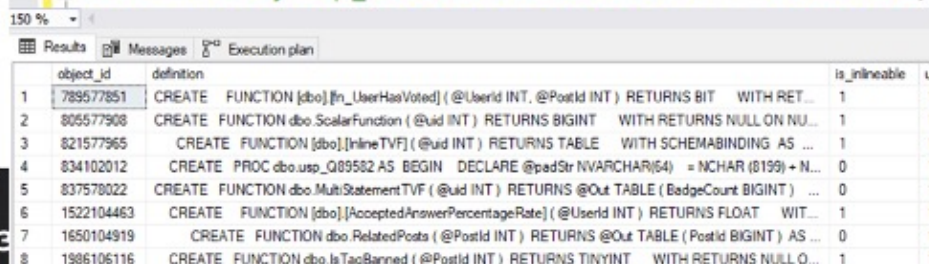
Set up a SQL Server 2019 server, restore your databases into it

Check the `is_inlineable` column in `sys.sql_modules`:

That doesn't mean it'll actually get inlined:

the decision is made at compile time, based on the query calling the UDF

```
/* Will SQL Server 2019 help? */
ALTER DATABASE CURRENT SET COMPATIBILITY_LEVEL = 150;
GO
SELECT * FROM sys.sql_modules WHERE definition LIKE '%Function%';
```



	object_id	definition	is_inlineable	u
1	789577851	CREATE FUNCTION [dbo].[fn_UserHasVoted] (@UserId INT, @PostId INT) RETURNS BIT WITH RET...	1	1
2	805577908	CREATE FUNCTION [dbo].[ScalarFunction] (@uid INT) RETURNS BIGINT WITH RETURNS NULL ON NU...	1	1
3	821577965	CREATE FUNCTION [dbo].[InlineTVF] (@uid INT) RETURNS TABLE WITH SCHEMABINDING AS ...	1	1
4	834102012	CREATE PROC [dbo].[usp_Q089582] AS BEGIN DECLARE @padStr NVARCHAR(64) = NCHAR (8199) + N...	0	1
5	837578022	CREATE FUNCTION [dbo].[MultiStatementTVF] (@uid INT) RETURNS @Out TABLE (BadgeCount BIGINT) ...	0	1
6	1522104463	CREATE FUNCTION [dbo].[AcceptedAnswerPercentageRate] (@UserId INT) RETURNS FLOAT WIT...	1	1
7	1650104919	CREATE FUNCTION [dbo].[RelatedPosts] (@PostId INT) RETURNS @Out TABLE (PostId BIGINT) AS ...	0	1
8	1986106116	CREATE FUNCTION [dbo].[IsTagBanned] (@PostId INT) RETURNS TINYINT WITH RETURNS NULL O...	1	1

2.2 p43

There are a huge list of limitations

<https://docs.microsoft.com/en-us/sql/relational-databases/user-defined-functions/scalar-udf-inlining>

A scalar T-SQL UDF can be inlined if all of the following conditions are true:

- The UDF is written using the following constructs:
 - `DECLARE`, `SET`: Variable declaration and assignments.
 - `SELECT`: SQL query with single/multiple variable assignments ¹.
 - `IF / ELSE`: Branching with arbitrary levels of nesting.
 - `RETURN`: Single or multiple return statements.
 - `UDF`: Nested/recursive function calls ².
 - Others: Relational operations such as `EXISTS`, `ISNULL`.
- The UDF does not invoke any intrinsic function that is either time-dependent (such as `GETDATE()`) or has side effects ³ (such as `NEWSEQUENTIALID()`).

- The UDF does not reference table variables or table-valued parameters.
- The query invoking a scalar UDF does not reference a scalar UDF call in its `GROUP BY` clause.
- The query invoking a scalar UDF in its select list with `DISTINCT` clause does not have `ORDER BY` clause.
- The UDF is not used in `ORDER BY` clause.
- The UDF is not natively compiled (interop is supported).
- The UDF is not used in a computed column or a check constraint definition.
- The UDF does not reference user-defined types.
- There are no signatures added to the UDF.
- The UDF is not a partition function.
- The UDF does not contain references to Common Table Expressions (CTEs).
- The UDF does not contain references to intrinsic functions that may alter the results when inlined (such as `@@ROWCOUNT`)⁴.
- The UDF does not contain aggregate functions being passed as parameters to a scalar UDF⁴.
- The UDF does not reference built-in views (such as `OBJECT_ID`)⁴.
- The UDF does not reference XML methods⁵.
- The UDF does not contain a `SELECT` with `ORDER BY` without a `TOP 1` clause⁵.
- The UDF does not contain a `SELECT` query that performs an assignment in conjunction with the `ORDER BY` clause (such as `SELECT @x = @x + 1 FROM table1 ORDER BY col1`)⁵.
- The UDF does not contain multiple `RETURN` statements⁶.
- The UDF is not called from a `RETURN` statement⁶.
- The UDF does not reference the `STRING_AGG` function⁶.
- The UDF does not reference remote tables⁷.
- The UDF-calling query does not use `GROUPING SETS`, `CUBE`, or `ROLLUP`⁷.
- The UDF-calling query does not contain a variable that is used as a UDF parameter for assignment (for example, `SELECT @y = 2, @x = UDF(@y)`)⁷.

Microsoft keeps adding more limitations

Since SQL Server 2019 released, most of the Cumulative Updates have removed inline function support for more and more functions:

<https://support.microsoft.com/en-us/help/4538581/fix-scalar-udf-inlining-issues-in-sql-server-2019>



FIX: Scalar UDF Inlining issues in SQL Server 2019

Applies to: SQL Server 2019 on Linux, SQL Server 2019 on Windows

Symptoms

User-Defined Functions (UDFs) that are implemented in Transact-SQL and that return a single data value are referred to as T-SQL Scalar User-Defined Functions (UDFs).

2.2 p46

This cumulative update also **blocks** inlining in the following scenarios:

- If the UDF references certain intrinsic functions (for example, @@ROWCOUNT) that may alter the results when inlined (added in Microsoft SQL Server 2019 CU2)
- When aggregate functions are passed as parameters to a scalar UDF (added in Microsoft SQL Server 2019 CU2)
- If the UDF references built-in views (for example: OBJECT_ID) (added in Microsoft SQL Server 2019 CU2)
- If the UDF uses XML methods (added in Microsoft SQL Server 2019 CU4)
- If the UDF contains a SELECT with ORDER BY without a "TOP 1" (added in Microsoft SQL Server 2019 CU4)
- If the SELECT query performs an assignment in conjunction with the ORDER BY clause (for example, SELECT @x = @x + 1 FROM table ORDER BY *column_name*) (added in Microsoft SQL Server 2019 CU4)
- If the UDF contains multiple RETURN statements (added in Microsoft SQL Server 2019 CU5)
- If the UDF is called from a RETURN statement (added in Microsoft SQL Server 2019 CU5)
- If the UDF references the STRING_AGG function (added in Microsoft SQL Server 2019 CU5)
- If the UDF definition references remote tables (added in Microsoft SQL Server 2019 CU6)
- If the UDF-calling query uses GROUPING SETS, CUBE, or ROLLUP (added in Microsoft SQL Server 2019 CU6)
- If the UDF-calling query contains a variable that is used as a UDF parameter for assignment (for example, SELECT @y=2, @x=UDF(@y)) (added in Microsoft SQL Server 2019 CU6)

Let's try our ScalarFunction example

It used to take ~8 minutes. In 2019 compat level: 2 minutes!

```
SELECT TOP 1000  
  u.DisplayName,  
  dbo.ScalarFunction(u.Id)  
FROM dbo.Users AS u;
```

150 %

Results Messages Execution plan

	DisplayName	(No column name)
1	Community	1
2	Jeff Atwood	328
3	Geoff Dalgas	90
4	Jarrod Dixon	130
5	Joel Spolsky	185

✓ Query executed successfully. SQL2019 (15.0 RTM) SQL2019\Administrator ... StackOverflow 00:01:56 1,000 rows

STATISTICS IO now shows what's going on

```
SELECT TOP 1000  
u.DisplayName,  
dbo.ScalarFunction(u.Id)  
FROM dbo.Users AS u;  
GO
```

But it's not good: that's a
hell of a lot of reads.

150 %

Results Messages Execution plan

SQL Server parse and compile time:
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server Execution Times:
CPU time = 0 ms, elapsed time = 0 ms.

SQL Server parse and compile time:
CPU time = 4 ms, elapsed time = 4 ms.

(1000 rows affected)

Table 'Worktable'. Scan count 1000, logical reads 83182561, physical re
Table 'Badges'. Scan count 1, logical reads 168260, physical reads 0, p
Table 'Users'. Scan count 1, logical reads 69, physical reads 0, page s

(1 row affected)

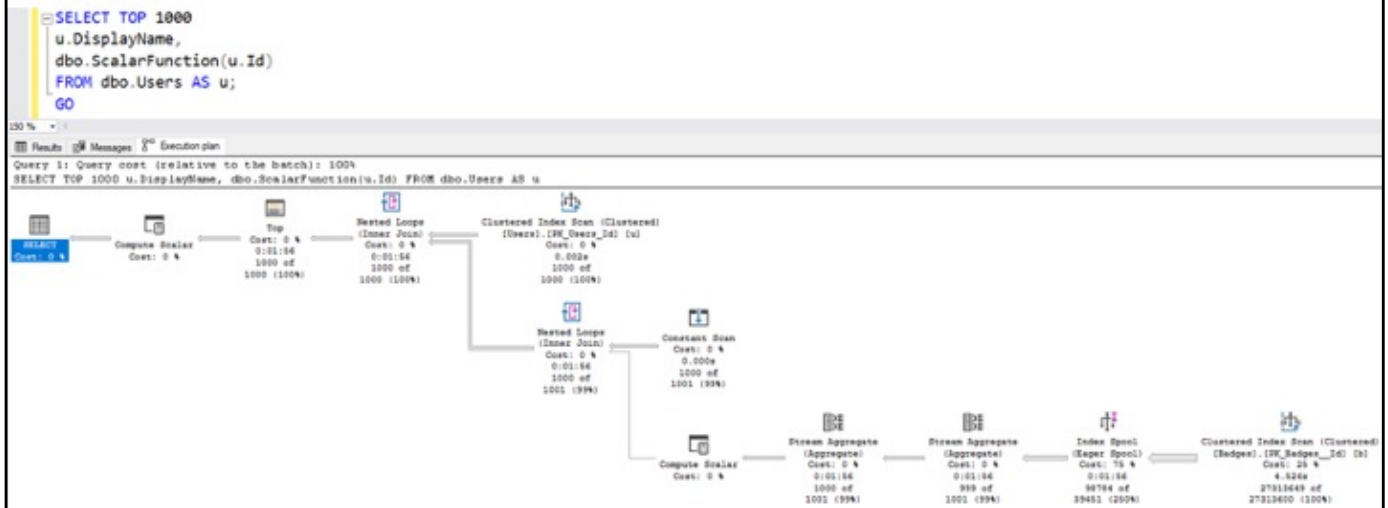
SQL Server Execution Times:
CPU time = 116157 ms, elapsed time = 116255 ms.

2.2 p49

The plan shows what's up, too.

The Badges table now shows up, but...

We get an Index Spool (Eager Spool) and no parallelism.



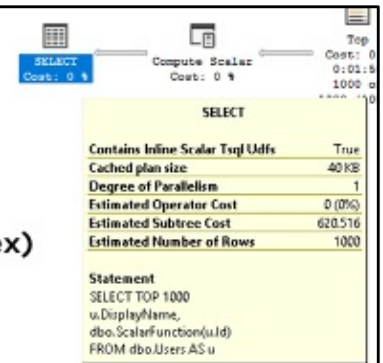
What went wrong

SQL Server 2019 inlined the function, but:

- Chose an index spool plan (that didn't suggest an index)
- Only used a single CPU core
(whereas 2017 let the function itself go parallel)
- And it wasn't due to low query cost either

If inlining doesn't work out for your function (some go slower), you can use **WITH INLINE = OFF** at the function level:

```
CREATE OR ALTER FUNCTION dbo.discount_price(@price DECIMAL(12,2), @discount DECIMAL(12,2))
RETURNS DECIMAL (12,2)
WITH INLINE = OFF
AS
BEGIN
    RETURN @price * (1 - @discount);
END
```



SELECT	
Contains Inline Scalar Udfs	True
Cached plan size	40 KB
Degree of Parallelism	1
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	620.516
Estimated Number of Rows	1000

Statement	
SELECT TOP 1000	
u.DisplayName,	
dbo.ScalarFunction(u.Id)	
FROM dbo.Users AS u	



**And you can still do way
better yourself.**

2.2 p52





Rewriting your UDFs

Switching to single-statement inline ones or APPLY

If you're facing MSTVFs that look like this:

```
-- Multi Statement Table Valued Function
CREATE FUNCTION dbo.MultiStatementTVF ( @uid INT )
RETURNS @Out TABLE ( BadgeCount BIGINT )
WITH SCHEMABINDING
AS
BEGIN
    INSERT INTO @Out (BadgeCount)
    SELECT COUNT_BIG(*) AS BadgeCount
    FROM    dbo.Badges AS b
    WHERE   b.UserId = @uid
    GROUP BY b.UserId;
    RETURN;
END;
```



And you call them like this:

```
SELECT TOP 1000
u.DisplayName,
mi.BadgeCount
FROM dbo.Users AS u
CROSS APPLY dbo.MultiStatementTVF(u.Id) AS mi
```



Then they'll fly if you make 'em look like this:

```
-- Inline Table Valued Function
CREATE FUNCTION dbo.InlineTVF ( @uid INT )
RETURNS TABLE
WITH SCHEMABINDING
AS
RETURN
    SELECT  COUNT_BIG(*) AS BadgeCount
    FROM    dbo.Badges AS b
    WHERE   b.UserId = @uid
    GROUP BY b.UserId;
```



Single-statement inline TVFs

If you can write it with a single `SELECT`, then SQL Server can inline it just like it was a view.

The good: you only have to change the function's contents, but not the way you call the function. The callers stay the same.

The bad: complex MSTVFs can turn into a hell of CASEs.

The ugly: if you can't get it into one statement, it won't perform. Your only options are things like stored procedures.



lol.exe is terminating unexpectedly

```
SELECT TOP 1000  
u.DisplayName,  
bi.BadgeCount  
FROM dbo.Users AS u  
CROSS APPLY dbo.InlineTVF(u.Id) AS bi
```



SEVEN SECONDS

Results Messages Execution plan

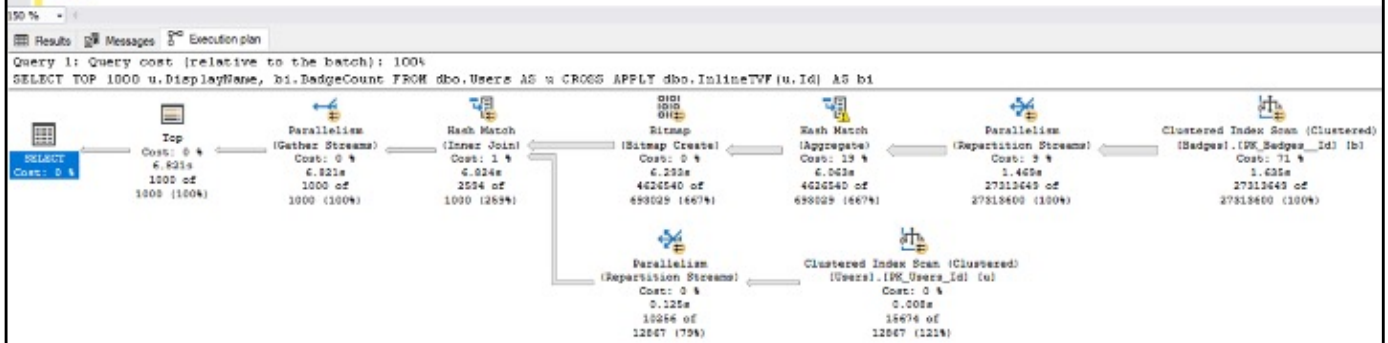
✓ Query executed...	SQL2019 (15.0 RTM)	SQL2019\Administrator ...	StackOverflow	00:00:07	1,000 rows
---------------------	--------------------	---------------------------	---------------	----------	------------

The plan gets way better

We get parallelism, and no index spool.

And if you tuned indexes, you could do even better.

```
SELECT TOP 1000  
u.DisplayName,  
bi.BadgeCount  
FROM dbo.Users AS u  
CROSS APPLY dbo.InlineTVF(u.Id) AS bi;  
GO
```



So many good benefits here

Parallelism

One call to the function

Done in seconds, not minutes

Accurate STATISTICS IO

Accurate query cost

2.2 p61



CROSS APPLY would work too

```
SELECT TOP 1000
    u.DisplayName ,
    bca.BadgeCount
FROM    dbo.Users AS u
CROSS APPLY ( SELECT b.UserId ,
                  COUNT_BIG(*) AS BadgeCount
                FROM    dbo.Badges AS b
                WHERE    u.Id = b.UserId
                GROUP BY b.UserId
            ) bca;
```

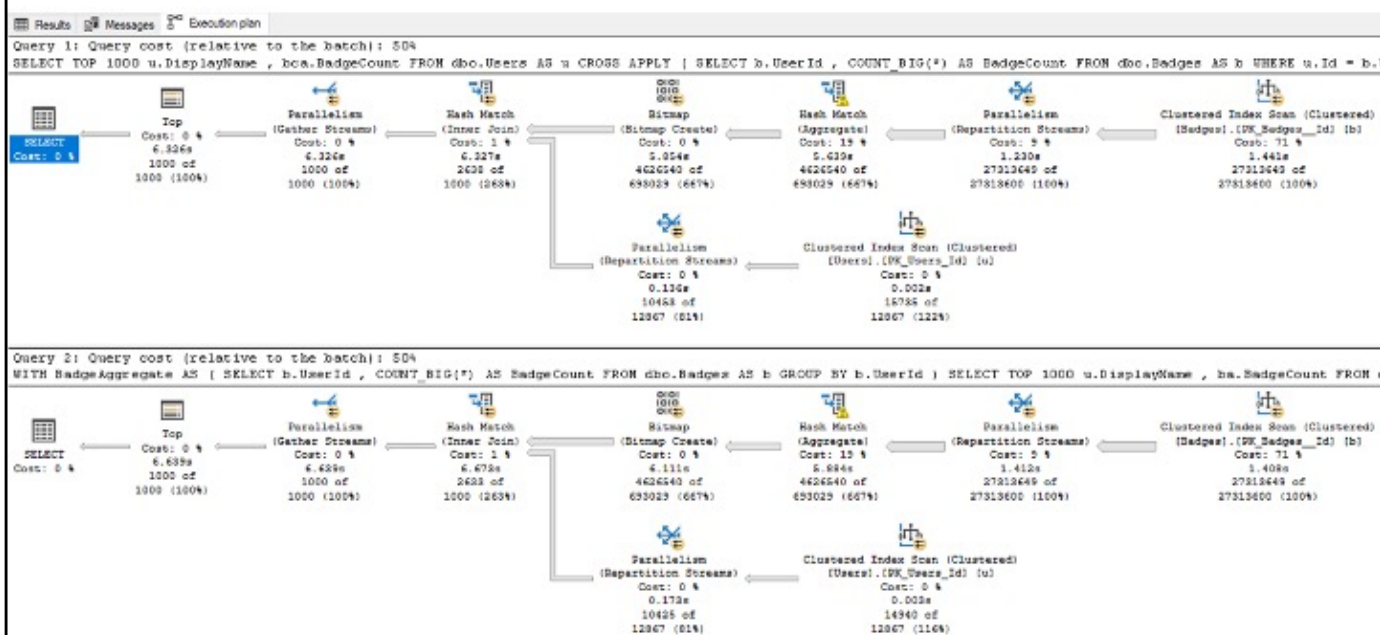


As would a CTE

```
WITH BadgeAggregate AS
( SELECT b.UserId ,
  COUNT_BIG(*) AS BadgeCount
  FROM    dbo.Badges AS b
  GROUP BY b.UserId )
SELECT TOP 1000
    u.DisplayName ,
    ba.BadgeCount
FROM    dbo.Users AS u
JOIN    BadgeAggregate AS ba
ON      ba.UserId = u.Id;
```



Same plan, same costs, etc.



But if you're facing scalar functions...

```
-- Scalar Function
CREATE FUNCTION dbo.ScalarFunction ( @uid INT )
RETURNS BIGINT
    WITH RETURNS NULL ON NULL INPUT,
    SCHEMABINDING
AS
BEGIN
    DECLARE @BCount BIGINT;
    SELECT @BCount = COUNT_BIG(*)
    FROM    dbo.Badges AS b
    WHERE   b.UserId = @uid
    GROUP BY b.UserId;
    RETURN @BCount;
END;
```



SQL Server 2017 & prior can't inline these.

Even if they only have one statement, they won't go inline.

That means you HAVE to get rid of them altogether, not just tune them.

Which is problematic since you have to touch every query that has 'em.

Thus the excitement about SQL Server 2019.





Finding the little buggers

Find queries executed frequently

```
EXEC sp_BlitzCache @SortOrder = 'xpm'  
GO  
EXEC sp_BlitzCache @SortOrder = 'executions'  
  
/* Sorts your plan cache by highest executions per  
minute or executions.  
  
When this is high, it's sometimes a sign of row-by-  
row functions being called, as we saw.  
  
http://www.brentozar.com/blitzcache/ */
```



Or you can rip through code...

```
/*You can look for bad functions like this*/

SELECT  SCHEMA_NAME(o.schema_id) AS [schema_name], o.name, o.type_desc
FROM    sys.objects AS o
WHERE   o.type IN ('FN', N'TF')
        AND o.is_ms_shipped = 0;

/*Plug the names in here to search stored proc text*/

SELECT  obj.name ,
        sc.text
FROM    sys.objects obj
INNER JOIN sys.syscomments sc
ON      sc.id = obj.object_id
WHERE   obj.type = 'P'
        AND sc.text LIKE '% FUNCTION NAME %'
```



Recap

Functions sound like a great idea at first

- Code reusability is a best practice everywhere else

But right now, functions cause problems in SQL Server

- Serializing execution
- Row by row execution

Inline table valued functions reduce these issues

- Treated like a view or CTE rather than a separate procedural task

