



**BRENT OZAR**  
UNLIMITED®

## Using Batches to Do a Lot of Work Without Blocking

3.3 p1







## When you have a lot to do, you have to make a choice.

Do it all at once in a transaction

- Good: easy to code, gets done quickly
- Bad: lock escalation means no one else can work

Or separate it into small batches

- Good: lets other people work alongside you
- Bad: harder to code, takes longer to run



**Sooner or later, you're gonna have to do the latter.**



**Or separate it into small batches**

- Good: lets other people work alongside you
- Bad: harder to code, takes longer to run



## Signs you need batching

A process takes hours to run

Other people need to keep querying while it runs

It can't be offloaded to another server,  
and you can't temporarily upsize the server

It doesn't need transactional consistency

You're willing to spend some time coding T-SQL



## Our Stack Overflow scenario

Let's say we need to delete:

- All of the Users who live in London  
(where Location = 'London, United Kingdom')
- All of their rows in related tables:  
Badges, Comments, Posts, Votes
- All of those tables have a UserId or OwnerUserId  
column, and we can join on that







```

--BEGIN TRAN;
--DELETE d
    FROM dbo.Badges d
    INNER JOIN dbo.Users u ON d.UserId = u.Id
    WHERE u.Location = 'London, United Kingdom';

--DELETE d
    FROM dbo.Comments d
    INNER JOIN dbo.Users u ON d.UserId = u.Id
    WHERE u.Location = 'London, United Kingdom';

--DELETE d
    FROM dbo.Posts d
    INNER JOIN dbo.Users u ON d.OwnerUserId = u.Id
    WHERE u.Location = 'London, United Kingdom';

--DELETE d
    FROM dbo.Votes d
    INNER JOIN dbo.Users u ON d.UserId = u.Id
    WHERE u.Location = 'London, United Kingdom';

/* No need for error checking here since it's in one tran: */
--DELETE
    FROM dbo.Users
    WHERE Location = 'London, United Kingdom';

/* Commenting this out only so we can see the effects: */
--COMMIT;

```

**The easy way:  
one big transaction**

## I have indexes to help

```
CREATE INDEX Location ON dbo.Users(Location);  
GO  
CREATE INDEX UserId ON dbo.Badges(UserId);  
GO  
CREATE INDEX UserId ON dbo.Comments(UserId);  
GO  
CREATE INDEX OwnerUserId ON dbo.Posts(OwnerUserId);  
GO  
CREATE INDEX UserId ON dbo.Votes(UserId);  
GO
```



## But one big transaction sucks.

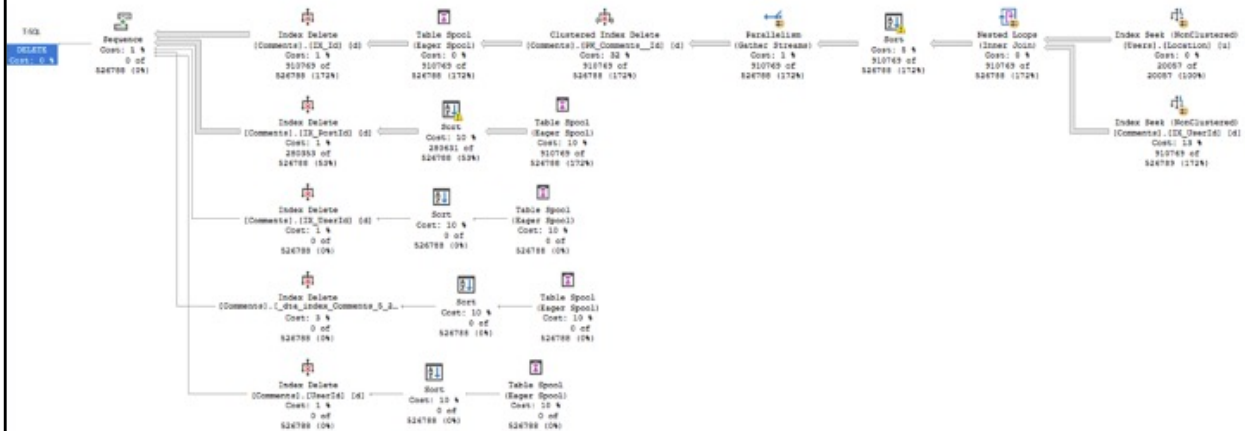
It takes a while to run.

While it's running:

- We're holding exclusive table-level locks on all the tables, not just one at a time
- The transaction log and the version store can grow out of control
- SQL Server deletes each index serially



# As the Comments delete runs...

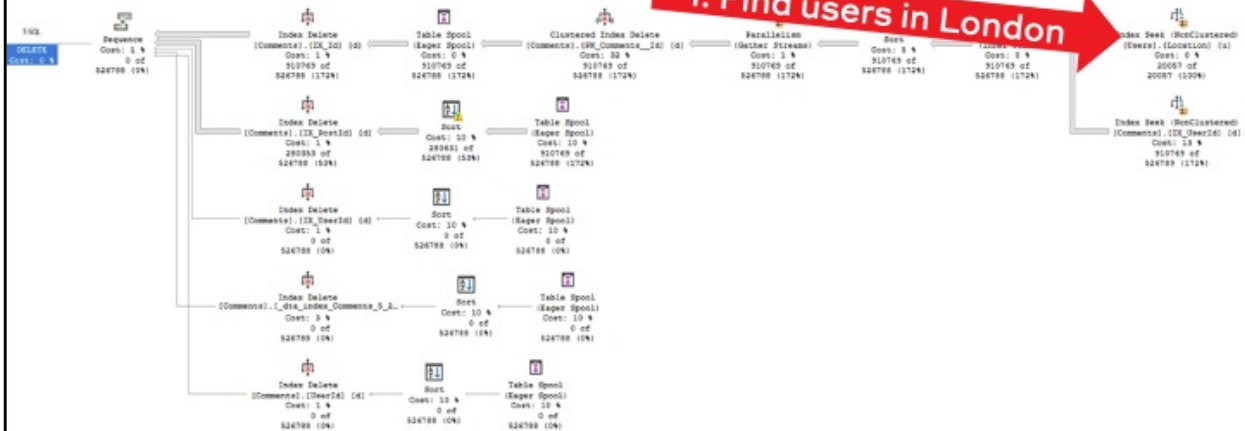


3.3 p13



# Step 1: Find the Users in London

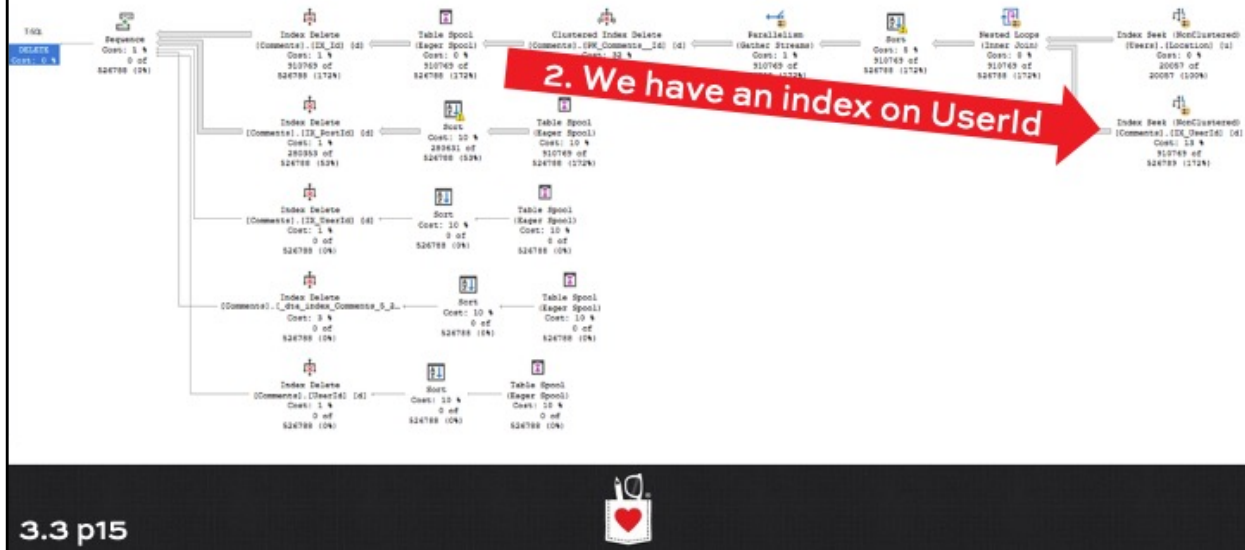
1. Find users in London



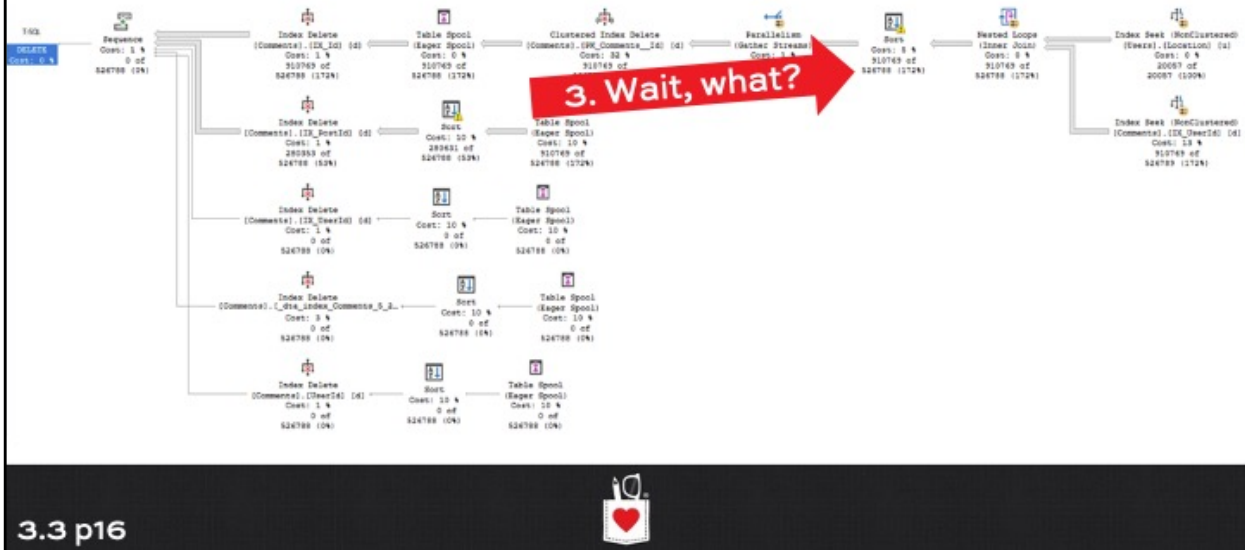
3.3 p14



## Step 2: Find their Comments



## Step 3: Sort them by Comment Id





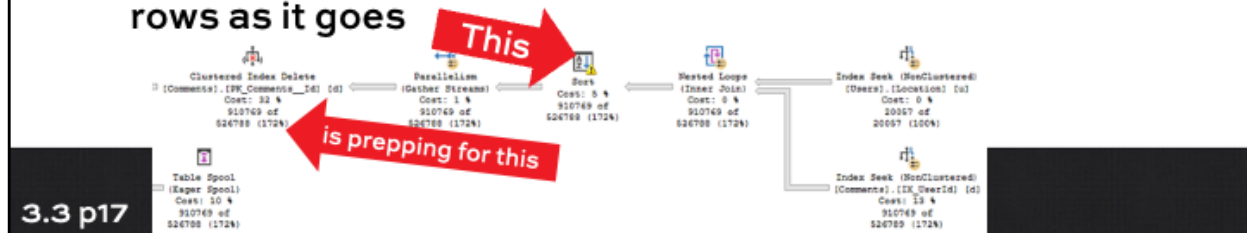
## The delete is going to be hard, so

It's going to involve deleting a "lot" of Comments

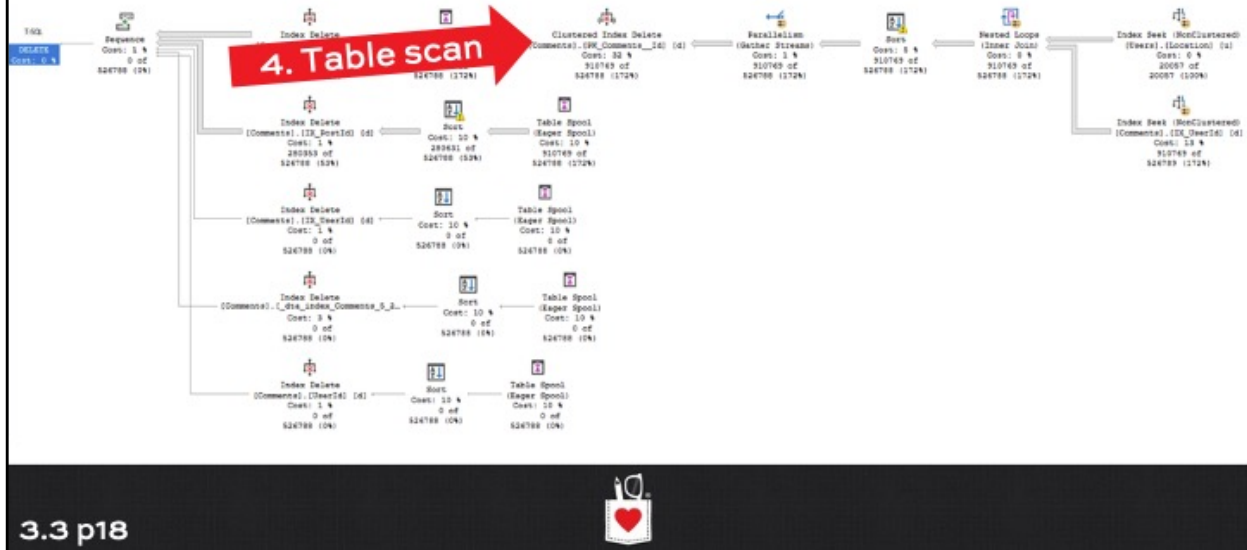
The clustered index of Comments is Id (not UserId)

To efficiently find them, SQL Server wants to sort all of the comments from Londoners, by Comment Id

That way it can scan the clustered index, deleting rows as it goes



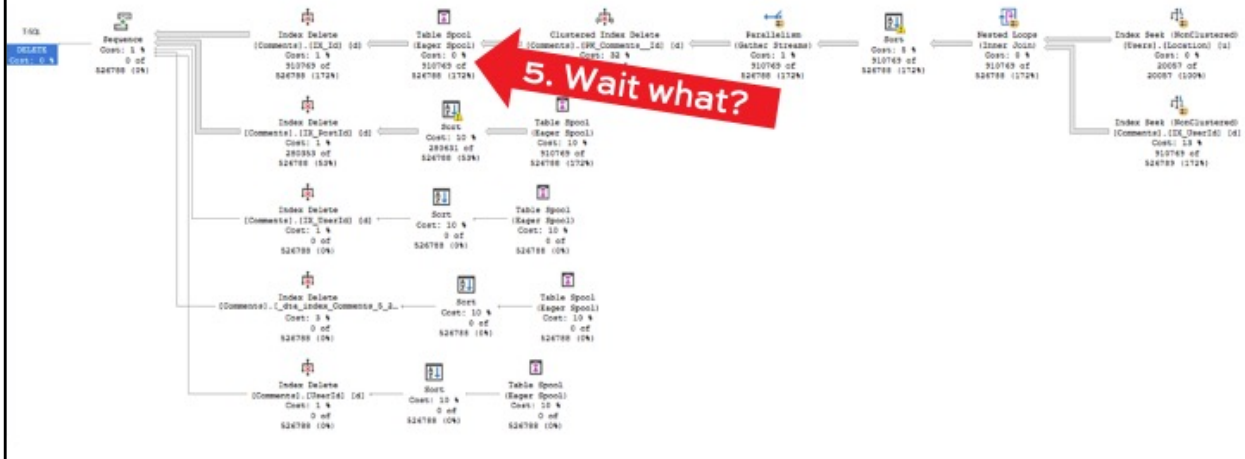
## Step 4: Scan the CX, delete rows



3.3 p18



## Step 5: Sort all London rows again



3.3 p19



## Deleting each index is hard too

It's going to involve deleting a "lot" of rows

So each index that we need to delete, SQL Server sorts all of the to-be-deleted rows by the index's keys

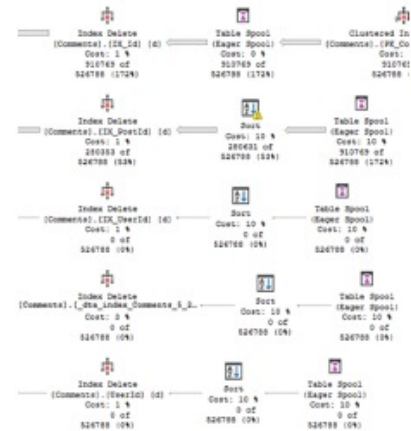
This happens one index at a time, in order



# All this work is single-threaded

Deletes, updates, and inserts are not normally parallelized:

<https://www.brentozar.com/archive/2020/09/update-insert-and-delete-are-not-normally-processed-in-parallel/>



## And we exclusively locked it all

As this query runs, it progressively grabs locks on more and more tables.

My favorite way to see it:

```
1 sp_WhoIsActive @get_locks = 1
```

150 %

Results Messages

	dd hh:mm:ss.mss	session_id	sql_text	locks
1	00 00:08:19.653	53	<?query -- BEGIN TRAN; DELETE d FROM dbo.Bad...	<Database name="StackOverflow"><Locks><Lock requ...

3.3 p22



```

1 <Database name="StackOverflow">
2   <Locks>
3     <Lock request_mode="S" request_status="GRANT" request_count="1" />
4   </Locks>
5   <Objects>
6     <Object name="(null)">
7       <Locks>
8         <Lock resource_type="EXTENT" page_type="*" request_mode="X" request_status="GRANT" request_count="803" />
9       </Locks>
10    </Object>
11    <Object name="Badges" schema_name="dbo">
12      <Locks>
13        <Lock resource_type="OBJECT" request_mode="X" request_status="GRANT" request_count="1" />
14      </Locks>
15    </Object>
16    <Object name="Comments" schema_name="dbo">
17      <Locks>
18        <Lock resource_type="OBJECT" request_mode="X" request_status="GRANT" request_count="1" />
19      </Locks>
20    </Object>
21    <Object name="Posts" schema_name="dbo">
22      <Locks>
23        <Lock resource_type="OBJECT" request_mode="X" request_status="GRANT" request_count="1" />
24        <Lock resource_type="PAGE" page_type="*" index_name="PK_Posts_Id" request_mode="X" request_status="GRANT" />
25      </Locks>
26    </Object>
27    <Object name="Users" schema_name="dbo">
28      <Locks>
29        <Lock resource_type="OBJECT" request_mode="X" request_status="GRANT" request_count="1" />
30      </Locks>
31    </Object>
32    <Object name="Votes" schema_name="dbo">
33      <Locks>
34        <Lock resource_type="OBJECT" request_mode="X" request_status="GRANT" request_count="1" />

```

3.3 p23





## Pseudocode

Create a table with a list of UserIds to delete

For each child table:

- Create a temp table with a list of its Ids to delete
- While rows exist in that list, delete 1,000 of them (staying well below lock escalation thresholds)



```

Part 1: gathering the keys. No exclusive locks for this. */
DROP TABLE IF EXISTS dbo.UsersToDelete;
DROP TABLE IF EXISTS dbo.BadgesToDelete;
DROP TABLE IF EXISTS dbo.CommentsToDelete;
DROP TABLE IF EXISTS dbo.PostsToDelete;
DROP TABLE IF EXISTS dbo.VotesToDelete;

CREATE TABLE dbo.UsersToDelete (Id INT PRIMARY KEY CLUSTERED);
INSERT INTO dbo.UsersToDelete (Id)
    SELECT Id
    FROM dbo.Users
    WHERE Location = 'London, United Kingdom';
GO

CREATE TABLE dbo.BadgesToDelete (Id INT PRIMARY KEY CLUSTERED);
INSERT INTO dbo.BadgesToDelete (Id)
    SELECT d.Id
    FROM dbo.Badges d
    INNER JOIN dbo.UsersToDelete u
        ON u.Id = d.UserId;

```

3.3 p26

## This is not a transaction.

This could be problematic if users are changing their locations while we work:

```
CREATE TABLE dbo.UsersToDelete (Id INT PRIMARY KEY CLUSTERED);  
INSERT INTO dbo.UsersToDelete (Id)  
    SELECT Id  
    FROM dbo.Users  
    WHERE Location = 'London, United Kingdom';
```

If users move into London, you could keep this table up to date with a trigger – but let's keep this simple.



## Same thing with the child tables.

If London users are continuing to add Badges, they won't show up in this table.

```
]CREATE TABLE dbo.BadgesToDelete (Id INT PRIMARY KEY CLUSTERED);  
]INSERT INTO dbo.BadgesToDelete (Id)  
    SELECT d.Id  
    FROM dbo.Badges d  
    INNER JOIN dbo.UsersToDelete u  
        ON u.Id = d.UserId;
```

Easy fix: just repeat the deletion process again at the end to catch any newly added Badges.



## Don't use a temp table for this.

A real table will:

- Be query-able from other sessions to check progress on the deletes while they run
- Persist after reboots & failovers, helpful for really long-running jobs
- Allow you to parallelize it across many sessions

```
CREATE TABLE dbo.BadgesToDelete (Id INT PRIMARY KEY CLUSTERED);
INSERT INTO dbo.BadgesToDelete (Id)
SELECT d.Id
FROM dbo.Badges d
INNER JOIN dbo.UsersToDelete u
ON u.Id = d.UserId;
```

3.3 p29

## Back to the big picture



Create a table with a list of UserIds to delete

For each child table:



Create a temp table with a list of its Ids to delete



- While rows exist in that list, delete 1,000 of them (staying well below lock escalation thresholds)



```

DECLARE @FirstId BIGINT = -9223372036854775807;

:WHILE EXISTS (
    SELECT del.Id
    FROM dbo.BadgesToDelete td
    INNER JOIN dbo.Badges del ON td.Id = del.Id)
:BEGIN
:    WITH ToBeDeleted AS (
        SELECT TOP 1000 td.Id
        FROM dbo.BadgesToDelete td
        WHERE td.Id >= @FirstId
        ORDER BY td.Id
    )
    DELETE d
        FROM ToBeDeleted tbd
        INNER JOIN dbo.Badges d ON tbd.Id = d.Id;

    /* If our FirstId filtered out all rows, reset it, something went wrong: */
    IF @@ROWCOUNT = 0 SET @FirstId = -9223372036854775807;

    /* Reset our low key for the next pass: */
:    SELECT TOP 1 @FirstId = td.Id
        FROM dbo.BadgesToDelete td
        INNER JOIN dbo.Badges del ON td.Id = del.Id
        WHERE td.Id >= @FirstId
        ORDER BY td.Id;
END;

```

3.3 p31

```
DECLARE @FirstId BIGINT = -9223372036854775807;  
  
:WHILE EXISTS (  
    SELECT del.Id  
    FROM dbo.BadgesToDelete td  
    INNER JOIN dbo.Badges del ON td.Id = del.Id)
```

When this doesn't return rows, we're done.

Very lightweight query: Ids only, low reads.



Now the locking is about to start, so we have to be careful about getting few rows, and touching as few indexes as possible.

```
3 WITH ToBeDeleted AS (  
    SELECT TOP 1000 td.Id  
    FROM dbo.BadgesToDelete td  
    WHERE td.Id >= @FirstId  
    ORDER BY td.Id  
)
```

The TOP 1000 with an ORDER BY means SQL Server knows exactly how many rows will come out, so it helps avoid lock escalation.

We don't have to join it to the Badges table to figure out which rows are left to delete.

**We take the 1,000 rows from the CTE:**

```
3  WITH ToBeDeleted AS (  
    SELECT TOP 1000 td.Id  
    FROM dbo.BadgesToDelete td  
    WHERE td.Id >= @FirstId  
    ORDER BY td.Id  
  )  
  DELETE d  
    FROM ToBeDeleted tbd  
    INNER JOIN dbo.Badges d ON tbd.Id = d.Id;
```

**And join directly to Badges using the Id.**

**SQL Server knows exactly how many rows will be involved, so this avoids lock escalation.**

Set the @FirstId using a lightweight read-only query, touching as few rows as possible, again to avoid table scans:

```
/* Reset our low key for the next pass: */  
3 SELECT TOP 1 @FirstId = td.Id  
   FROM dbo.BadgesToDelete td  
   INNER JOIN dbo.Badges del ON td.Id = del.Id  
   WHERE td.Id >= @FirstId  
   ORDER BY td.Id;  
END;
```

3.3 p35

```

DECLARE @FirstId BIGINT = -9223372036854775807;

:WHILE EXISTS (
    SELECT del.Id
    FROM dbo.BadgesToDelete td
    INNER JOIN dbo.Badges del ON td.Id = del.Id)
:BEGIN
:    WITH ToBeDeleted AS (
        SELECT TOP 1000 td.Id
        FROM dbo.BadgesToDelete td
        WHERE td.Id >= @FirstId
        ORDER BY td.Id
    )
    DELETE d
        FROM ToBeDeleted tbd
        INNER JOIN dbo.Badges d ON tbd.Id = d.Id;

    /* If our FirstId filtered out all rows, reset it, something went wrong: */
    IF @@ROWCOUNT = 0 SET @FirstId = -9223372036854775807;

    /* Reset our low key for the next pass: */
:    SELECT TOP 1 @FirstId = td.Id
        FROM dbo.BadgesToDelete td
        INNER JOIN dbo.Badges del ON td.Id = del.Id
        WHERE td.Id >= @FirstId
        ORDER BY td.Id;
END;

```

3.3 p36

```
1 DECLARE @firstId BIGINT = -9223372036854775807;
2
3 WHILE EXISTS (
4     SELECT del.Id
5     FROM dbo.BadgesToDelete td
6     INNER JOIN dbo.Badges del ON td.Id = del.Id)
7 BEGIN
8     WITH ToBeDeleted AS (
9         SELECT TOP 1000 td.Id
10        FROM dbo.BadgesToDelete td
11        WHERE td.Id >= @firstId
12        ORDER BY td.Id
13    )
14    DELETE d
15    FROM ToBeDeleted tbd
16    INNER JOIN dbo.Badges d ON tbd.Id = d.Id;
17
18    /* If our firstId filtered out all rows, reset it, something went wrong: */
19    IF @@ROWCOUNT = 0 SET @firstId = -9223372036854775807;
20
21    /* Reset our low key for the next pass: */
22    SELECT TOP 1 @firstId = td.Id
23    FROM dbo.BadgesToDelete td
24    INNER JOIN dbo.Badges del ON td.Id = del.Id
25    WHERE td.Id >= @firstId
26    ORDER BY td.Id;
27
28 END;
```

100 %

Messages

(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)  
(1000 rows affected)

## Run it...

This can be parallelized:

- Each table in its own session
- Or multiple ToBeDeleted tables, broken across sessions

Let's check the locks...



## Run sp\_WhoIsActive repeatedly

The locks will be changing every few milliseconds:

```
1  sp_WhoIsActive @get_locks = 1
2  GO 10
```

150 %

	dd hh:mm:ss.mss	session_id	sql_text	locks
1	00:00:00.01.573	54	<?query - WHILE EXISTS ( SELECT del.id FROM ...	<Database name="StackOverflow"><Locks><Lock requ...
<				
1	00:00:00.02.527	54	<?query - WHILE EXISTS ( SELECT del.id FROM ...	<Database name="StackOverflow"><Locks><Lock requ...
<				
1	00:00:00.03.437	54	<?query - WHILE EXISTS ( SELECT del.id FROM ...	<Database name="StackOverflow"><Locks><Lock requ...
<				
1	00:00:00.04.557	54	<?query - WHILE EXISTS ( SELECT del.id FROM ...	<Database name="StackOverflow"><Locks><Lock requ...
<				

3.3 p38

## If we did our job right...

There won't be any eXclusive table-level locks, only row-level or page-level locks:

```
<Database name="StackOverflow">
  <Locks>
    <Lock request_mode="S" request_status="GRANT" request_count="1" />
  </Locks>
  <Objects>
    <Object name="Badges" schema_name="dbo">
      <Locks>
        <Lock resource_type="KEY" index_name="_dta_index_Badges_5_2105058535_K3_K2_K4" request_mode="X" request_status="GRANT" request_count="52" />
        <Lock resource_type="KEY" index_name="IX_Id" request_mode="U" request_status="GRANT" request_count="69" />
        <Lock resource_type="KEY" index_name="IX_Id" request_mode="X" request_status="GRANT" request_count="189" />
        <Lock resource_type="KEY" index_name="IX_UserId" request_mode="X" request_status="GRANT" request_count="72" />
        <Lock resource_type="KEY" index_name="PK_Badges_Id" request_mode="X" request_status="GRANT" request_count="179" />
        <Lock resource_type="KEY" index_name="UserId" request_mode="X" request_status="GRANT" request_count="19" />
        <Lock resource_type="OBJECT" request_mode="IS" request_status="GRANT" request_count="1" />
        <Lock resource_type="PAGE" page_type="" index_name="_dta_index_Badges_5_2105058535_K3_K2_K4" request_mode="IX" request_status="GRANT" request_count="1" />
        <Lock resource_type="PAGE" page_type="" index_name="IX_Id" request_mode="IS" request_status="GRANT" request_count="1" />
        <Lock resource_type="PAGE" page_type="" index_name="IX_UserId" request_mode="IX" request_status="GRANT" request_count="7" />
        <Lock resource_type="PAGE" page_type="" index_name="PK_Badges_Id" request_mode="IX" request_status="GRANT" request_count="7" />
      </Locks>
    </Object>
    <Object name="BadgesToDelete" schema_name="dbo">
      <Locks>
        <Lock resource_type="OBJECT" request_mode="IS" request_status="GRANT" request_count="1" />
        <Lock resource_type="PAGE" page_type="" index_name="PK_BadgesToDelete_32146C07B41C08C1" request_mode="IS" request_status="GRANT" request_count="1" />
      </Locks>
    </Object>
  </Objects>
</Database>
```

3.3 p39

## ENHANCE

X = exclusive locks

X = OK on keys, but not okay on entire objects.

```
<Lock resource_type="KEY" index_name="_dta_index_Badges_5_2105058535__K3_K2_K4" request_mode="X" request_status="GRANT" request_count="1" />
<Lock resource_type="KEY" index_name="IX_Id" request_mode="U" request_status="GRANT" request_count="1" />
<Lock resource_type="KEY" index_name="IX_Id" request_mode="X" request_status="GRANT" request_count="1" />
<Lock resource_type="KEY" index_name="IX_UserId" request_mode="X" request_status="GRANT" request_count="1" />
<Lock resource_type="KEY" index_name="PK_Badges__Id" request_mode="Y" request_status="GRANT" request_count="1" />
<Lock resource_type="KEY" index_name="UserId" request_mode="IS" request_status="GRANT" request_count="1" />
<Lock resource_type="OBJECT" request_mode="IS" request_status="GRANT" request_count="1" />
<Lock resource_type="PAGE" page_type="*" index_name="_dta_index_Badges_5_2105058535__K3_K2_K4" request_mode="X" request_status="GRANT" request_count="1" />
<Lock resource_type="PAGE" page_type="*" index_name="IX_Id" request_mode="IS" request_status="GRANT" request_count="1" />
<Lock resource_type="PAGE" page_type="*" index_name="IX_UserId" request_mode="IX" request_status="GRANT" request_count="1" />
<Lock resource_type="PAGE" page_type="*" index_name="PK_Badges__Id" request_mode="IX" request_status="GRANT" request_count="1" />
```

OK

IS = Intent Shared

3.3 p40





## This avoids blocking.

```
<Lock resource_type="KEY" index_name="UserId" request_mode="X" request_status="GRANT" request_count="1" />  
<Lock resource_type="OBJECT" request_mode="IS" request_status="GRANT" request_count="1" />
```

Key = row-level

IS = Intent Shared

We take out exclusive locks on the deleted rows, but:

- Only those 1,000 rows (not hitting lock escalation to lock the entire table)
- We're releasing those 1,000 locks as soon as our delete finishes



## This is easy to get wrong.

If you try to do all this in a transaction,  
you'll escalate to table-level locking.

If you do this with row-level lock hints,  
your SQL Server can run out of memory.

If you don't have the supporting indexes,  
you'll hit a table scan and lock escalation.

Start with my code, modify it as needed,  
and test it in a dev server with prod scale data.



## How to test it

Run the batches just like you plan to run in prod.

You don't need the same server horsepower:  
your code should behave the same on any server size.

While it runs, run `sp_WhoIsActive @get_locks GO 10`.

Examine the Locks column, and if you see this, you're  
hitting lock escalation – tweak the code/indexes:


```
<Object name="Badges" schema_name="dbo">
```

```
<Locks>
```

```
<Lock resource_type="OBJECT" request_mode="X">
```



X = BAD



## You have the choice. The answer isn't always batching.

Do it all at once in a transaction

- Good: easy to code, gets done quickly
- Bad: lock escalation means no one else can work

**Or** separate it into small batches

- Good: lets other people work alongside you
- Bad: harder to code, takes longer to run

Batching isn't easier: it's hard. It's one of your tools.



## My scenario was deletes.

But this same thought process holds true for:

- Trickle loading data: track which data needs to be inserted in a table
- Mass updates: track which rows need to be updated, and what changes need to be made



## When you do need to do batching:

Avoid the ~5,000 row lock escalation threshold:

<https://www.littlekendra.com/2017/04/03/which-locks-count-toward-lock-escalation/>

Design your batching code to run on the clustered indexes of the tables involved.

Use the fast-ordered-delete technique of CTEs or views so SQL Server understands how many rows will be involved:

<https://www.brentozar.com/archive/2018/04/how-to-delete-just-some-rows-from-a-really-big-table/>



## Code is yours, and there's no lab.

The full demo scripts are in this video's page.

There's no test on this in the labs because:

- This is about processes that take a long time
- You're probably not going to need to jump into this technique next week
- When you *\*do\** need it, you'll need to customize it based on your own database schema anyway

