



BRENT OZAR
UNLIMITED®

PAGEIOLATCH: Reading from Disk

Also relevant for folks who have
CXCONSUMER, CXPACKET, and LATCH_EX
as their top combo.

1.2 p1

Preparation

Set your max memory to 1GB, parallelism to defaults:

```
EXEC sys.sp_configure N'max server memory (MB)', N'1024';  
EXEC sys.sp_configure N'cost threshold for parallelism', N'5';  
EXEC sys.sp_configure N'max degree of parallelism', N'0';
```

Restart your SQL Server

Drop all nonclustered indexes on
StackOverflow.dbo.Comments

Set up a new window for sp_BlitzFirst:

```
sp_BlitzFirst @ExpertMode = 1, @Seconds = 30
```




Simple Select

```
SELECT *  
FROM StackOverflow.dbo.Comments  
WHERE UserId = 26837
```



PAGEIOLATCH: Reading the data file from disk



1.2 p5



Data you have to read
/ How fast you read
Minimum query runtime

1.2 p6



Reading millions of 8KB pages

```
SELECT * FROM dbo.Comments WHERE UserId = 26837;
```

150 %



Results



Messages

(150 rows affected)
Table 'Comments'. Scan count 9, logical reads 2765260,
physical reads 13, read-ahead reads 2762692,
lob logical reads 0, lob physical reads 0, lob read-ahead reads 0.

2,765,260 pages * 8KB each / 1,024 = 21.6GB

1.2 p7



So how do we tune this?

**Data you have to read
/ How fast you read
Minimum query runtime**

1.2 p9



How to reduce PAGEIOLATCH

1. Tune indexes, looking at missing indexes:
`sp_BlitzIndex @GetAllDatabases = 1`
2. Tune queries:
`sp_BlitzCache @SortOrder = 'reads'`
3. Add more memory
4. Make storage faster



What index could you make?

```
/* Our query: */  
SELECT *  
FROM dbo.Comments  
WHERE UserId = 26837
```



Fixing it with indexes

```
SELECT * FROM dbo.Comments  
WHERE UserId = 26837
```

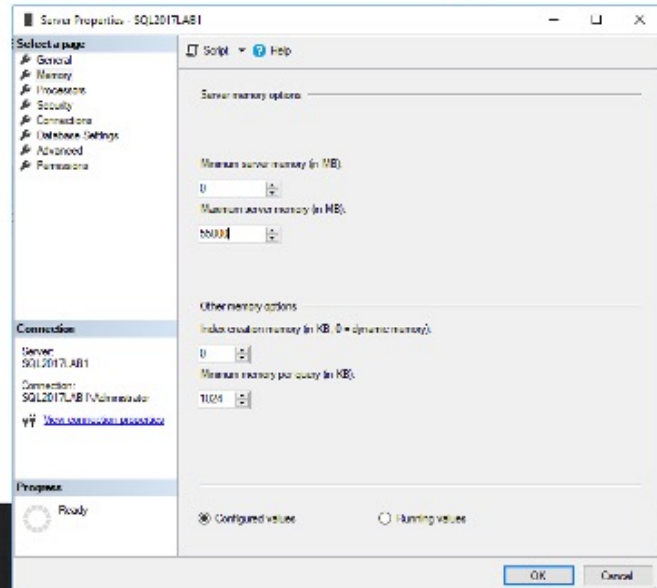
```
/* Our index: */  
CREATE INDEX IX_UserId ON  
dbo.Comments (UserId);
```



Before you do that, add memory.

Otherwise, trust me,
the index build takes
forever.

My default: leave 4GB
or 10% free for the OS,
whichever is greater.



1.2 p13

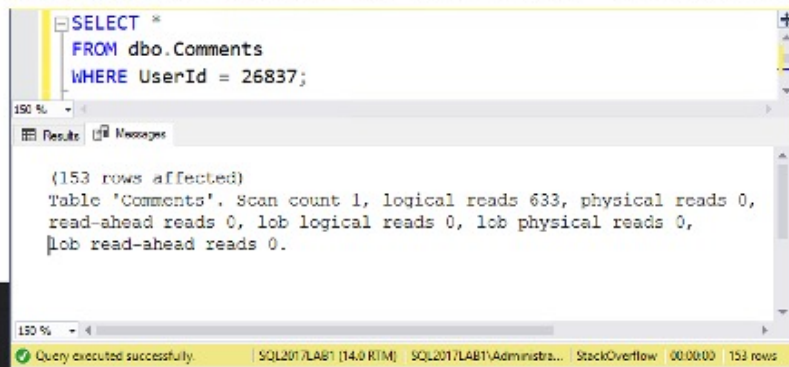
With the index, we read less pages

Under 1K pages down from millions, and finishes in a second

Frees up our memory for other queries

Reduces the workload on our slow storage

We don't have to cover the whole query – just one column helps



The screenshot shows a SQL Server Enterprise Manager window. The top pane displays a query: `SELECT * FROM dbo.Comments WHERE UserId = 26837;`. The bottom pane shows the execution plan and results. The results pane indicates that 153 rows were affected. The execution plan shows a single index seek operation on the 'Comments' table, with a scan count of 1, logical reads of 633, and physical reads of 0. The status bar at the bottom indicates that the query was executed successfully.

```
SELECT *  
FROM dbo.Comments  
WHERE UserId = 26837;
```

Results Messages

(153 rows affected)
Table 'Comments'. Scan count 1, logical reads 633, physical reads 0,
read-ahead reads 0, lob logical reads 0, lob physical reads 0,
lob read-ahead reads 0.

Query executed successfully. SQL2017/LAB1 [14.0 RTM] SQL2017/LAB1/Administra... StackOverflow 00:00:00 153 rows

1.2 p14

With the index, we read less pages

Pattern	Sample End	Seconds Sample	wait_type	wait_category	Wait Time (Seconds)	Per Core Per Second	Signal Wait Time (Seconds)	Percent Signal Waits	Number of Waits	Average Per Wait
1 WAIT STATS	2018-02-13 22:33:30.870294 +00:00	25	ASYNC_NETWORK_IO	Network IO	0.4	0.0	0.0	0.0	22	18.9
2 WAIT STATS	2018-02-13 22:33:30.870294 +00:00	25	MEMORY_ALLOCATION_EXT	Memory	0.0	0.0	0.0	0.0	5391	0.0
3 WAIT STATS	2018-02-13 22:33:30.870294 +00:00	25	PAGELATCH_SH	Buffer IO	0.0	0.0	0.0	0.0	19	0.5
4 WAIT STATS	2018-02-13 22:33:30.870294 +00:00	25	PREEMPTIVE_OS_WRITEFILE	Preemptive	0.0	0.0	0.0	0.0	1	1.0

Pattern	Sample Time	Seconds Sample	File Name	Drive	# Reads/Writes	MB Read/Written	Avg Stall (ms)	File physical name
1 PHYSICAL READS	2018-02-13 22:33:30.885887 +00:00	25	tempdev (ROWS)	Z:	3	0.2	0	Z:\MSSQL\DATA\tempdev.mdf
2 PHYSICAL READS	2018-02-13 22:33:30.885887 +00:00	25	MSDBdev (ROWS)	D:	8	0.5	0	D:\MSSQL14\MSSQLSERVER\MSSQL\DATA\MSDBdev.mdf
3 PHYSICAL READS	2018-02-13 22:33:30.885887 +00:00	25	StackOverflow_1 (ROWS)	Z:	30	2.4	0	Z:\MSSQL14\MSSQLSERVER\MSSQL\DATA\StackOverflow_1.mdf

Waits down under 1 second.

The quality of the storage simply no longer matters.

Storage stalls disappear because less work means faster average response time per task.



How to reduce PAGEIOLATCH

1. ~~Tune indexes, looking at missing indexes:~~

~~`sp_BlitzIndex @GetAllDatabases = 1`~~

2. Tune queries:

`sp_BlitzCache @SortOrder = 'reads'`

3. Add more memory
4. Make storage faster

But if we can't tune indexes



I'd rather tune indexes first.

The right indexes make *lots* of queries faster.

Tuning individual queries can take a lot longer.



Can we “fix” this query?

```
SELECT * FROM dbo.Comments  
WHERE UserId = 26837
```

1. Don't select columns you don't need
2. Don't select rows you don't need (paginate)
3. Cache in the application layer

We cover these in Mastering Query Tuning.



How to reduce PAGEIOLATCH

1. ~~Tune indexes, looking at missing indexes:~~

~~`sp_BlitzIndex @GetAllDatabases = 1`~~

2. ~~Tune queries:~~

~~`sp_BlitzCache @SortOrder = 'reads'`~~

3. Add more memory

But if we can't fix those...

4. Make storage faster



Dr. Phil says,
**you don't solve
money problems
with money.**

1.2 p20



Dr. Brent says,
**you don't solve
storage problems
with storage.**

1.2 p21



Dr. Brent says,
**you solve
storage problems
with memory.**

1.2 p22



Systems Performance by Brendan Gregg

1 CPU cycle	0.3 ns
Level 1 cache access	0.9 ns
Level 2 cache access	2.8 ns
Level 3 cache access	12.9 ns
Main memory access	120 ns
Solid-state disk I/O	50-150 μ s
Rotational disk I/O	1-10 ms
Internet: SF to NYC	40 ms
Internet: SF to UK	81 ms
Internet: SF to Australia	183 ms

1.2 p23



Systems Performance by Brendan Gregg

1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access	120 ns	6 min
Solid-state disk I/O	50-150 μ s	2-6 days
Rotational disk I/O	1-10 ms	1-12 months
Internet: SF to NYC	40 ms	4 years
Internet: SF to UK	81 ms	8 years
Internet: SF to Australia	183 ms	19 years

1.2 p24



How much memory do we need?

dbo.Comments table: about 22GB

But that's just this one table, and not:

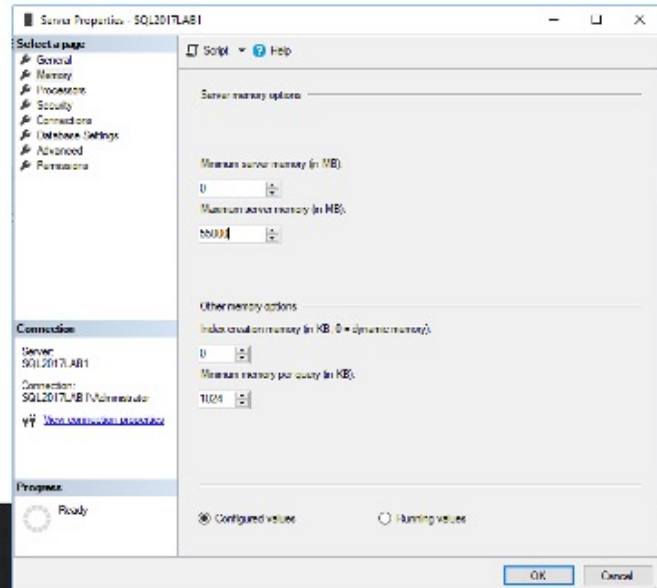
- Other tables we query
- Caching execution plans
- Query workspace memory for sorts, joins, etc.



Take off the handcuffs

If you temporarily restricted memory, crank it back up to high.

In our lab, this lets us cache the entire table in memory.



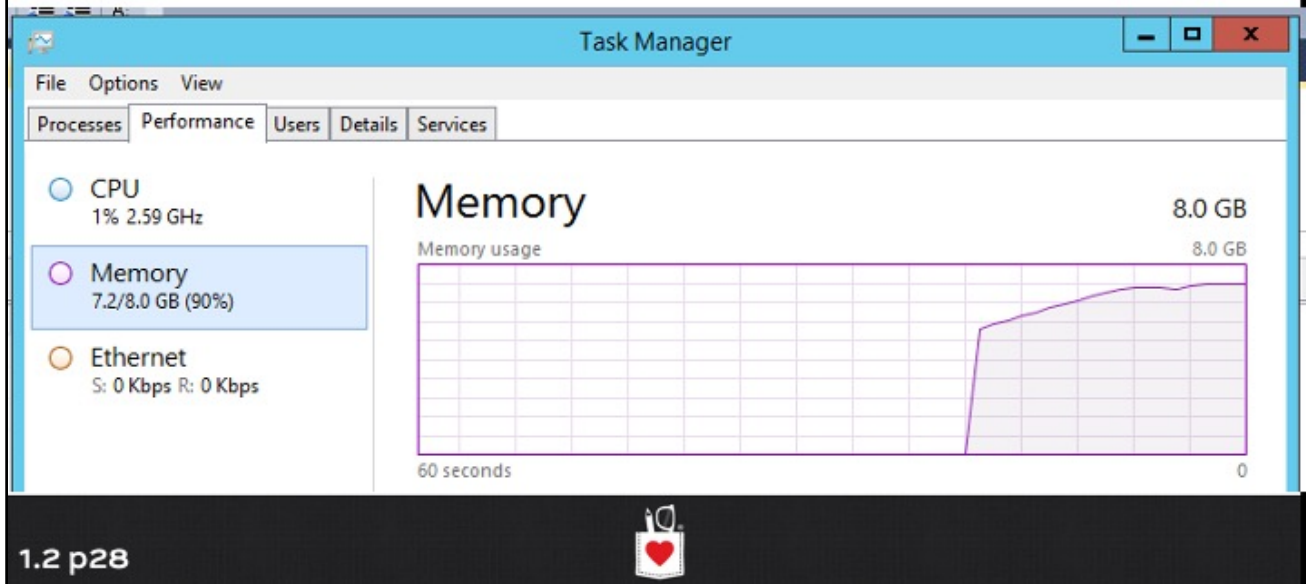
1.2 p26

Then run the query a few times.

```
/* Our query: */  
SELECT *  
FROM dbo.Comments  
WHERE UserId = 26837  
GO 5
```



SQL starts using more memory



The art of fixing it with RAM

You don't have to cache the whole table in memory.

But the more you cache, the less storage has to work.

1.2 p29



How much RAM to throw at it

Physical boxes:

- 2012-2014 Standard: put 96GB RAM in the box
- 2016+ Standard: put 144GB in the box
- Enterprise Edition: 256GB

Virtual machines: use the Tequila Technique

Cloud: dictated by your # of cores



1.2 p30



How to reduce PAGEIOLATCH

1. ~~Tune indexes, looking at missing indexes:~~

~~`sp_BlitzIndex @GetAllDatabases = 1`~~

2. ~~Tune queries:~~

~~`sp_BlitzCache @SortOrder = 'reads'`~~

3. ~~Add more memory~~

4. Make storage faster

But if we can't fix those...



Let's do the math.

1.2 p32



Reading millions of 8KB pages

```
SELECT * FROM dbo.Comments WHERE UserId = 26837;
```

150 %



Results



Messages

(150 rows affected)

Table 'Comments'. Scan count 9, logical reads 2765260,
physical reads 13, read-ahead reads 2762692,
lob logical reads 0, lob physical reads 0, lob read-ahead reads 0.

2,765,260 pages * 8KB each / 1,024 = 21.6GB

1.2 p33



If the users want it in 2 seconds...

How much data do we need to read?	21,600 MB
How long will the query take?	<u>2 seconds</u>
Divide the two to get throughput required:	10,800 MB per second



Larger example, data warehouse

How much data do we need to read?	250,000 MB (250GB)
How many seconds can we take?	<u>30 seconds</u>
Divide the two to get throughput required:	8,333 MB per second



Benchmark your current storage

CrystalDiskMark: BrentOzar.com/go/cdm

Get the zip edition without ads, save it to a share

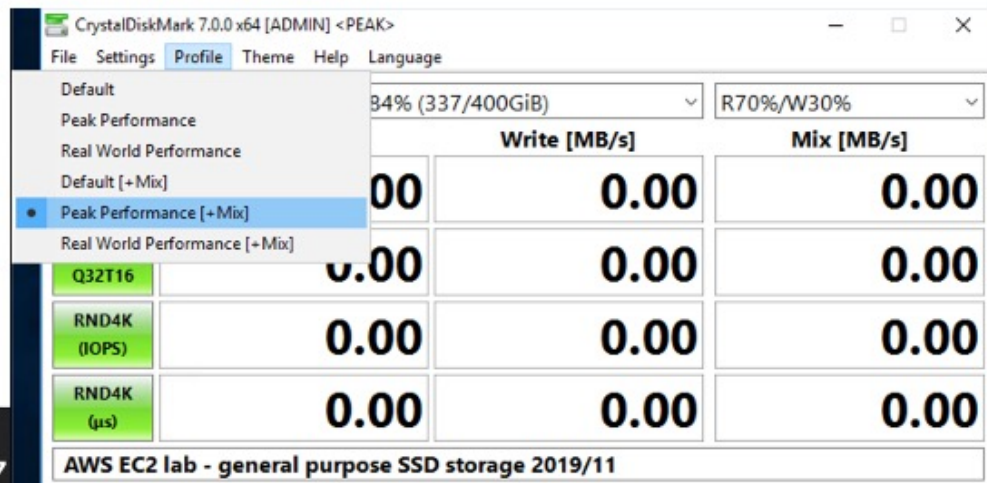


1.2 p36



Setting it up

Click Profile, Peak Performance + Mix



1.2 p37

Bonus points: choose your mix

CrystalDiskMark 7.0.0 x64 [ADMIN] <PEAK>

File Settings Profile Theme Help Language

All 1 1GiB S: 84% (337/400GiB)


	Read [MB/s]	Write [MB/s]	
SEQ1M Q8T1	0.00	0.00	R70%/W30%
RND4K Q32T16	0.00	0.00	R10%/W90%
RND4K (IOPS)	0.00	0.00	R20%/W80%
RND4K (μs)	0.00	0.00	R30%/W70%
			R40%/W60%
			R50%/W50%
			R60%/W40%
			R70%/W30%
			R80%/W20%
			R90%/W10%
			0.00

AWS EC2 lab - general purpose SSD storage 2019/11

1.2 p38



Settings across the top



1 = number of tests (1 = quick test, 9 = many tests)

1GiB = test file size (1GiB = quick test, 64GiB = in-depth test that won't just get cached)

S: = the drive letter or folder that you want to test



Interpreting the results

CrystalDiskMark 7.0.0 x64 [ADMIN] <PEAK>

File Settings Profile Theme Help Language

All 1 1GiB Z: 25% (450/1770GiB) R70%/W30%

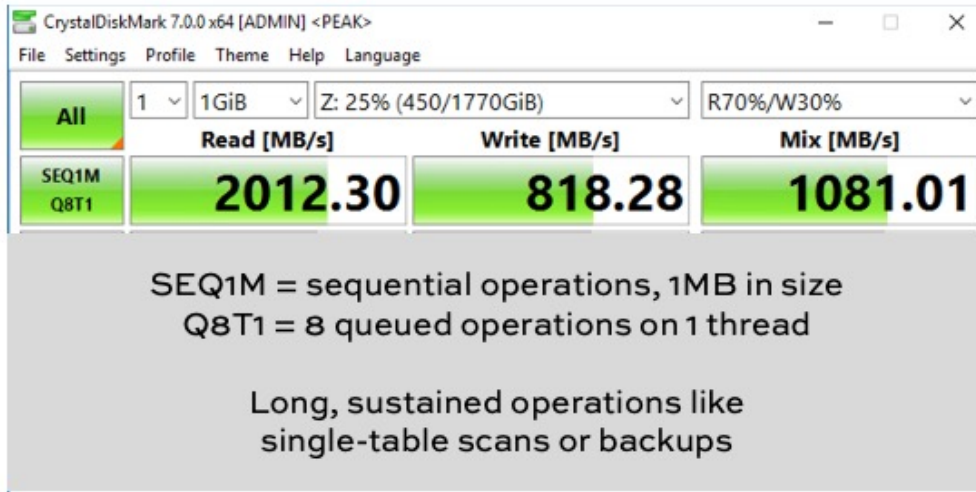
	Read [MB/s]	Write [MB/s]	Mix [MB/s]
SEQ1M Q8T1	2012.30	818.28	1081.01
RND4K Q32T16	1192.14	782.25	948.13
RND4K (IOPS)	291050.54	190979.00	231476.32
RND4K (μs)	1638.58	2625.50	2205.37

AWS EC2 lab - ephemeral local storage 2019/11

1.2 p40



Interpreting the results



1.2 p41



Interpreting the results

RND 4K= random operations, tiny 4KB size
Q32T16 = 32 queued operations on 16 threads: busy!

Kinda like an active, busy OLTP server or TempDB

RND4K Q32T16	1192.14	782.25	948.13
RND4K (IOPS)	291050.54	190979.00	231476.32
RND4K (µs)	1638.58	2625.50	2205.37

AWS EC2 lab - ephemeral local storage 2019/11



Slower storage

CrystalDiskMark 7.0.0 x64 [ADMIN] <PEAK>

File Settings Profile Theme Help Language

All

1

1GiB

S: 84% (337/400GiB)

R70%/W30%

	Read [MB/s]	Write [MB/s]	Mix [MB/s]
SEQ1M Q8T1	214.29	214.51	270.92
RND4K Q32T16	12.90	12.88	13.00
RND4K (IOPS)	3149.41	3143.31	3173.83
RND4K (µs)	159771.71	159968.11	158372.80

AWS EC2 lab - general purpose SSD storage 2019/11

1.2 p43



Click File, Copy, and paste into Notepad

```
-----
CrystalDiskMark 7.0.0 x64 (C) 2007-2019 hiyohiyo
Crystal Dew World: https://crystalmark.info/
-----
* MB/s = 1,000,000 bytes/s [SATA/600 = 600,000,000 bytes/s]
* KB = 1000 bytes, KiB = 1024 bytes

[Read]
Sequential 1MiB (Q= 8, T= 1): 214.294 MB/s [ 204.4 IOPS] < 38947.38 us>
Random 4KiB (Q= 32, T=16): 12.900 MB/s [ 3149.4 IOPS] <159771.71 us>

[Write]
Sequential 1MiB (Q= 8, T= 1): 214.515 MB/s [ 204.6 IOPS] < 38829.04 us>
Random 4KiB (Q= 32, T=16): 12.875 MB/s [ 3143.3 IOPS] <159968.11 us>

[Mix] Read 70%/Write 30%
Sequential 1MiB (Q= 8, T= 1): 270.915 MB/s [ 258.4 IOPS] < 30766.38 us>
Random 4KiB (Q= 32, T=16): 13.000 MB/s [ 3173.8 IOPS] <158372.80 us>

Profile: Peak
Test: 1 GiB (x1) [Interval: 5 sec] <DefaultAffinity=DISABLED>
Date: 2019/11/23 19:34:10
OS: Windows Server 2016 Datacenter (Full installation) [10.0 Build 14393] (x64)
Comment: AWS EC2 lab - general purpose SSD storage 2019/11
```

If the users want it in 2 seconds...

How much data do we need to read?	21,600 MB
How long will the query take?	<u>2 seconds</u>
Divide the two to get throughput required:	10,800 MB per second



This explains our performance

How much data do we need to read?	21,600 MB
But if the query is taking...	<u>11 seconds</u>
Then our throughput is only around...	1,963 MB per second
If the users want it to finish in...	5 seconds
Then our storage needs to read this fast...	21,600 MB <u>/ 5 seconds</u> 4,320 MB/sec



This simple formula drives everything.

1.2 p47



**Data you have to read
/ How fast you read
Minimum query runtime**

1.2 p48



Red herrings

Buffer pool extensions

In-memory OLTP (Hekaton)

Columnstore indexes for transactional workloads

Fast disks connected over shared network storage

1.2 p49



Simple SAN path (incl. cloud)



Server



HBA/NIC



Cable



Switch



Cable



Shared Storage Gear

1.2 p50



Path Speeds

1Gb iSCSI

2Gb Fiber Channel

4Gb Fiber Channel

8Gb Fiber Channel

10Gb iSCSI

Bandwidth Reference

Fastest Time to
Transfer 1GB

Theoretical
Max Bandwidth

USB 2.0	17 Seconds	60 MB/Sec	_____
1 Gb iSCSI	8 Seconds	125 MB/Sec	=====
4 Gb Fibre Channel	2.4 Seconds	425 MB/Sec	██████████
USB 3.0	1.7 Seconds	600 MB/Sec	██████████
SATA Revision 3.0 (6Gb)	1.3 Seconds	750 MB/Sec	██████████
8 Gb Fibre Channel	1.2 Seconds	850 MB/Sec	██████████
10 Gb iSCSI	800 Milliseconds	1.25 GB/Sec	██████████

1.2 p51

SQL Server handles more than you think

Maximum Consumption Rate (MCR):

How fast each core can consume data from storage

Starting point: 200MB/sec
of sequential reads per core

- 2 quad cores: 1,600MB/sec
- 4 quad cores: 3,200MB/sec
- 4 ten cores: 8,000MB/sec

Bandwidth Reference

	Fastest Time to Transfer 1GB	Theoretical Max Bandwidth	
USB 2.0	17 Seconds	60 MB/Sec	▬
1 Gb iSCSI	8 Seconds	125 MB/Sec	▬
4 Gb Fibre Channel	2.4 Seconds	425 MB/Sec	▬
USB 3.0	1.7 Seconds	600 MB/Sec	▬
SATA Revision 3.0 (6Gb)	1.3 Seconds	750 MB/Sec	▬
8 Gb Fibre Channel	1.2 Seconds	850 MB/Sec	▬
10 Gb iSCSI	800 Milliseconds	1.25 GB/Sec	▬

1.2 p52



But your VMs get a lot less



Server



HBA/NIC



Cable



Switch



Cable



Shared Storage Gear

Because we pack
a lot of clowns in
this car.

1.2 p53



How to reduce PAGEIOLATCH

1. Tune indexes, looking at missing indexes:
`sp_BlitzIndex @GetAllDatabases = 1`
2. Tune queries:
`sp_BlitzCache @SortOrder = 'reads'`
3. Add more memory
4. Make storage faster

