



**BRENT OZAR**  
UNLIMITED®

## Poison Wait: THREADPOOL

Like Deadpool, but much less funny.

3.1 p1

## Agenda

We already covered running a lot of normal queries:  
**SOS\_SCHEDULER\_YIELD**

Running an abnormal query: blocking

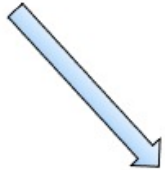
Running an abnormal query AND a lot of normal ones:  
**THREADPOOL**



## Flashback to CPU scheduling

**What's Running Now**

**SELECT \***  
**FROM dbo.Restaurants**  
*(By Brent)*



**What's Waiting (Queue)**

**SELECT \***  
**FROM dbo.SoccerClubs**  
*(By Richie)*

**SELECT \***  
**FROM dbo.Resorts**  
*(By Erika)*

3.1 p3



But the instant my query needs something that isn't in cache, like I need to wait for a locked page or I need to wait for something to come back from disk, I go to the back of the line.

## SOS\_SCHEDULER\_YIELD

A task uses CPU for up to 4 milliseconds straight

At that point, it yields the CPU scheduler

- Counts as 1 SOS\_SCHEDULER\_YIELD wait
- The clock starts on the wait time
- The clock stops when it can get back on a CPU

The longer the average wait time,  
the worse off your CPUs are.



## Some admins: “\\_(ツ)\_/” at SOS\_SCHEDULER\_YIELD

“Our CPU load is high, but SQL still works.”

“We use virtualization, and we expect high CPU.”

“We paid a lot for these licenses, and we want to use them, not have them sitting around idle.”

And that’s okay...until it’s not.

3.1 p5



Let's run an  
*ab*normal query.

3.1 p6



## Take out a lock

```
BEGIN TRAN  
UPDATE dbo.Users  
SET Reputation = 1000000  
WHERE Id = 26837;
```



## When COUNT(\*) tries to run...

It starts counting rows, but ...

It gets to User Id 26837 and stops

Which means it stops consuming CPU

But it's still consuming worker threads  
while it waits to move forward





## sp\_BlitzFirst warns you.

sp\_BlitzFirst

150 %

Results Messages

	Priority	FindingsGroup	Finding	QueryText	Details
1	0	sp_BlitzFirst 2018-0...	From Your Community Volunteers	NULL	<a href="#">Click To See Details - We hope you</a>
2	1	Query Problems	Long-Running Query Blocking Others	BEGIN TRAN UPDATE dbo.Users SET Reputation = 1...	<a href="#">Click To See Details - Query in Sta</a>
3	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
4	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
5	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
6	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
7	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
8	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
9	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
10	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
11	1	Query Problems	Long-Running Query Blocking Others	SELECT COUNT(*) FROM dbo.Users;	<a href="#">Click To See Details - Query in Sta</a>
12	50	Query Problems	Plan Cache Erased Recently	NULL	<a href="#">Click To See Details - The oldest q</a>
13	200	Wait Stats	CXPACKET	NULL	<a href="#">Click To See Details - For 3126 sec</a>
14	200	Wait Stats	LCK_M_S	NULL	<a href="#">Click To See Details - For 48 secur</a>

# CPU is idle – no one is working.

## CPU

Intel(R) Core(TM) i7-6920HQ CPU @ 2.90GHz

% Utilization over 60 seconds

100%



Utilization	Speed	Maximum speed:	2.90 GHz
2%	2.90 GHz	Sockets:	4
Processes	Threads	Virtual processors:	4
49	1122	Virtual machines:	Yes
	21214	L1 cache:	N/A

3.1 p10



## Blocking usually looks like this.

We have ugly queries running, and normally they need CPU, but they also need locks.

Our monitoring tools would alert us that we have large numbers of long-blocked queries.

We can still run diagnostic queries just fine.

In most shops, this is where alarm bells get raised.

*But not everywhere.*



Let's run an  
*ab*normal query,  
plus LOTS of  
normal queries.

3.1 p12



## Leave the blocking query running

```
BEGIN TRAN  
UPDATE dbo.Users  
SET Reputation = 1000000  
WHERE Id = 26837;
```



## Start a big load test

Stop SQLQueryStress with **150 threads**.

What do your CPUs look like now?

What are your wait stats?

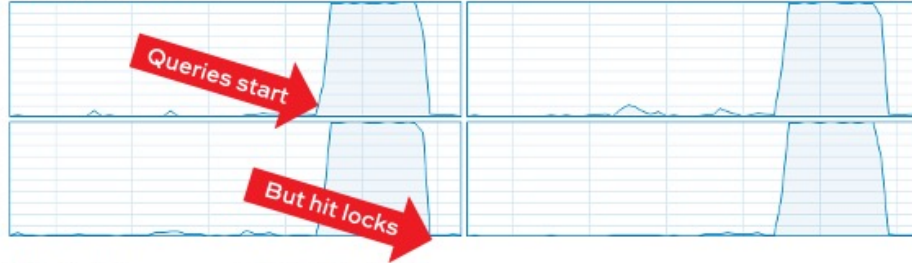


## CPU goes to 100%, then drops.

### CPU

Intel(R) Core(TM) i7-6920HQ CPU @ 2.90GHz

% Utilization over 60 seconds



Utilization	Speed	Maximum speed:	2.90 GHz
2%	2.90 GHz	Sockets:	4
Processes	Threads	Virtual processors:	4
49	1290	Virtual machines:	Yes
	Handles	LT cache:	N/A
	21666		

3.1 p15



## Simple DMV queries are failing

```
/* Overall workers */
SELECT s.cpu_id, w.scheduler_address, COUNT(*) AS workers
FROM sys.dm_os_workers w
INNER JOIN sys.dm_os_schedulers s ON w.scheduler_address = s.scheduler_address
WHERE s.status = 'VISIBLE ONLINE'
GROUP BY s.cpu_id, w.scheduler_address
ORDER BY s.cpu_id, w.scheduler_address;
```

100 %

Messages

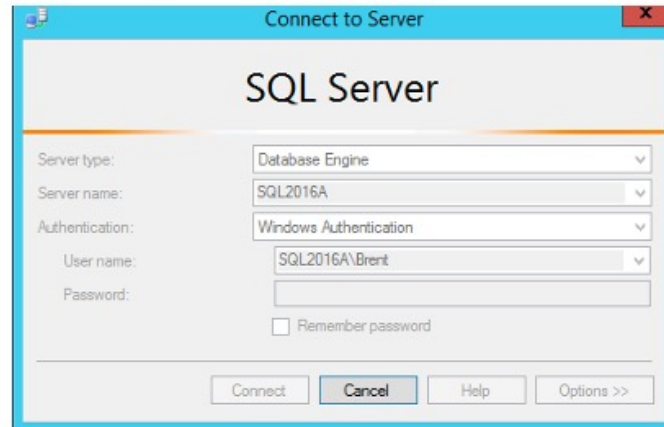
An error occurred while executing batch.  
Error message is: ExecuteReader requires an open and available Connection. The connection's current state is closed.

3.1 p16





**And when you try to reconnect...**



The screenshot shows the 'Connect to Server' dialog box. The title bar reads 'Connect to Server'. The main heading is 'SQL Server'. The dialog contains the following fields and options:

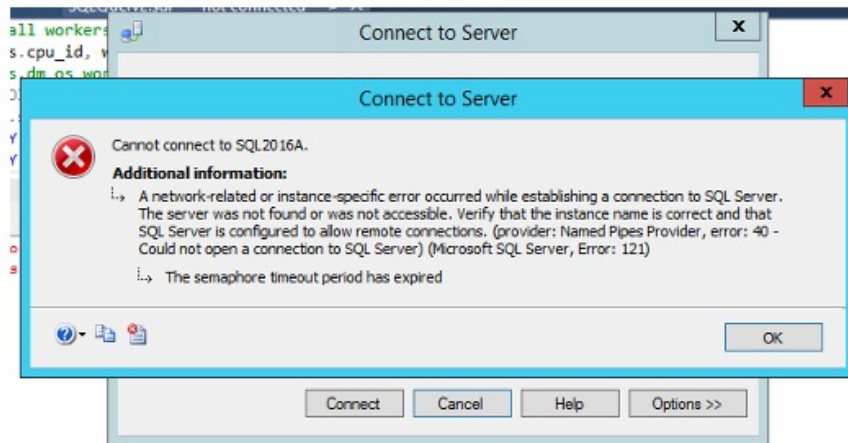
- Server type: Database Engine (dropdown)
- Server name: SQL2016A (dropdown)
- Authentication: Windows Authentication (dropdown)
- User name: SQL2016A\Brent (dropdown)
- Password: (text field)
- ☐ Remember password

At the bottom, there are four buttons: 'Connect', 'Cancel', 'Help', and 'Options >>'.

3.1 p17



## We start getting login failures



3.1 p18



## But Task Manager says we're idle.

### CPU

Intel(R) Core(TM) i7-6920HQ CPU @ 2.90GHz

% Utilization over 60 seconds

100%



Utilization	Speed	Maximum speed:	2.90 GHz
1%	2.90 GHz	Sockets:	4
Processes	Threads	Handles	Virtual processors:
48	1257	21789	4
		Virtual machines:	Yes
		L1 cache:	N/A

3.1 p19





It's like this.



## This is THREADPOOL.

We have ugly queries running, and normally they need CPU, but they also need locks.

Our monitoring tools can't connect to SQL Server: if we look at historical data, it looks like just holes of missing data.

You rarely see the THREADPOOL wait directly, but this is what it feels like: being unable to connect.

Let's dig deeper.



## How to dig deeper

At startup, the SQLOS sets aside:

- A scheduler (but not a CPU core)
- A worker thread mapped to that scheduler
- A small amount of RAM
- A dedicated port to connect to this scheduler

It's called the Dedicated Admin Connection (DAC).

Only usable by one session at a time

Learn more: [BrentOzar.com/go/dac](https://BrentOzar.com/go/dac)



## The Dedicated Admin Connection (DAC)

Reserves CPU resources for an administrator

- Only allows single threaded queries  
(this isn't for maintenance, it's just for emergencies)
- Only one person, one session can use it at a time

You need to enable this for remote access to use it

- "Remote" mean access over TCP/IP

This is a simple `sp_configure` change and does not require a restart



## Enable the remote DAC

```
/* look for pending changes before you go farther*/  
SELECT *  
FROM sys.configurations  
WHERE value <> value_in_use;  
GO  
  
/* This sets the value */  
exec sp_configure 'remote admin connections', 1  
GO  
/* This make it take effect */  
RECONFIGURE  
GO
```

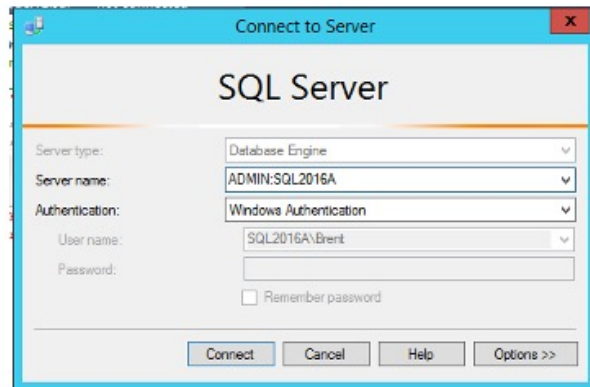




## Connecting with the DAC

Use the prefix “admin:”

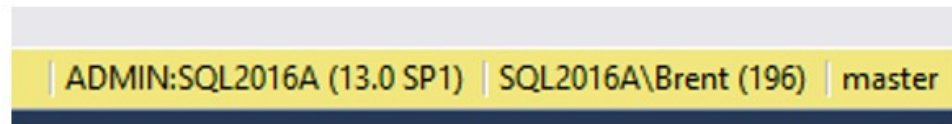
You may get an error.  
That’s Intellisense,  
trying to connect in a  
secret session, and  
failing miserably.



## How to tell you got the DAC

Look at the bottom right of your session window!

It should show that you're connected to admin:instancename if it worked



If someone else forgot to disconnect, you can find out who they are by using the query at <http://BrentOzar.com/go/DAC>



## Now you can troubleshoot.

Check sys.dm\_os\_workers: we do indeed have many, and we've hit (or exceeded) max worker threads

```
/* Overall workers */
SELECT s.cpu_id, w.scheduler_address, COUNT(*) AS workers
FROM sys.dm_os_workers w
INNER JOIN sys.dm_os_schedulers s ON w.scheduler_address = s.scheduler_address
WHERE s.status = 'VISIBLE ONLINE'
GROUP BY s.cpu_id, w.scheduler_address
ORDER BY s.cpu_id, w.scheduler_address;
```

100 %

Results Messages

	cpu_id	scheduler_address	workers
1	0	0x000000C6E0180040	137
2	1	0x000000C6E01A0040	135
3	2	0x000000C6E01C0040	136
4	3	0x000000C6E01E0040	138



## What's going on?

We have a lot of queries waiting on locks, but specifically how many?

sp\_whoIsActive

100 %

Results Messages

id	hh:mm:ss.mss	session_id	sql_text	login_name	wait_info	tempdb_allocations	tempdb_current	blocking_session_id	reads
1	00:00:20.55.496	58	<?query -- BEGIN TRAN, UPDATE dbo.Users SET Re...	SQL2016A\Brent	NULL	0	0	NULL	24
2	00:00:17.55.006	60	<?query -- SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1074957ms)LCK_M_S	0	0	56	4
3	00:00:17.55.006	61	<?query -- SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1074957ms)LCK_M_S	0	0	56	4
4	00:00:17.55.006	56	<?query -- SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1074950ms)LCK_M_S	0	0	58	0
5	00:00:17.54.870	62	<?query -- SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1074843ms)LCK_M_S	0	0	56	0
6	00:00:17.54.800	63	<?query -- SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1074777ms)LCK_M_S	0	0	56	0
7	00:00:17.54.766	64	<?query -- SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1074746ms)LCK_M_S	0	0	56	0
8	00:00:17.54.666	65	<?query -- SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1074639ms)LCK_M_S	0	0	56	0

Query executed successfully. ADMIN:SQL2016A (13.0 SP1) SQL2016A\Brent (196) master 138 rows

Not 150 or 151

3.1 p28

## We ran out of worker threads.

My 4-core VM's max: 512.

150 queries \* 4 threads each:  
need at least 600.

So not all 150 can run, and  
that's why new queries can't  
start either.

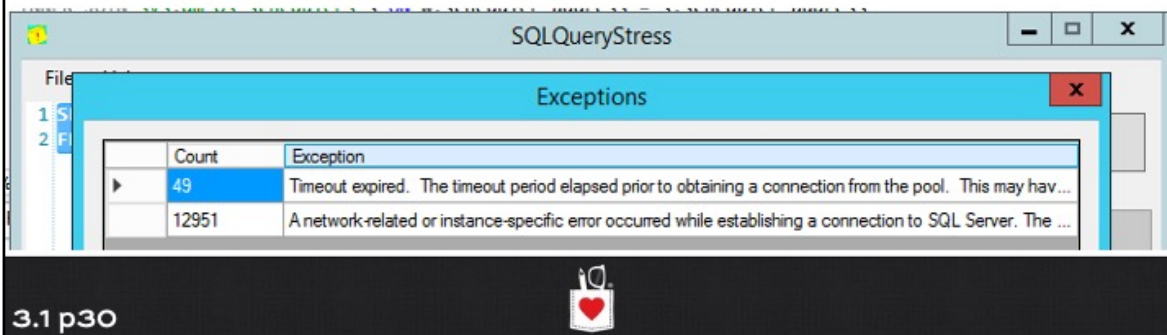
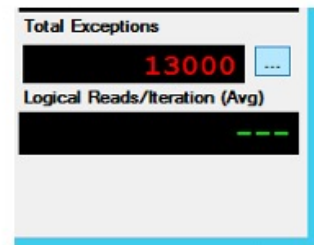
Number of CPUs	32-bit computer	64-bit computer
<= 4 processors	256	512
8 processors	288	576
16 processors	352	704
32 processors	480	960
64 processors	736	1472
128 processors	4224	4480
256 processors	8320	8576



## Check SQLQueryStress.

It's throwing thousands of exceptions.

Click on the ... button to see details.



## THREADPOOL symptoms

Normally, the server is alright. Not great, just alright.

But every now and then:

- We have a lot of active worker threads
- For some reason, they can't make progress (it's usually due to a blocking query)
- SQL Server runs out of available worker threads and starts recording THREADPOOL waits
- New queries can't get started, so
- Monitoring tools can't get data
- The server just feels frozen – but CPU use is low



## That wait is poison

That's why we report any occurrences of **THREADPOOL** wait as an urgent poison wait in tools like **sp\_Blitz**. If you're seeing *any* of it, the server is unusable during that time.

sp\_BlitzFirst @SinceStartup = 1

00 %

Results Messages

	Pattern	Sample Ended	Hours Sample	wait_type	Wait Time (Hours)	Per Core Per Hour
1	WAIT STATS	2017-02-23 09:12:54.3986015 -05:00	5.1	CXPACKET	92.3	0.0
2	WAIT STATS	2017-02-23 09:12:54.3986015 -05:00	5.1	LCK_M_S	8.6	0.0
3	WAIT STATS	2017-02-23 09:12:54.3986015 -05:00	5.1	SOS_SCHEDULER_YIELD	0.4	0.0
4	WAIT STATS	2017-02-23 09:12:54.3986015 -05:00	5.1	THREADPOOL	0.1	0.0

3.1 p32





## How to react when it strikes

Connect to the DAC

Run `sp_WhoIsActive` and look for a blocking firestorm

Collect the query, plan, and app sources for the lead blocker and the blocked queries

sp\_WhoIsActive @get\_plans = 1

100 %

Results Messages

	dd hh:mm:ss.mss	session_id	sql_text	login_name	wait_info	tempdb_allocations	tempdb_current	blocking_session_id
1	00:00:23.34.050	58	<?query - BEGIN TRAN UPDATE dbo.Users SET Re...	SQL2016A\Brent	NULL	0	0	NULL
2	00:00:20:33.570	60	<?query - SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1233535ms)LCK_M_S	0	0	56
3	00:00:20:33.570	61	<?query - SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1233535ms)LCK_M_S	0	0	56
4	00:00:20:33.570	56	<?query - SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1233538ms)LCK_M_S	0	0	58
5	00:00:20:33.434	62	<?query - SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1233421ms)LCK_M_S	0	0	56
6	00:00:20:33.364	63	<?query - SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1233349ms)LCK_M_S	0	0	56
7	00:00:20:33.330	64	<?query - SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1233324ms)LCK_M_S	0	0	56
8	00:00:20:33.330	55	<?query - SELECT COUNT(*) FROM StackOverflow.d...	SQL2016A\Brent	(1233313ms)LCK_M_S	0	0	56

## Should you kill the lead blocker?

That's a judgment call: rollbacks are single-threaded.

We don't judge here.

But we don't believe all queries have the right to life.

3.1 p34



# How to fix **THREADPOOL** long-term

3.1 p35



## Variables in the equation

Server-level:	Number of cores	4
	Number of worker threads SQL Server creates by default	512
	Cost Threshold for Parallelism	5
	How parallel a query will go	MAXDOP
Query-level:	Number of simultaneous queries	?
	Cost of those queries	?
	Lock requirements	?



## Adding more cores

Licensing makes this expensive, but it has two effects:

1. The amount of available CPU time goes up
2. The amount of worker threads may go up if we add enough cores

Number of CPUs	32-bit computer	64-bit computer
<= 4 processors	256	512
8 processors	288	576
16 processors	352	704
32 processors	480	960
64 processors	736	1472
128 processors	4224	4480
256 processors	8320	8576



## Will going to 8 cores help?

If we spend on licensing and hardware, worker threads goes up to 576.


That's not gonna be enough to run  $150 \times 4$  threads.

And the added CPU power won't help since we're waiting around on blocking.

Number of CPUs	32-bit computer	64-bit computer
<= 4 processors	256	512
8 processors	288	576
16 processors	352	704
32 processors	480	960
64 processors	736	1472
128 processors	4224	4480
256 processors	8320	8576



## Variables in the equation

Server-level:	Number of cores	4
How about this? 	Number of worker threads SQL Server creates by default	512
	Cost Threshold for Parallelism	5
	How parallel a query will go	MAXDOP
Query-level:	Number of simultaneous queries	?
	Cost of those queries	?
	Lock requirements	?



## Max Worker Threads

This `sp_configure` setting lets you override the number of worker threads SQL Server allows.

We could raise it into the thousands, but:

- Our blocked queries won't finish faster
- We're only buying a small amount of time
- If we get more end user `SELECT` loads, we're still gonna hit `THREADPOOL`
- We can run into memory starvation issues





## It gets even worse

What happens when our hundreds of simultaneous **SELECT** queries are actually working?

Each query will get even less CPU time per second, because when it yields the CPU scheduler, it's going to have to wait even longer to get back on.

When you see **THREADPOOL** waits, this isn't the fix.



## Variables in the equation

Server-level:	Number of cores	4
	Number of worker threads SQL Server creates by default	512
How about these?	Cost Threshold for Parallelism	5
	How parallel a query will go	MAXDOP
Query-level:	Number of simultaneous queries	?
	Cost of those queries	?
	Lock requirements	?



## Our queries are going parallel.

We could either:

- Raise Cost Threshold for Parallelism high enough that small queries don't go parallel, or
- Set Max Degree of Parallelism so they use less worker threads (or even only 1 thread)

To try it, you'll need to stop the load test first (because in-flight queries aren't affected.)



## MAXDOP 1 kinda fixes it.

Queries go single-threaded, and use less worker threads.

But they're still blocked. Users won't be happy.

The server is just now responsive again, and our monitoring tools show things like blocking.

When there's no blocking, our queries will SUCK.

This isn't really the right long-term fix  
(but can be good emergency duct tape.)



## Variables in the equation

Server-level:	Number of cores	4
	Number of worker threads SQL Server creates by default	512
	Cost Threshold for Parallelism	5
	How parallel a query will go	MAXDOP
Query-level:	Number of simultaneous queries	?
	Cost of those queries	?
	Lock requirements	?

*How about this?*



## True story

Our application runs on about a dozen app servers.

It's not fast enough, so we spun up more. About 300.

Every 5 minutes, they look for work in a queue table.

The queue table got big enough that the look-for-work query started going parallel.

Every 5 minutes, boom: SQL Server seems frozen.



## True solution

“Can we put an index on that queue table?”

“No, it’s a vendor app.”

“Can you do me a favor and  
shut down 290 of the app servers?”

“You’d better be right about this.”

“Wow, it’s amazing!  
So fast! Much wow!”

3.1 p47



## Variables in the equation

Server-level:	Number of cores	4
	Number of worker threads SQL Server creates by default	512
	Cost Threshold for Parallelism	5
	How parallel a query will go	MAXDOP
Query-level:	Number of simultaneous queries	?
How about this?	Cost of those queries	?
	Lock requirements	?

3.1 p48





## Reducing query costs

Run `sp_BlitzIndex`, and look for missing indexes with high impact (there's usually a very active table with no indexes on it at all, or a desperately needed index)

Run `sp_BlitzCache @SortOrder = 'reads'` and look at queries run frequently that desperately need indexes (and you may have to hand-tune the indexes here)

That usually makes `THREADPOOL` disappear.  
And it's free! This is my favorite fix. (No, I'm not free.)



## Variables in the equation

Server-level:	Number of cores	4
	Number of worker threads SQL Server creates by default	512
	Cost Threshold for Parallelism	5
	How parallel a query will go	MAXDOP
Query-level:	Number of simultaneous queries	?
	Cost of those queries	?
	Lock requirements	?

How about this?



## Reducing blocking issues

Finding the lead blocker during THREADPOOL issues is really tough.

Monitoring systems don't help because they can't gather data during these windows either.

When the SQL Server seems frozen:

- Connect to the DAC
- Run `sp_WhoIsActive` to find the lead blocker
- No, doing this in an Agent job won't work



## When you find the lead blocker

Can we make this transaction shorter?  
(Is it doing multiple things in one long block?)

Can we make this transaction faster?  
(Does it need an index to find the right rows?)

If those fail, can a different isolation level help?  
(We cover RCSI and snapshot in the LCK module.)





# Recap



## THREADPOOL means...

We've run out of worker threads, new queries can't start

SQL Server seems frozen, and monitoring tools fail

Windows OS CPU metrics seem totally fine, idle

Connect with the DAC, find the lead blocker

To fix it: tune indexes, queries to avoid the blocking storm  
and reduce parallelism requirements

Don't throw CPU or worker threads at it:  
it's expensive and ineffective under most circumstances

