



**BRENT OZAR**  
UNLIMITED®

# Poison Wait: RESOURCE\_SEMAPHORE

So low on memory, we can't even start a query

3.2 p1

# We often talk about these two memory pools in SQL Server

## Execution Plan cache

(metadata about how to run a query, stored in memory to be reused)

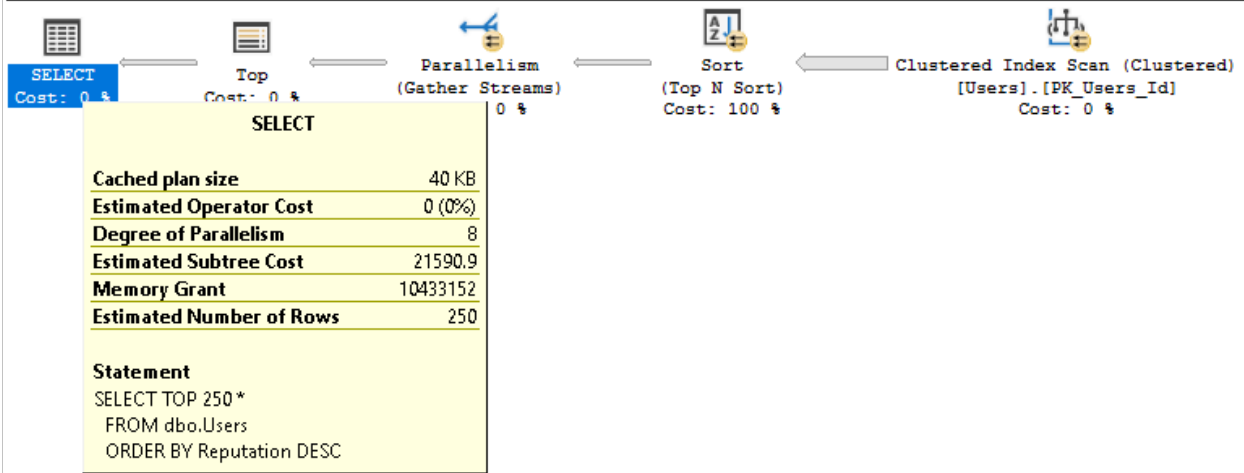
## Data cache aka “Buffer Pool”

(Data pages from your tables which have been read by queries, stored in memory to be reused)



## Queries also need memory grants

Query 1: Query cost (relative to the batch): 100%  
SELECT TOP 250 \* FROM dbo.Users ORDER BY Reputation DESC



## Memory grants

This is sometimes called “query workspace memory”

Queries need memories for things like:

- Join operators (these may secretly build temp objects that need a good chunk of memory)
- Sort operators (same thing)
- Parallelism

SQL Server estimates how much it needs for this before the query starts



## Actual Plan details

Granted memory =  
How much workspace memory  
the query actually got (in KB)

Requested memory =  
Ideal memory grant (KB)

Required memory =  
The *minimum* grant the query  
needed to get started (KB)

Max used memory =  
Just what it sounds like (KB)

\* SQL Server 2012 and higher

Estimated Subtree Cost	21590.9
Memory Grant	10433152
MemoryGrantInfo	
DesiredMemory	70376896
GrantedMemory	10433152
GrantWaitTime	0
MaxQueryMemory	10433152
MaxUsedMemory	1273360
RequestedMemory	10433152
RequiredMemory	5632
SerialDesiredMemory	70371728
SerialRequiredMemory	512
Optimization Level	FULL



# Query workspace memory has to come from somewhere

## Execution Plan cache

(metadata about how to run a query, stored in memory to be reused)

## Data cache

aka “Buffer Pool”

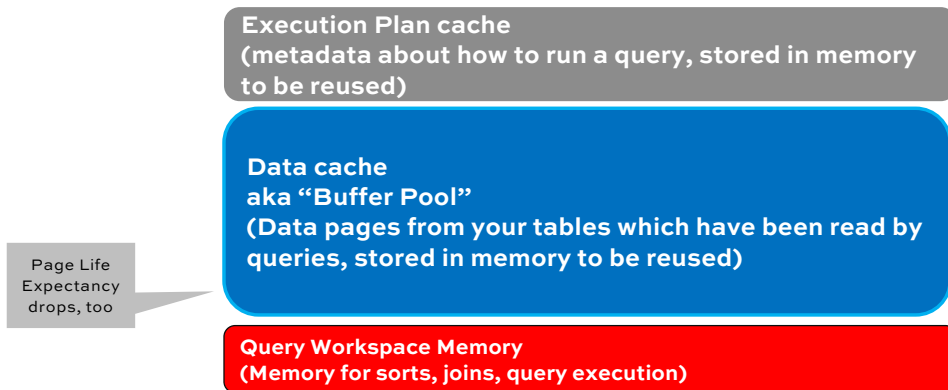
(Data pages from your tables which have been read by queries, stored in memory to be reused)

## Query Workspace Memory

(Memory for sorts, joins, query execution)



## If it needs to grow, it “steals” memory from the Buffer Pool



So does the execution plan cache, but whatever.

# And it can get bad. Really bad.

Execution Plan cache

Data cache  
aka "Buffer Pool"

Query Workspace Memory  
(Memory for sorts, joins, query execution)

3.2 p8



So does the execution plan cache, but whatever.



# Let's see it.

3.2 p9



## No skimping here.

Give your SQL Server plenty of memory.

```
EXEC sys.sp_configure N'max server memory (MB)', N'55000';  
EXEC sys.sp_configure N'cost threshold for parallelism', N'5';  
EXEC sys.sp_configure N'max degree of parallelism', N'0';  
GO  
RECONFIGURE  
GO  
USE StackOverflow;  
GO  
DropIndexes;  
GO
```



## BIG DATA

Let's change the data types on the Users table and make them big.

```
ALTER TABLE dbo.Users  
    ALTER COLUMN DisplayName NVARCHAR(400);  
ALTER TABLE dbo.Users  
    ALTER COLUMN Location NVARCHAR(1000);  
ALTER TABLE dbo.Users  
    ALTER COLUMN WebsiteUrl NVARCHAR(2000);
```

This will happen instantly: it's a metadata-only change.

The table's not actually getting bigger.













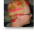











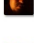













Secure | <https://stackoverflow.com/users?tab=Reputation&filter=all>

Questions Developer Jobs Tags Users Search...

11,625 +30 +56

Users reputation new users voters editors moderators

Type to find users:

 Jon Skeet Reading, United Kingdom 1007k • 512 • 7451 • 8120 c#, java, .net	 BalusC Amsterdam, Netherlands 781k • 266 • 2914 • 3043 java, jsp, jstl	 Darin Dimitrov Sofia, Bulgaria 781k • 199 • 2851 • 2639 c#, asp.net-mvc, asp.net-mvc-3	 VonC France 748k • 249 • 2294 • 2673 git, eclipse, github
 Hans Passant Madison, WI 739k • 101 • 1179 • 1896 c#, .net, workflows	 Marc Gravell Forest of Dean, United Kingdom 718k • 179 • 1996 • 2431 c#, .net, linq	 CommonWare Pennsylvania, United States 703k • 124 • 1719 • 1768 android, java, android-inet	 Gordon Linoff New York, United States 638k • 24 • 230 • 327 sql, mysql, sql-server
 SLaks New Jersey 626k • 128 • 1532 • 1869 c#, javascript, .net	 Martijn Pieters Cambridge, United Kingdom 622k • 105 • 1888 • 1967 python, python-3.x, python-2.7	 Greg Hewgill Christchurch, New Zealand 604k • 130 • 973 • 1124 git, python, c++	 T.J. Crowder United Kingdom 602k • 103 • 1943 • 1152 javascript, jquery, java
 Quentin United Kingdom 591k • 63 • 873 • 955 javascript, html, css	 paxdiablo 747k • 150 • 1141 • 1573 c, c++, bash	 Alex Martelli Sunnyvale, CA 603k • 107 • 975 • 1237 python, list, multidimensional-array	 dasblinkenlight United States 602k • 81 • 870 • 1082 c#, java, c++
 CMS Guatemala 535k • 145 • 811 • 797 javascript, jquery, c#	 Mark Byers Denmark 530k • 111 • 1255 • 1290 c#, python, mysql	 marc_s Bern, Switzerland 528k • 115 • 1016 • 1179 sql-server, sql, c#	 Ignacio Vazquez-Abrams 255k • 98 • 963 • 1094 python, django, bash
 JaredPar Redmond, WA 523k • 102 • 1022 • 1310 c#, .net, c++	 Jonathan Leffer California, USA 519k • 81 • 602 • 844 c, bash, unix	 Gufla Västervik, Sweden 514k • 68 • 508 • 829 c#, javascript, jquery	 Nick Craver Winston-Salem, NC 496k • 105 • 1155 • 1070 jquery, javascript, html
 Eric Lippert Seattle, WA 488k • 136 • 983 • 1833 c#, .net, c++	 David Heffernan Ulverston, United Kingdom 487k • 31 • 722 • 1104 depth, c++, c#	 JB Nizet Saint-Etienne, France 486k • 42 • 746 • 908 java, hibernate, jpa	 Felix Kling Sunnyvale, CA 486k • 105 • 790 • 822 javascript, scss, jquery, jquery
 unutbu 480k • 83 • 998 • 1092 python, pandas, numpy	 anubhava Bangalore, India 470k • 40 • 252 • 326 regex, bash, htaccess	 darius New York, NY 468k • 133 • 813 • 897 java, jquery, algorithm	 Gumbo Lim, Germany 465k • 80 • 636 • 732 php, json, content-type
 Stephen C Australia 463k • 63 • 500 • 849 java, multithreading, string	 BohClock Singapore 460k • 110 • 1077 • 1140 css, css-selectors, html	 Charles Bailey United Kingdom 456k • 67 • 522 • 586 git, c++, gpgignore	 Pascal Thivent 451k • 101 • 810 • 1040 java, maven-2, hibernate

weekly / monthly / quarterly reputation leagues

1 2 3 4 5 ... 224103 next

## Top users page

Let's build a bad query to get this page's data:

<https://stackoverflow.com/users?tab=Reputation&filter=all>

## The query is simple

```
SELECT TOP 250 *  
FROM dbo.Users  
ORDER BY Reputation DESC;
```

Test it. How fast is it?

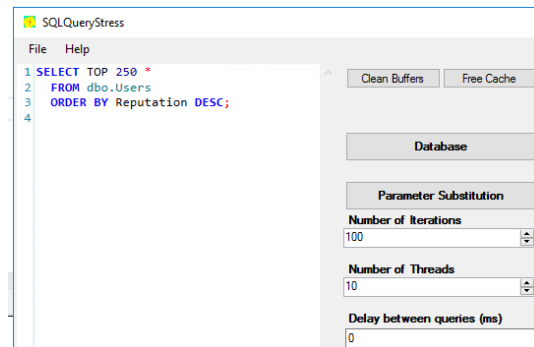
How many reads does it do?

How much CPU time does it take?



## Let's stress test it.

But just a little stress:  
run it in SQLQueryStress  
on just 10 threads,  
100 iterations per thread.



## A new wait type appears!

sp\_BlitzWho

Results							Messages
	run_date	elapsed_time	session_id	database_name	query_text	wait_info	
1	2018-02-15 00:37:17.460	0:00:00:10:000	62	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	CXPACKET (128 ms)	
2	2018-02-15 00:37:17.460	0:00:00:10:000	63	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	CXPACKET (11941 ms)	
3	2018-02-15 00:37:17.460	0:00:00:09:000	64	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	CXPACKET (21 ms)	
4	2018-02-15 00:37:17.460	0:00:00:09:000	65	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	RESOURCE_SEMAPHORE (39716 ms)	
5	2018-02-15 00:37:17.460	0:00:00:09:000	66	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	RESOURCE_SEMAPHORE (39632 ms)	
6	2018-02-15 00:37:17.460	0:00:00:09:000	67	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	RESOURCE_SEMAPHORE (39276 ms)	
7	2018-02-15 00:37:17.460	0:00:00:09:000	68	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	RESOURCE_SEMAPHORE (38988 ms)	
8	2018-02-15 00:37:17.460	0:00:00:09:000	69	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	RESOURCE_SEMAPHORE (38756 ms)	
9	2018-02-15 00:37:17.460	0:00:00:09:000	70	StackOverflow	SELECT TOP 250 * FROM dbo.Users ORDER BY R...	RESOURCE_SEMAPHORE (38584 ms)	

3.2 p15



# ExpertMode shows grants, too

Requires current patches of 2012 & newer

`sp_BlitzWho @ExpertMode = 1`

Results									
Messages									
grant_time	requested_memory_kb	grant_memory_kb	is_request_granted	required_memory_kb	query_memory_grant_used_memory_kb	ideal_memory_kb	wait_order	wait_time_ms	next_candidate_for_memory_grant
Feb 15 2018 12:37AM	10433152	10433152	Memory Request Granted	5632	775608	70376896	N/A	N/A	N/A
Feb 15 2018 12:37AM	10433152	10433152	Memory Request Granted	5632	2312	70376896	N/A	N/A	N/A
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	0	20984	Yes
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	2	20891	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	1	20906	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	3	10531	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	5	10281	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	4	10531	No

3.2 p16





## Performance counters show info

‘Pending’ memory grants indicate  
RESOURCE\_SEMAPHORE / workspace memory  
grant waits

```
SELECT object_name, counter_name, instance_name, cntr_value
FROM sys.dm_os_performance_counters
where
(counter_name like 'Active memory grants count%' or
counter_name like 'Pending memory grants count%' or
counter_name like 'Max request memory grant (KB)%')
and instance_name like 'default%'
and cntr_type=65792
ORDER by object_name, counter_name, instance_name
GO
```

	object_name	counter_name	instance_name	cntr_value
1	SQLServer:Resource Pool Stats	Active memory grants count	default	3
2	SQLServer:Resource Pool Stats	Pending memory grants count	default	5
3	SQLServer:Workload Group Stats	Max request memory grant (KB)	default	781384

## What we saw

A procedure wanted a high memory grant

- The amount it wanted was more than was allowed by default for a single request (based on default configuration), so it had to take a reduced grant

When we ran threads running 8 of this query simultaneously...

- SQL Server limited the number that could execute at once
- The waiting queries showed RESOURCE\_SEMAPHORE wait

Queries did complete, but they were very slow

- In the real world, this can easily cause application timeouts and the feeling that the SQL Server just isn't working



# What causes big memory grants?

## Queries processing a lot of data

- As data grows, queries have to do more work for hash joins, sorts, parallelism, etc.

## Running lots of queries

- As your user-base grows, more people using the SQL Server can run more large queries at the same time

## Optimization problems

- Joins with functions in them are very difficult for SQL Server to estimate and can be prone to high over-estimation
- Linked server queries (particularly to other databases) can vastly over-estimate the data returned from the remote side

## Bugs in SQL Server

- We ran into one case where SQL Server 2005 transactional replication was vastly over-estimating the rows used on internal queries



# How do I know if I'm having this problem?

3.2 p20



## At the server level

Wait stats (sys.dm\_os\_wait\_stats DMV)

- You see RESOURCE\_SEMAPHORE waits
- Even low values are a sign that things are sometimes going really wrong

Performance counters

- SQLServer: Memory Manager – Memory Grants Pending
- You want this counter to be at 0
- > 0 means someone has to wait to get a workspace memory grant and is being queued



## At the DMV details level

- Currently executing memory grants
- Historic memory grants (and reduced grants)

This info is in two DMVs:

- `sys.dm_resource_governor_resource_pools`
- `sys.dm_resource_governor_workload_groups`

You get the info even if...

- You haven't configured resource governor specially
- You're using Standard Edition



## Shown in sp\_BlitzWho

Also called by sp\_BlitzFirst @ExpertMode = 1

Shows currently running queries, like  
sp\_WhoIsActive, but with more tuning details:

```
sp_BlitzWho @ExpertMode = 1
```

Results									
grant_time	requested_memory_kb	grant_memory_kb	is_request_granted	required_memory_kb	query_memory_grant_used_memory_kb	ideal_memory_kb	wait_order	wait_time_ms	next_candidate_for_memory_grant
Feb 15 2018 12:37AM	10433152	10433152	Memory Request Granted	5632	775608	70376896	N/A	N/A	N/A
Feb 15 2018 12:37AM	10433152	10433152	Memory Request Granted	5632	2312	70376896	N/A	N/A	N/A
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	0	20984	Yes
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	2	20891	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	1	20906	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	3	10531	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	5	10281	No
Memory Not Granted	10433152	NULL	N/A	5632	NULL	70376896	4	10531	No

3.2 p23



## Shown in sp\_BlitzCache

@SortOrder = 'memory grant'

@SortOrder = 'unused grant'

```
sp_BlitzCache @SortOrder = 'memory grant'
```

150 %

Results Messages

	Average Rows	Minimum Memory Grant KB	Maximum Memory Grant KB	Minimum Used Grant KB	Maximum Used Grant KB	Average Max Memory Grant	Min Spills	Max Spills	Total Spills	Avg Spills
1	0.00	10433152	10433152	1273360	1273360	10433152.00	0	0	0	0.00
2	0.00	10433152	10433152	1273440	1273440	10433152.00	0	0	0	0.00
3	0.00	10433152	10433152	1273352	1273424	186306.2857	0	0	0	0.00
4	1.00	175016	175016	2888	2960	58338.6667	0	0	0	0.00
5	1.5714	167496	167496	2152	2224	23928.00	0	0	0	0.00
6	0.00	20768	20768	2576	2576	20768.00	0	0	0	0.00
7	0.00	19208	19208	1584	1584	19208.00	0	0	0	0.00
8	0.00	3984	3984	776	776	3984.00	0	0	0	0.00
-	-	-	-	-	-	-	-	-	-	-

3.2 p24





## And we show grants as warnings

### “Unused Memory Grant” warning

	Query Type	Warnings
YY R...	Statement	Parallel, Multiple Plans, Plan created last 4hrs, Expensive Sort, Row Goals
YY R...	Statement	Parallel, Multiple Plans, Plan created last 4hrs, Expensive Sort, Row Goals
YY R...	Statement	Parallel, Multiple Plans, Plan created last 4hrs, Expensive Sort, Row Goals
ALE...	Statement	Compilation Timeout, Plan Warnings, Function Join, Forced Serialization, Unused Memory Grant, Plan created last 4hrs, Table Variables, Long Running With Lo
ALE...	Statement	Compilation Timeout, Plan Warnings, Function Join, Forced Serialization, Unused Memory Grant, Plan created last 4hrs, Table Variables, Table Scans, Long Ru
bl.cr...	Statement	Compilation Timeout, Plan Warnings, Implicit Conversions, Function Join, Forced Serialization, Plan created last 4hrs, Table DML, Row Goals, MSTVFs
AS [I...	Statement	Compilation Timeout, Plan Warnings, Implicit Conversions, Function Join, Forced Serialization, Unused Memory Grant, Plan created last 4hrs, Row estimate mism
sp.cr...	Statement	Compilation Timeout, Plan Warnings, Implicit Conversions, Function Join, Forced Serialization, Plan created last 4hrs, Row estimate mismatch, MSTVFs
tbl I...	Statement	Forced Serialization, Unused Memory Grant, Plan created last 4hrs
tbl I...	Statement	Forced Serialization, Unused Memory Grant, Plan created last 4hrs

3.2 p25



## What can fix the issue?

1. Tuning queries and/or indexes
2. Taming the problem with Resource Governor (Enterprise Edition only)
3. Adding more memory



## 1) Tuning queries and indexes

This is the best long term solution

Sometimes queries are vastly OVER estimating their memory grant

- This can be due to code problems
- This can be due in part to bad statistics

Find the queries running when this wait starts happening

Look at which ones have the big memory grant, and tune them

Tools to help:

- **sp\_BlitzCache @SortOrder = 'memory grant'**
- sp\_BlitzWho shows queries running now
- A monitoring tool that trends wait stats, running queries



## 2) Resource Governor

You can ...

- Classify queries into groups
- Control memory allocations by group

But:

- Wouldn't a better way to raise the amount of memory grant available be to add more memory in a vast majority of cases?
- If we're classifying queries, are we making them *faster*?

Downsides:

- If you do this wrong, you make everything slower
- This is an Enterprise feature, more expensive than memory



### 3) Adding more memory

Check if your instance is under provisioned

Memory is a cheap way to improve performance

It's not just important for caching data pages

It's also vitally important to make sure your queries can do sorts, joins, etc



# Recap

3.2 p30



## RESOURCE\_SEMAPHORE

Queries need a memory grant to start running

SQL Server can't grant it due to bad memory pressure

Find the queries causing it:

- `sp_BlitzCache @SortOrder = 'memory grant'`
- `sp_BlitzWho`

Tuning queries & indexes is usually the cheapest fix

Got less than 32-64GB RAM? May need to add some.

